Introduction to Refinement Checking and FDR

Andrzej Filinski DIKU

Extreme Multiprogramming 12 January 7, 2014

Overview

- Brief introduction to refinement checking
 - An alternative alphabet model for CSP
 - Generalized concurrent composition
 - Refining specifications to implementations
 - trace refinements
 - failure refinements
- ► A look at the FDR2/FDR3 tool
 - From Oxford Computing Lab / Formal Systems Europe
 - Free academic version (for Unix-like platforms only)
- ▶ Not a requirement to use FDR for exam
 - It will not construct the required derivations for you
 - But may (or may not) be useful to sanity-check your results
 Beware of bugs in the above code; I have only proved it
 correct, not tried it. Donald E. Knuth

CSP as a machine-processable notation

- ▶ CSP can be used as foundation of a programming notation.
 - ▶ Often entails considerable syntactic and semantic adjustments.
 - Especially if embedding CSP inside existing language.
 - ▶ → programming track of the course.
- ► Also a tool for machine-supported specification and verification of general concurrent systems.
 - hardware design, communication protocols, parallel-algorithm skeletons, ...
- ▶ Machine-readable notation (CSP_M) very close to CSP book
 - ▶ Plain-ASCII syntax: $P \upharpoonright Q$ for $P \sqcap Q$, P(x) for P_x , etc.
 - Some (but surprisingly few) restrictions on allowable forms of process definitions.
 - Main difference: somewhat different formal treatment of alphabets and concurrent composition.

Alphabet models

- ▶ In traditional CSP (Hoare), all processes must have explicitly specified alphabets.
 - ► Alphabets don't usually matter for actual process behavior, except in concurrent composition (||).
 - ► This formulation leads to particularly simple equational laws for ||.
- ▶ In alternative formulation (Roscoe), alphabets are only explicitly mentioned in concurrent composition (and a few other places, e.g. CHAOS_A)
 - More convenient for some purposes.
 - But muddies waters a bit, because concurrency can now by itself introduce some nondeterminism.
- ▶ Generalized concurrent composition: $P \parallel Q$.
 - ► Concrete syntax: P [| B |] Q, much nicer layout.
 - Processes synchronize (only) on events from set B.

Generalized concurrent composition

- $ightharpoonup P \parallel Q$ like $P \parallel Q$ in Hoare CSP, but
 - Event in B can happen iff both P and Q are willing to participate
 - (and then both processes advance)
 - ► Event not in B can happen iff either P or Q willing to participate
 - (and then only the participating process advances)
 - ▶ If both are willing, the event happens *twice* (not necessarily consecutively).
- ▶ $P \parallel Q$ (in Hoare) corresponds to taking $B = \alpha P \cap \alpha Q$
- ▶ $P \parallel Q$ (interleaving) corresponds to taking $B = \{\}$.
- ▶ Other choices for *B* do not necessarily correspond to anything simple in Hoare CSP.
- ► See full details in Roscoe: *The Theory and Practice of Concurrency*, Chapter 2.

Other significant differences between book and CSP_M notation

► No explicit syntactic construct for enumerated choice; use general choice instead:

$$(x \to P \mid y \to Q) = (x \to P) \square (y \to Q).$$

- Somewhat different treatment of divergence and CHAOS.
 - Shouldn't be an issue as long as all processes are guarded
- ▶ Allows indexed nondeterministic choice, potentially infinitary: $\prod_{x \in A} P(x)$.
- A few others, but mainly in parts of CSP we did not cover. See FDR manual
- ► [Live demo 1: HW1 in FDR; equivalence and deadlock checking]

Trace refinement

- ▶ Already have notion of satisfaction, P sat $\phi(tr)$.
 - ▶ Says that every (finite) trace tr of P satisfies logical formula ϕ .
 - lacktriangledown ϕ expressed in "mathematics" (sets, sequences, functions, etc.).
 - Very general, but often hard to verify mechanically without substantial human assistance.
- ► A simpler notion of specification satisfaction is often sufficient: *trace refinement*.
 - ▶ Set of allowable traces expressed as *specification* process *S*.
 - Set of actual traces determined by implementation process 1.
 - ▶ Relation $S \sqsubseteq_T I$, "I trace-refines S",

$$S \sqsubseteq_{\mathrm{T}} I \iff traces(I) \subseteq traces(S)$$

- Motivation: refining nondeterministic specification into (more) deterministic implementation
 - ▶ Picking one concrete behavior from space of allowable ones.

Trace refinement disallows patently incorrect behaviors

Ex: Specification and implementations of vending machines

```
SVM = coin \rightarrow ((sprite \rightarrow SVM) \sqcap (pepsi \rightarrow STOP))

I1 = coin \rightarrow sprite \rightarrow I1

I2 = coin \rightarrow pepsi \rightarrow STOP

I3 = coin \rightarrow ((sprite \rightarrow I3) \sqcap (pepsi \rightarrow STOP))

I4 = coin \rightarrow I4

I5 = coin \rightarrow sprite \rightarrow pepsi \rightarrow STOP

I6 = coin \rightarrow pepsi \rightarrow I6
```

- ▶ Easy to check that for each of I1, I2, I3, have $SVM \sqsubseteq_T I$.
- ▶ But each of /4, /5, /6 has at least one trace that is not a valid trace of SVM.
- ► [Live demo 2: Trace-refinement checking]

Trace refinement is not always enough

Obvious problem: a deadlocked machine doesn't actively do anything wrong.

$$S = coin \rightarrow ((sprite \rightarrow S) \sqcap (pepsi \rightarrow STOP))$$

 $I1 = coin \rightarrow sprite \rightarrow STOP$
 $I2 = STOP$
 $I3 = coin \rightarrow STOP$

All these I also satisfy $S \sqsubseteq_T I$.

▶ More subtle problem: traces don't distinguish between internal and external choice (recall: $traces(P \sqcap Q) = traces(P \sqcap Q)$.)

$$SE = coin \rightarrow ((sprite \rightarrow SE) \square (pepsi \rightarrow STOP))$$

 $I1 = coin \rightarrow sprite \rightarrow I1$
 $I2 = coin \rightarrow ((sprite \rightarrow I2) \sqcap (pepsi \rightarrow STOP))$

Both / satisfy, $SE \sqsubseteq_T$ /, but neither actually allows customer to choose beverage.

Failures

- Recall: refusals(P) = all sets B s.t. P may deadlock if environment only offers B.
- ▶ **Def.** A failure of a process P is a pair (s, B), s.t. $s \in traces(P)$ and $B \in refusals(P / s)$.
- ▶ A failure (despite the name) does not indicate an *error* in *P*.
 - Simply means that P may (possibly correctly) deny a specific sequence of requests from environment.
 - ▶ E.g., $(\langle \rangle, \{sprite\})$ should be a failure of $VM = coin \rightarrow sprite \rightarrow VM$ (do not dispense free sprites!).
- ▶ **Def.** failure-refinement: $S \sqsubseteq_F I \Leftrightarrow failures(I) \subseteq failures(S)$.
- ▶ Stronger than trace-refinement: If $S \sqsubseteq_F I$ then also $S \sqsubseteq_T I$:
 - 1. Suppose $s \in traces(I)$
 - 2. then $(s, \{\}) \in failures(I)$ [since $\{\} \in refusals(P) \text{ always}$]
 - 3. hence $(s, \{\}) \in failures(S)$ [by assumption that $S \sqsubseteq_F I$)
 - 4. so in particular $s \in traces(S)$.

Failure refinement and deadlocks

Example

$$S = coin \rightarrow sprite \rightarrow S$$

 $I = coin \rightarrow STOP$

Traces:

$$traces(S) = \{\langle \rangle, \langle coin \rangle, \langle coin, sprite \rangle, \langle coin, sprite, coin \rangle, ...\}$$
$$traces(I) = \{\langle \rangle, \langle coin \rangle\}$$

Have $traces(I) \subseteq traces(S)$, so $S \sqsubseteq_T I$

► Failures:

$$failures(S) = \{(\langle \rangle, \{\}), (\langle \rangle, \{sprite\}), (\langle coin \rangle, \{\}), (\langle coin \rangle, \{coin\}), (\langle coin, sprite \rangle, \{\}), (\langle coin, sprite \rangle, \{sprite\}), ...\}$$

$$failures(I) = \{(\langle \rangle, \{\}), (\langle \rangle, \{sprite\}), (\langle coin \rangle, \{\}), (\langle coin \rangle, \{coin\}), (\langle coin \rangle, \{sprite\}), (\langle coin \rangle, \{coin, sprite\})\}$$

Here, e.g., $(\langle coin \rangle, \{sprite\})$ in failures(I), but not in failures(S), so $S \not\sqsubseteq_F I$.

Failure refinement and choices

- ▶ **Recall:** $refusals(P \sqcap Q) = refusals(P) \cup refusals(Q);$ $refusals(P \sqcap Q) = refusals(P) \cap refusals(Q).$
- **Example**: $(A = \{coin, sprite, pepsi\})$

$$\begin{array}{l} \textit{I1} = \textit{coin} \rightarrow \textit{sprite} \rightarrow \textit{STOP} \\ \textit{I2} = \textit{coin} \rightarrow ((\textit{sprite} \rightarrow \textit{STOP}) \sqcap (\textit{pepsi} \rightarrow \textit{STOP})) = \textit{SI} \\ \textit{I3} = \textit{coin} \rightarrow ((\textit{sprite} \rightarrow \textit{STOP}) \sqcap (\textit{pepsi} \rightarrow \textit{STOP})) = \textit{SE} \\ \textit{failures}(\textit{I1}) = \{(\langle \rangle, B) \mid B \subseteq \{s, p\}\} \cup \{(\langle c \rangle, B) \mid B \subseteq \{c, p\}\} \cup \{(\langle c, s \rangle, B) \mid B \subseteq \{s, p\}\} \cup \{(\langle c \rangle, B) \mid B \subseteq \{c, p\}\} \cup \{(\langle c \rangle, B) \mid B \subseteq \{s, p\}\} \cup \{(\langle c, s \rangle, B) \mid s \in \{s, p\}, B \subseteq A\} \\ \textit{failures}(\textit{I3}) = \{(\langle \rangle, B) \mid B \subseteq \{s, p\}\} \cup \{(\langle c \rangle, B) \mid B \subseteq \{c\}\} \cup \{(\langle c, s \rangle, B) \mid s \in \{s, p\}, B \subseteq A\} \\ \end{array}$$

So for all I, have $SI \sqsubseteq_F I$; but $SE \not\sqsubseteq_F I1$, $SE \not\sqsubseteq_F I2$.

[Live demo 3: Failure-refinement checking]

FDR

- Under continuous development since 1990; lots of other features.
 - "Concurrency workbench" for experimental development.
- Also includes systematic treatment of divergences
 - ▶ Name from Failures—Divergence Refinement
 - ▶ Implementation may be allowed to diverge, if spec does
 - Formal treatment is a bit nastier, when recursion isn't necessarily guarded.
- Can also have parameterized processes, channel communications.
 - For small alphabets/state sets, works as expected.
 - ▶ For large or infinite, need to be careful.
 - Verification may "time out" because it needs to check a huge (or infinite) number of cases
- Details in FDR manual.

Conclusion

- ▶ No more theory lectures!
 - Last programming lecture on Friday
- Course evaluation (handled outside of Absalon now): please complete!
- Final exam: will be out on Monday.