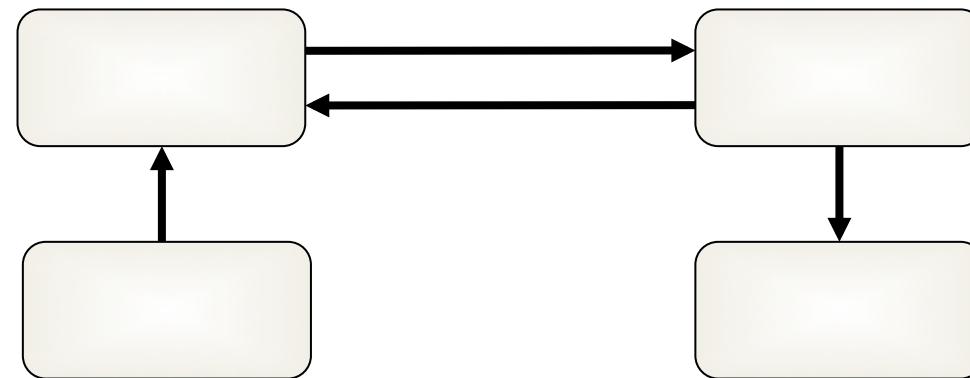


Distribution and Mobility

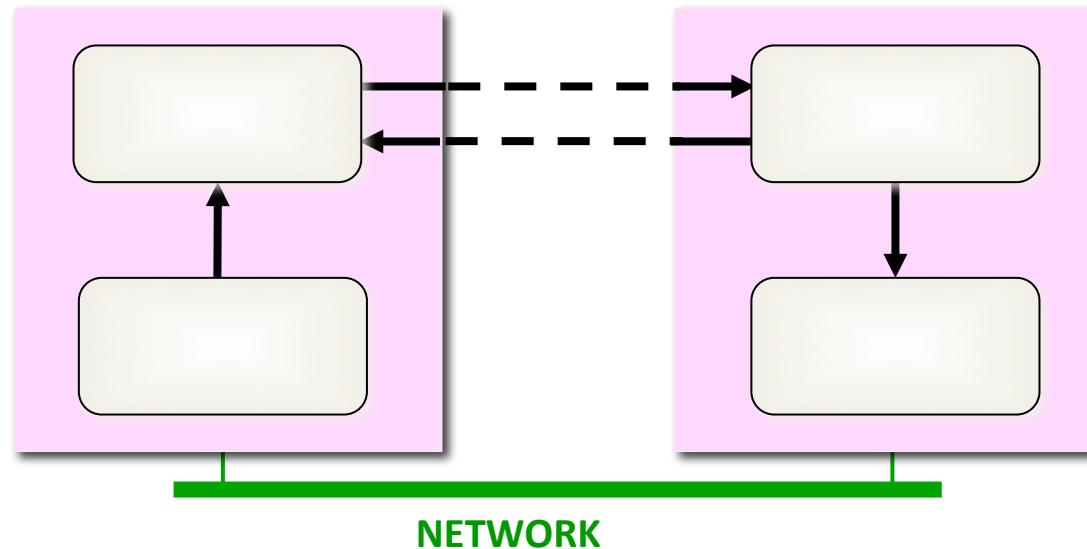
Distributed JCSP

- Want to use the same model for concurrent processes *whether or not they are on the same machine*:



Distributed JCSP

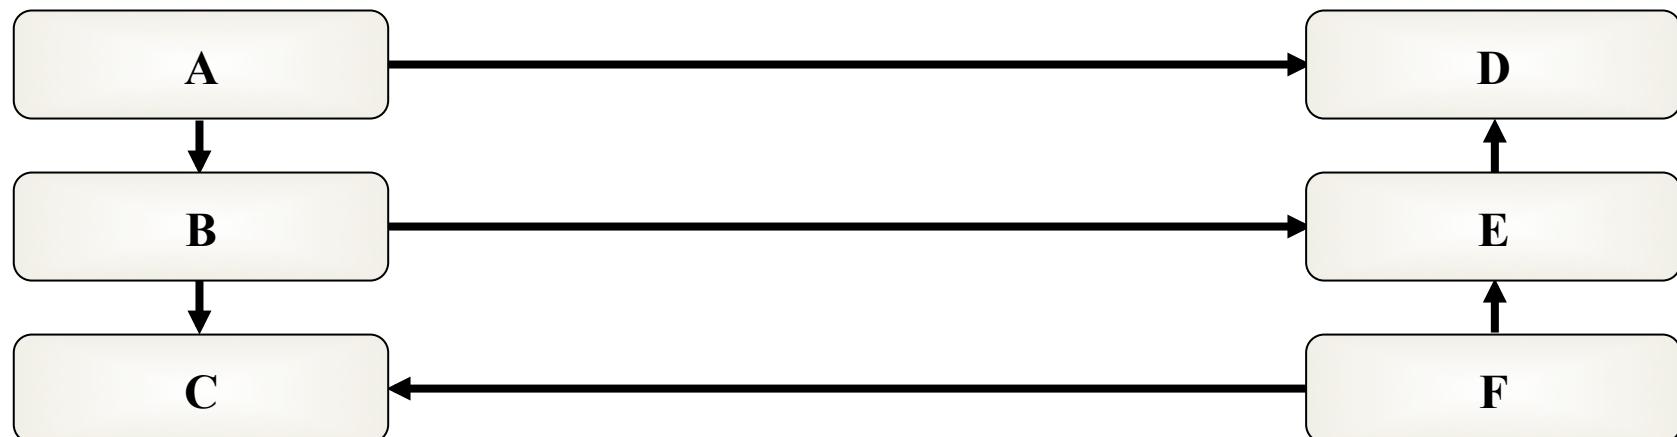
- Want to use the same model for concurrent processes *whether or not they are on the same machine*:



- Processes on different processing *nodes* communicate via *virtual channels*.

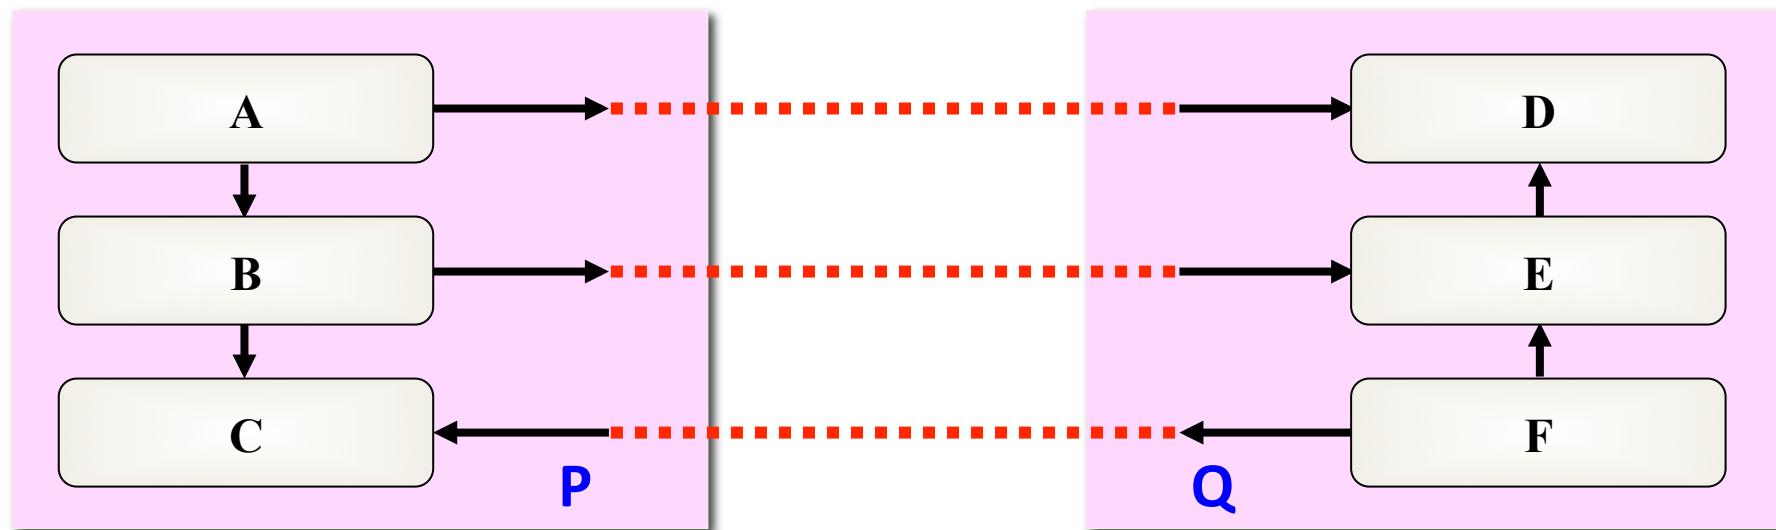
Logical Network

- Suppose a system contains processes **A, B, C, D, E** and **F**, communicating as shown below.
- There may be other processes and communication channels (but they are not relevant here).
- Suppose we want to distribute these processes over two processors ...



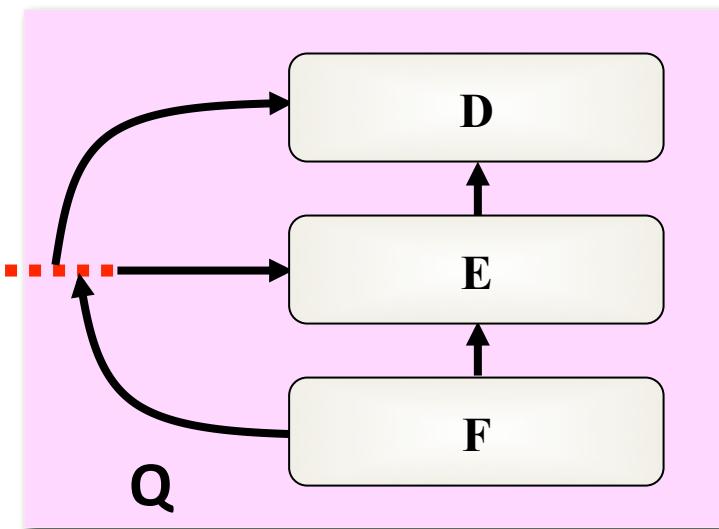
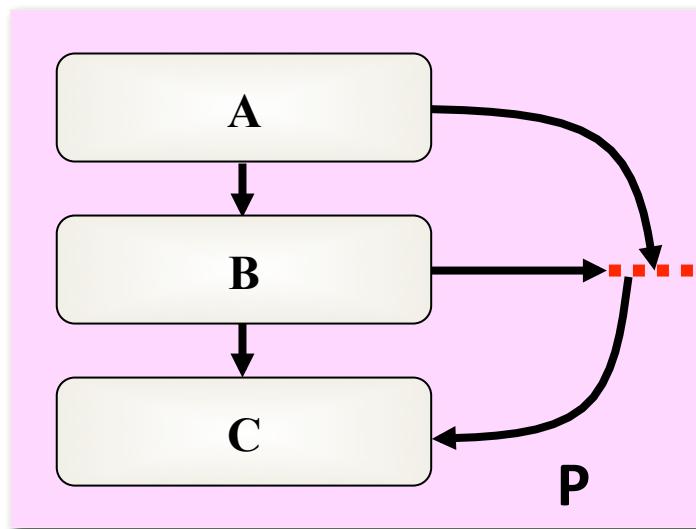
Physical Network

- Suppose we want to distribute these processes over two processors (**P** and **Q**, say) ...
- We could set up separate network links ...
- Or, since links may be a scarce resource, we could multiplex over a shared link ...



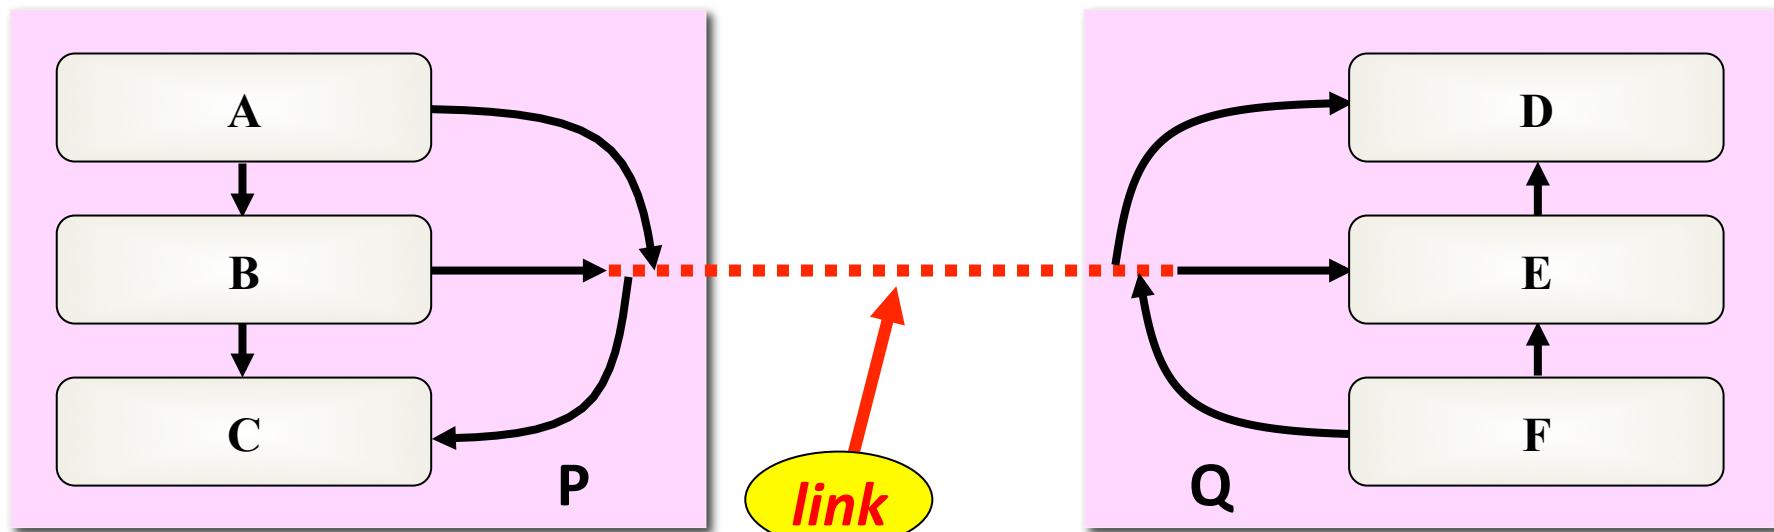
Physical Network

- Suppose we want to distribute these processes over two processors (**P** and **Q**, say) ...
- We could set up separate network links ...
- Or, since links may be a scarce resource, we could multiplex over a shared link ...



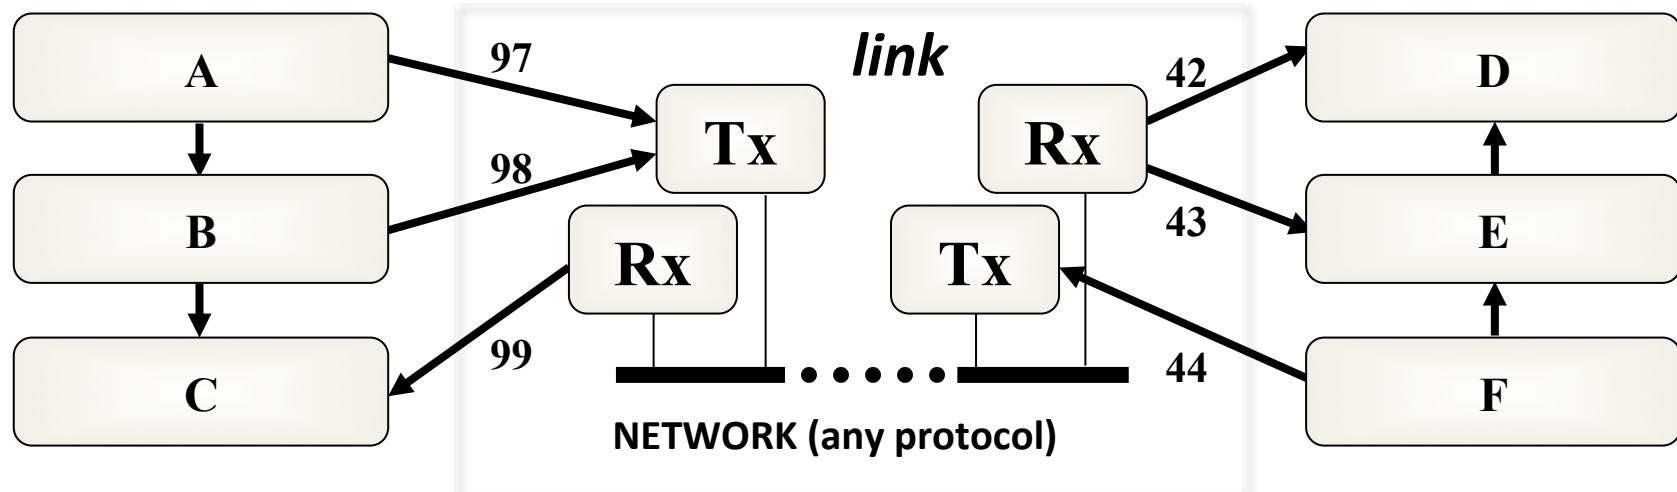
JCSP Links

- A connection between two processing nodes (*JVMs* in the context of **JCSP**) is called a *link*.
- Multiple channels between two nodes may use the same link – data is multiplexed in both directions.
- Links can ride on any network infrastructure (*TCP/IP, Firewire, 1355, ...*).

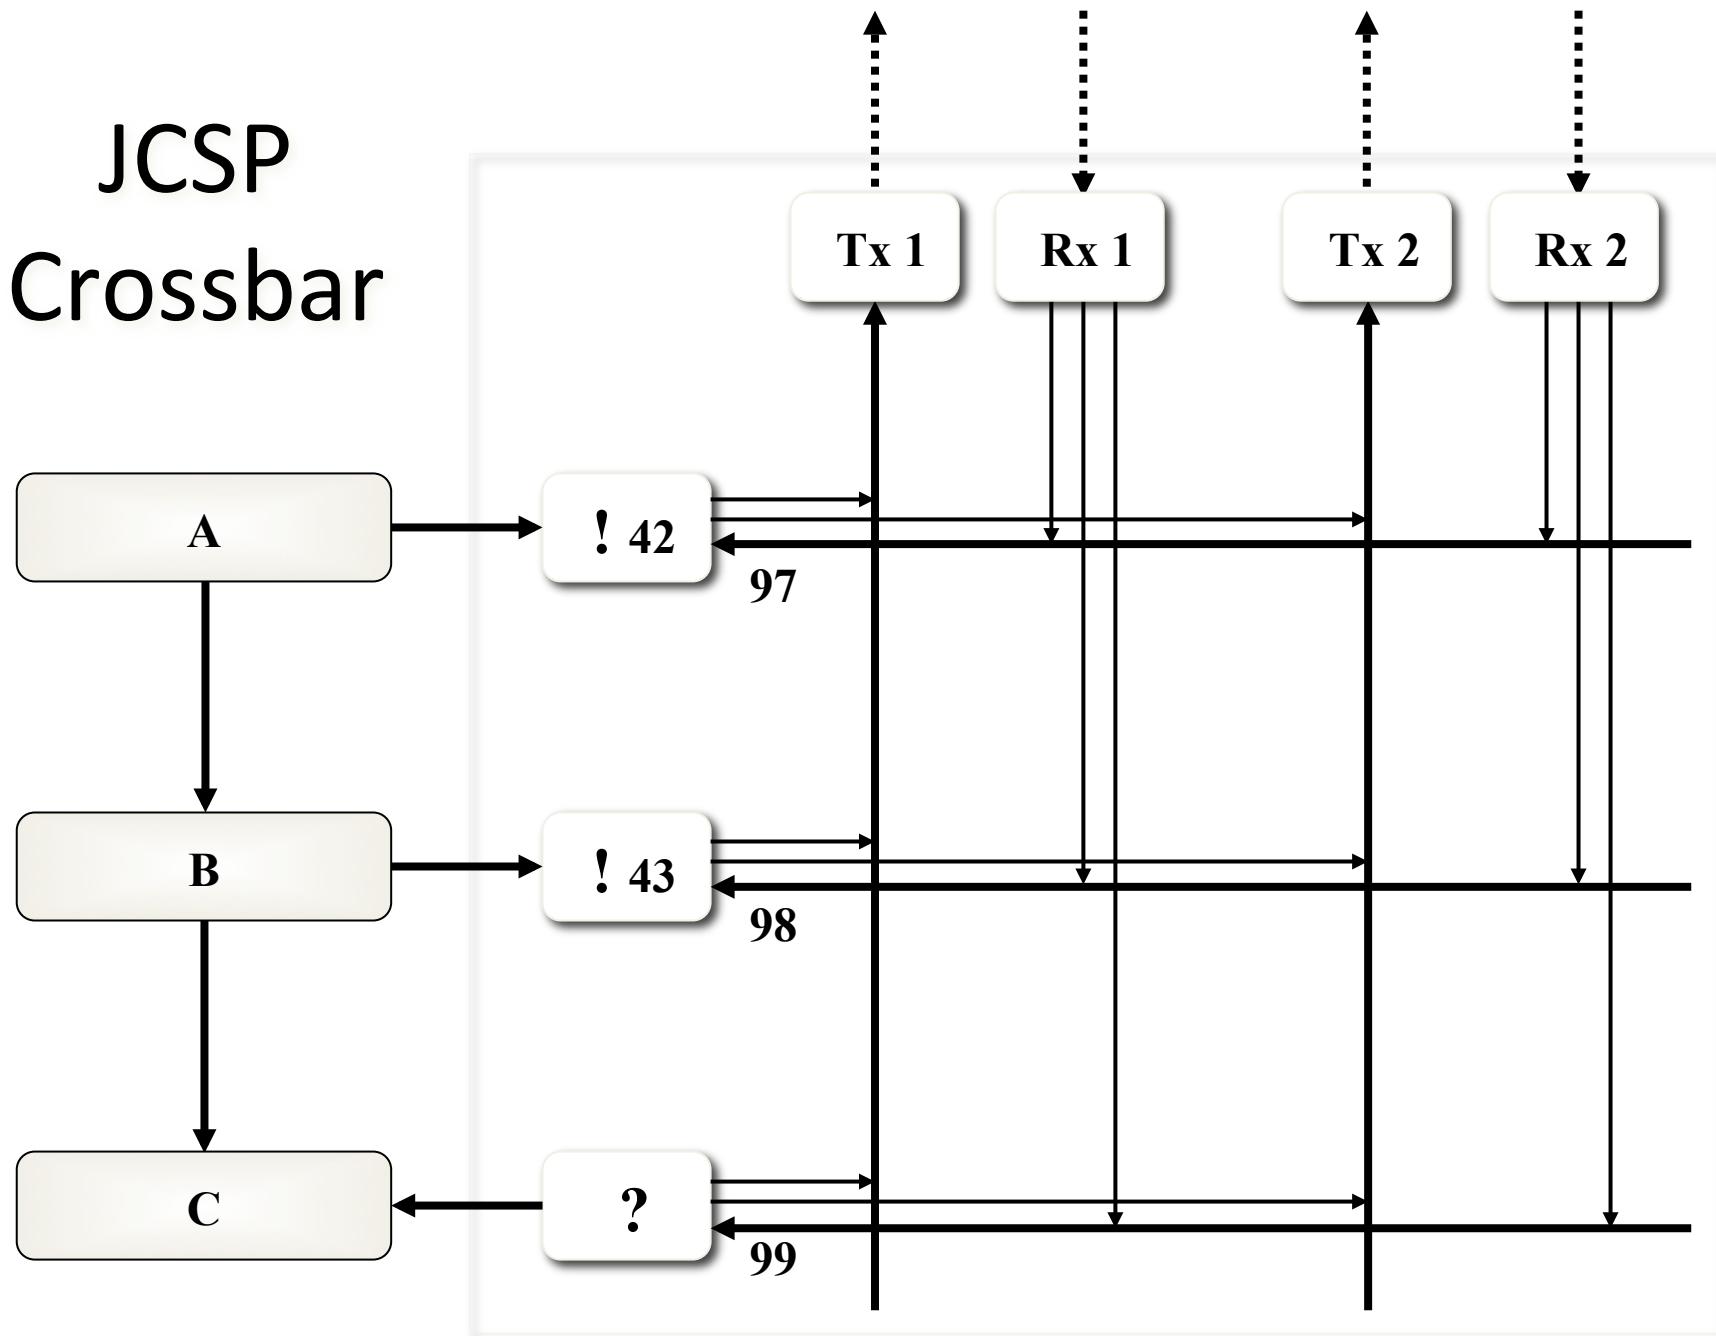


JCSP Links

- Each end of a (e.g. *TCP/IP*) network channel has a network address (e.g. *<IP-address, port-number>*) and **JCSP *virtual-channel-number*** (see below).
- **JCSP** uses the channel-numbers to multiplex and de-multiplex data and acknowledgements.
- The ***JCSP.net*** programmer sees none of this.

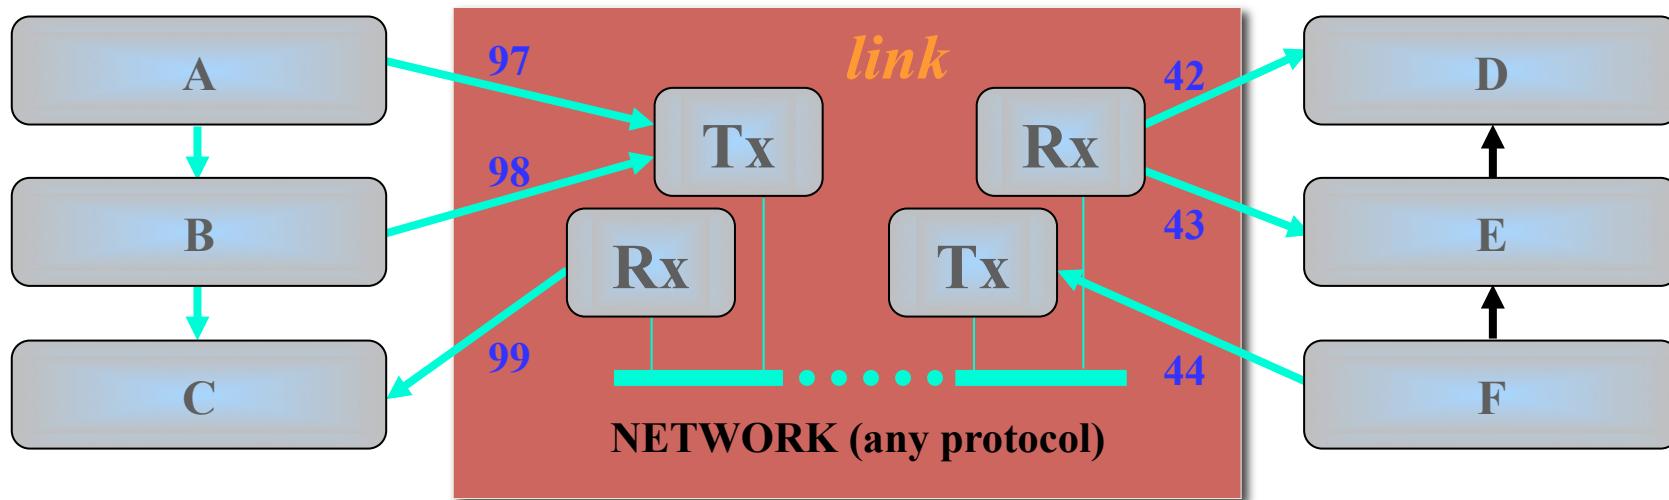


JCSP Crossbar



JCSP Links

- Each end of a (e.g. *TCP/IP*) network channel has a network address (e.g. *<IP-address, port-number>*) and JCSP *virtual-channel-number* (see below).
- JCSP uses the channel-numbers to multiplex and de-multiplex data and acknowledgements.
- **The *JCSP.net* programmer sees none of this.**



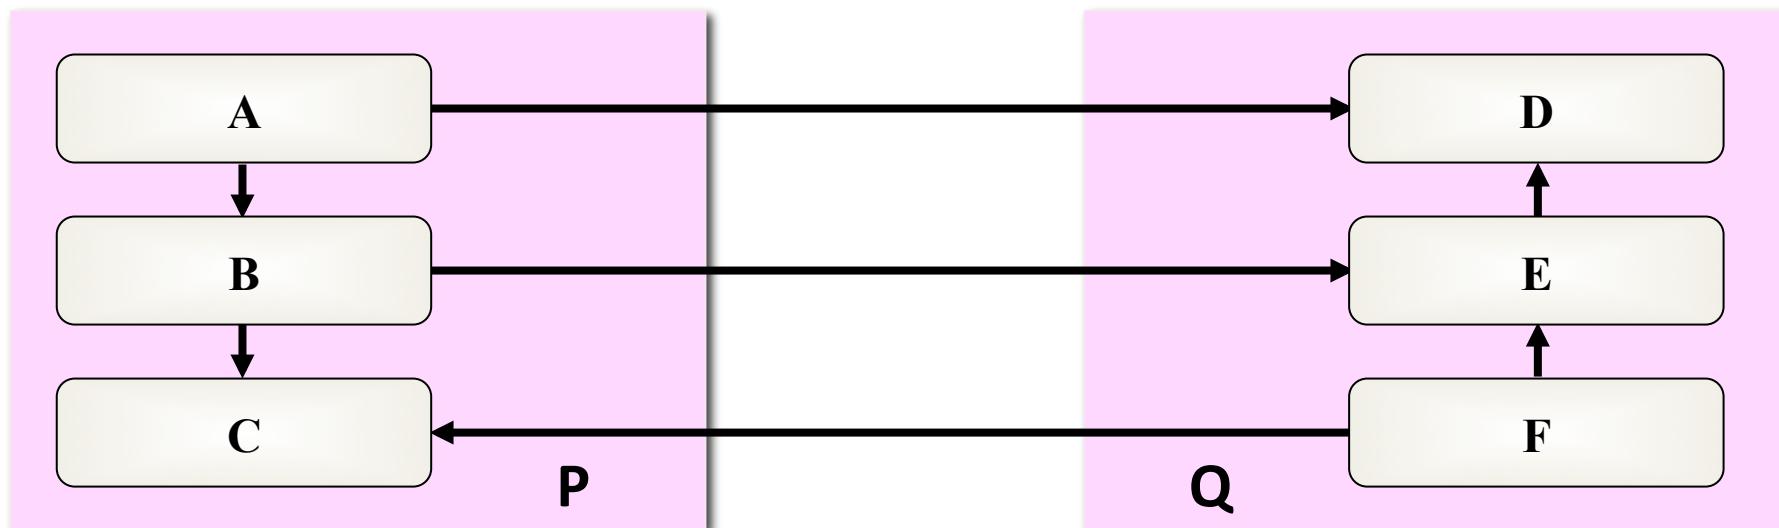
JCSP Networks

- The *JCSP.net* programmer just sees this.
- Channel synchronisation semantics for network channels are *exactly the same* as for internal ones.
- Buffered network channels can be *streamed* - i.e. network acks can be saved through *windowing*.



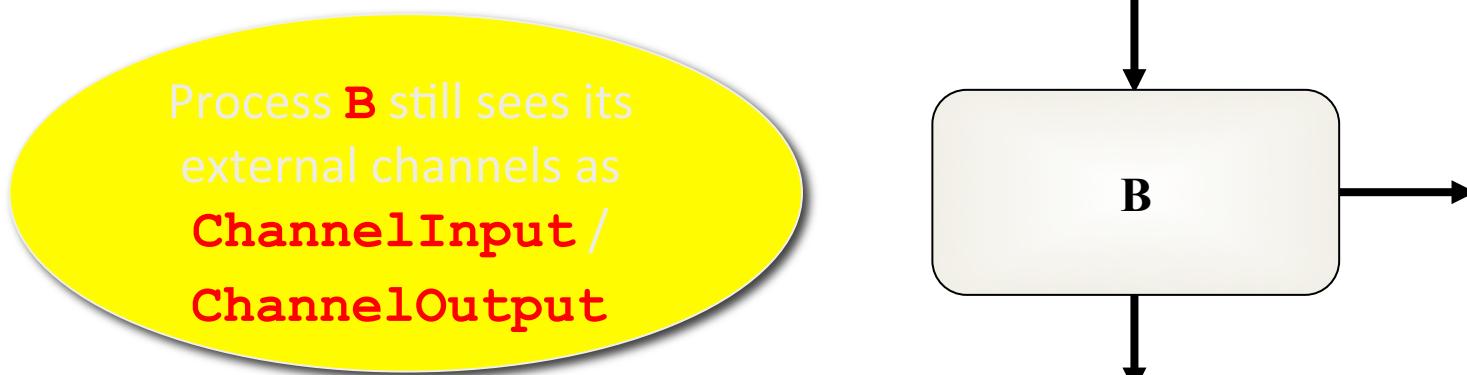
JCSP Networks

- However, there is one important semantic difference between a network channel and a local channel.
- Over *local* channels, objects are passed by *reference* (which leads to *race hazards* if careless).
- Over *network* channels, objects are passed by *copying* (currently, using Java *serialization*).



JCSP Networks

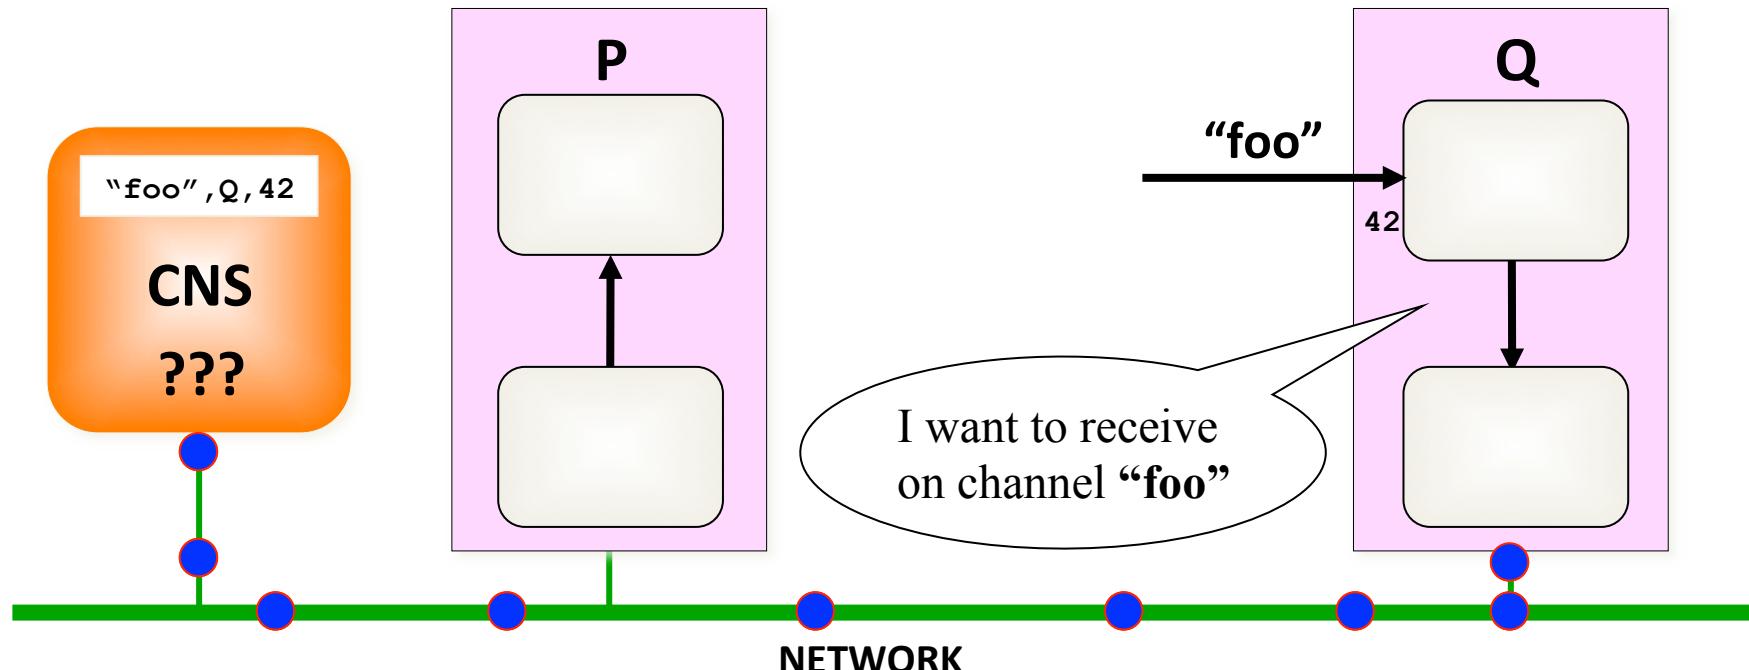
- That semantic difference will not impact correctly designed **JCSP** systems (i.e. those free from *race hazards*).
- With that *caveat*, **JCSP** processes are *blind* as to whether they are connected to local or network channels.



- One other *caveat* - currently, only **Serializable** objects are copied over network channels - sorry!

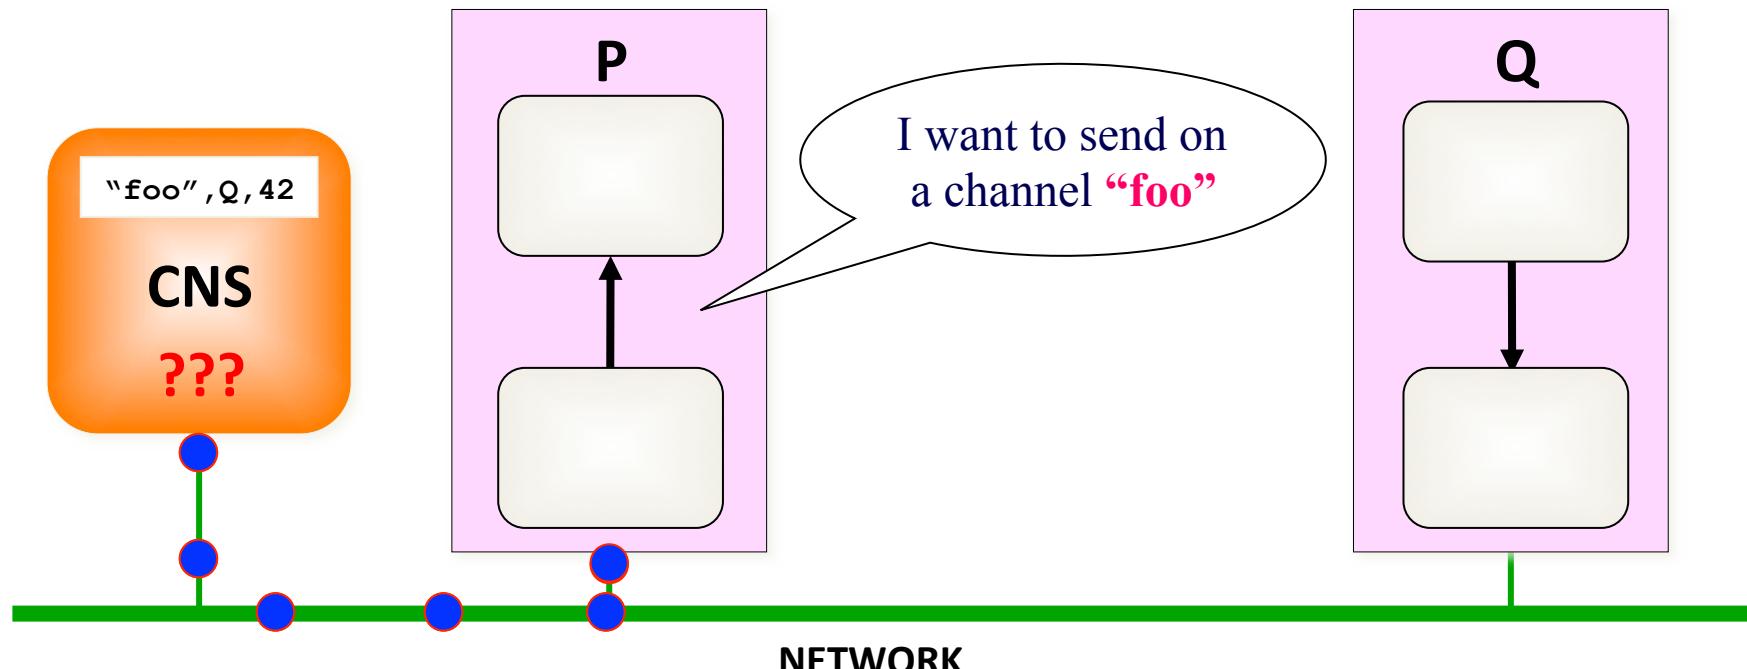
Establishing Network Channels

- Network channels may be connected by the **JCSP** *Channel Name Server* (CNS).
- Channel *read ends* register names with the CNS.



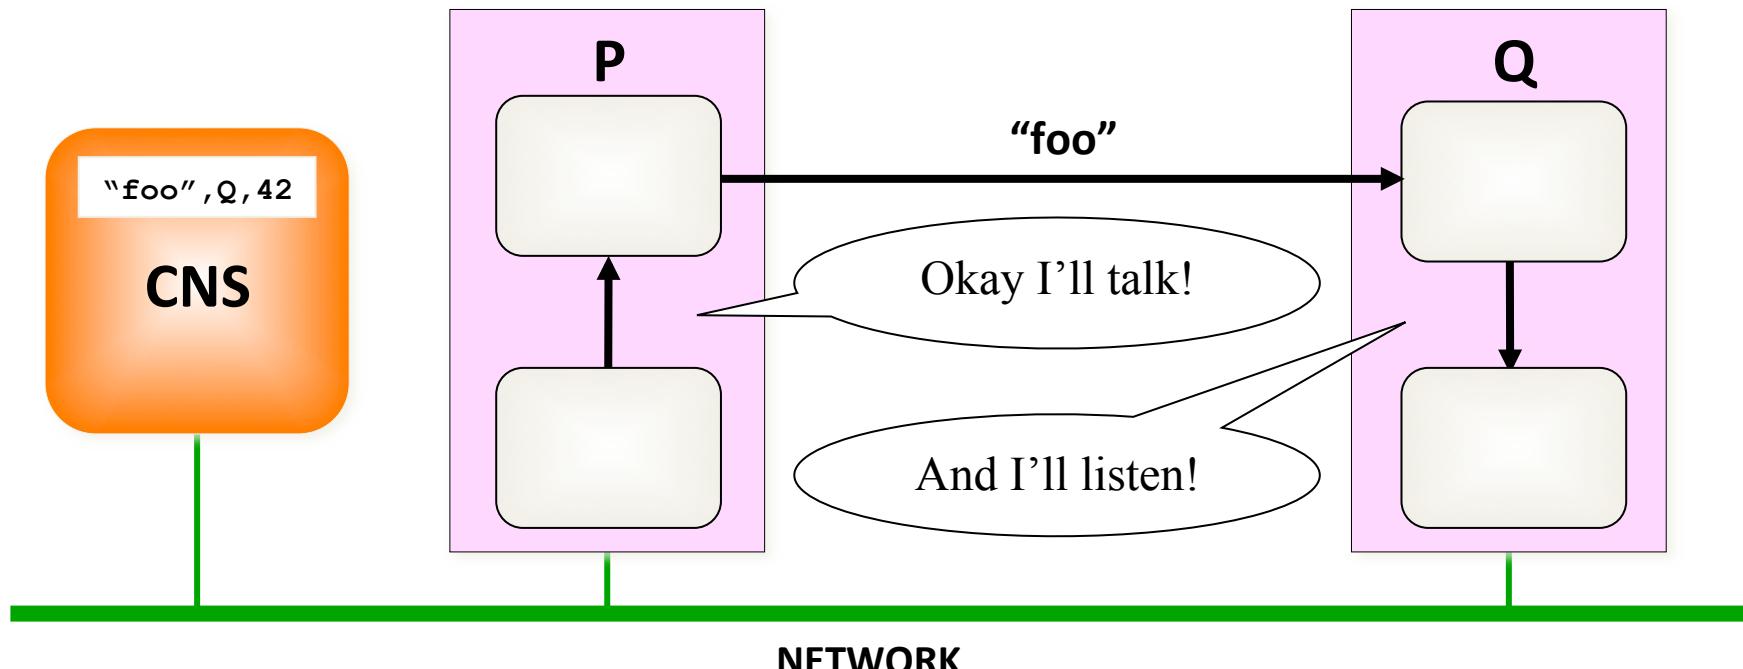
Establishing Network Channels

- Network channels may be connected by the JCSP ***Channel Name Server*** (CNS).
- Channel *read ends* register names with the CNS.
- Channel *write ends* ask CNS about names.



Establishing Network Channels

- Network channels may be connected by the JCSP ***Channel Name Server*** (CNS).
- Channel *read ends* register names with the CNS.
- Channel *write ends* ask CNS about names.



Using Distributed JCSP

- On each machine, do this once:

```
Node.getInstance().init(); // use default CNS
```

- On the **CMU** machine:

```
One2NetChannel out = new One2NetChannel ("ukc.foo");  
new Producer (out);
```



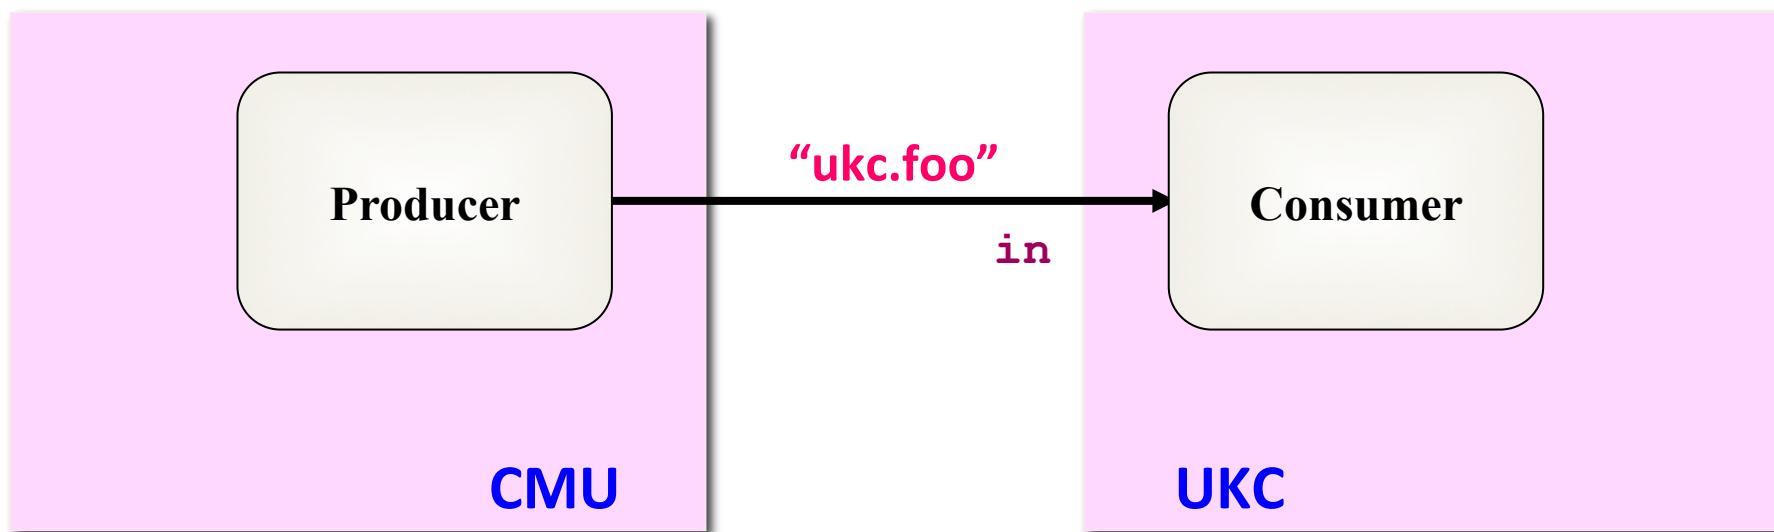
Using Distributed JCSP

- On each machine, do this once:

```
Node.getInstance().init(); // use default CNS
```

- On the **UKC** machine:

```
Net2OneChannel in = new Net2OneChannel ("ukc.foo");  
new Consumer (in);
```



Using Distributed JCSP

```
One2NetChannel out = new One2NetChannel ("ukc.foo");
```

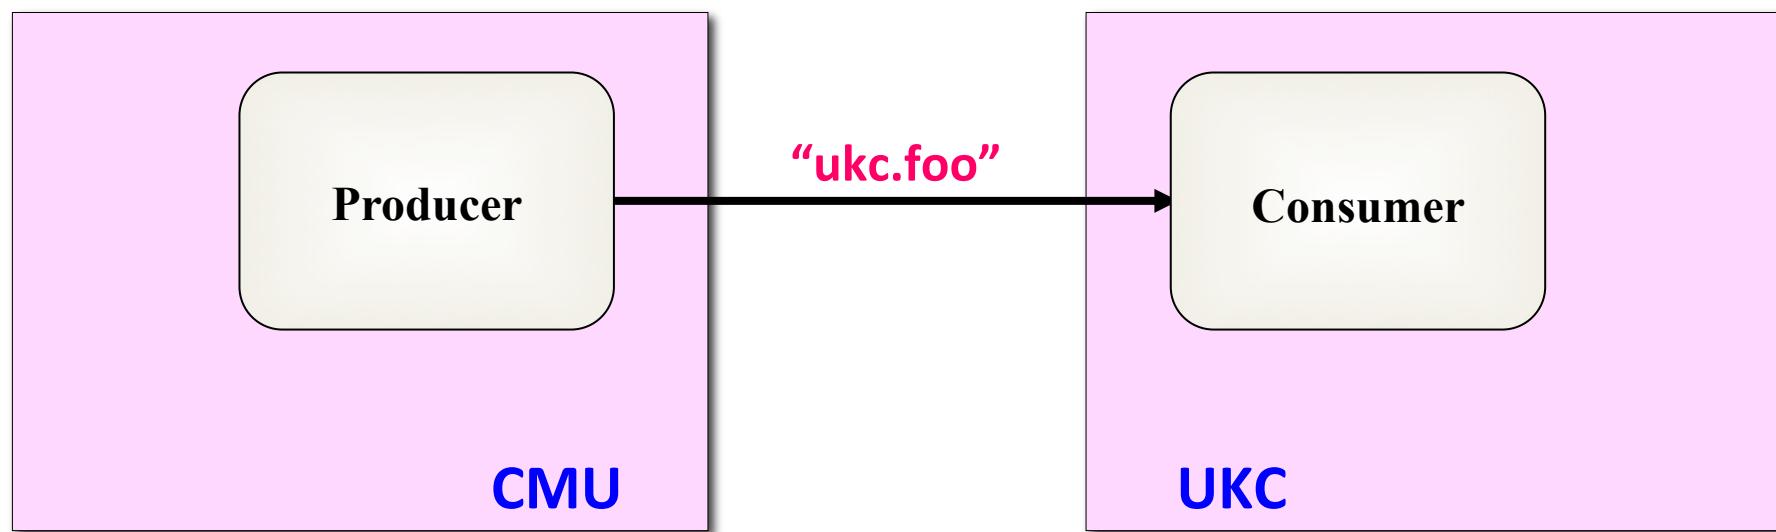
Named network *output* channel construction
blocks until the name is registered by a reader

```
Net2OneChannel in = new Net2OneChannel ("ukc.foo");
```

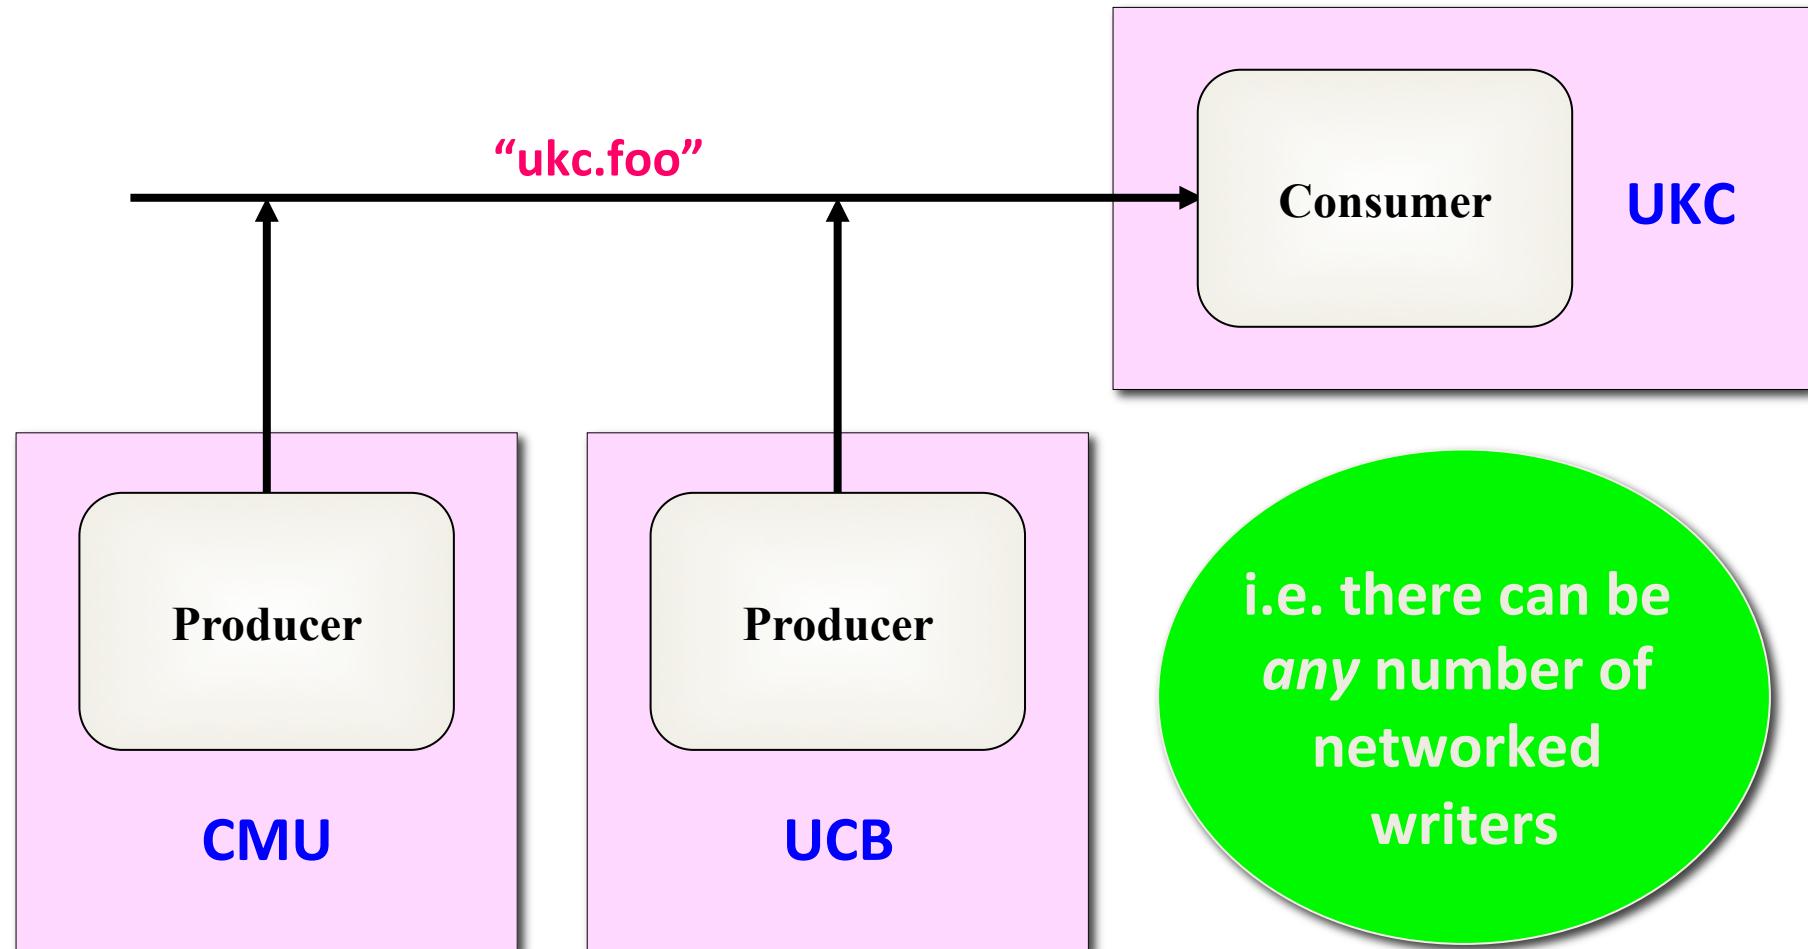
Named network *input* channel construction
registers the name with the CNS (will fail if
already registered)

Using Distributed JCSP

- User processes just have to agree on (or find out) *names* for the channels they will use to communicate.
- User processes do not have to know where each other is located (e.g. *IP-address / port-number / virtual-channel-number*).

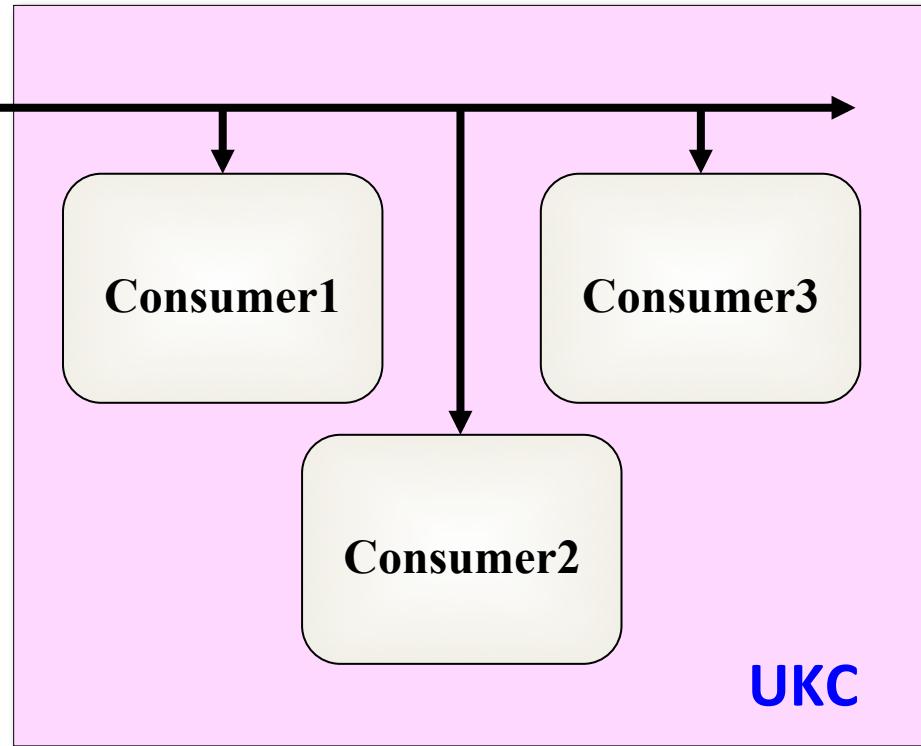


Network Channels are *Any-1*



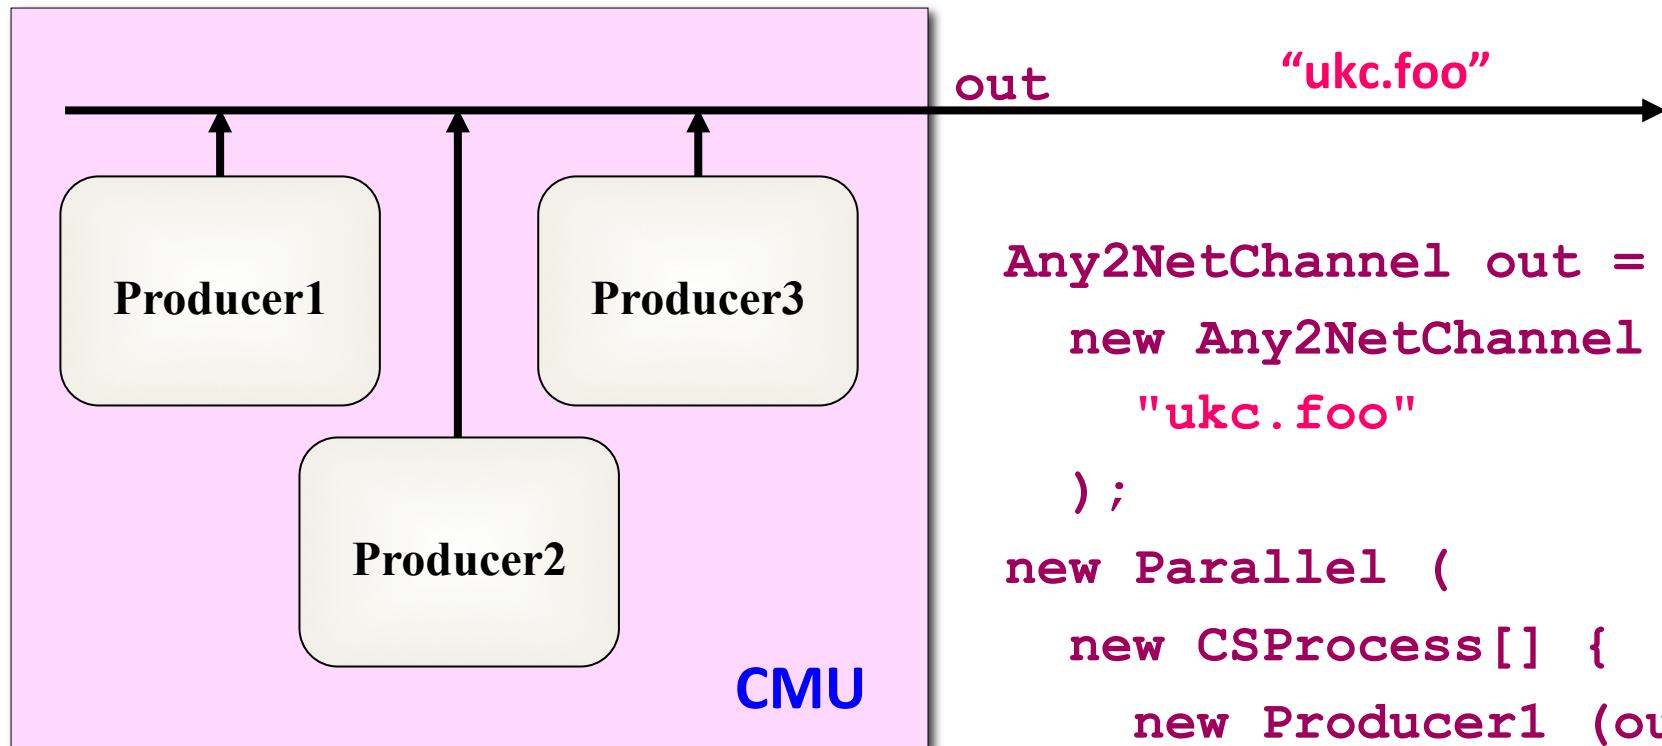
Net-Any Channels

```
“ukc.foo”  
-----  
in  
  
Net2AnyChannel in =  
    new Net2AnyChannel (  
        "ukc.foo"  
    );  
  
new Parallel (  
    new CSProcess[] {  
        new Consumer1 (in) ,  
        new Consumer2 (in) ,  
        new Consumer2 (in)  
    }  
).run ();
```



i.e. within a node, there
can be *any* number of
readers

Any-Net Channels



i.e. within a node, there
can be *any* number of
writers

```
Any2NetChannel out =  
    new Any2NetChannel (  
        "ukc.foo"  
    );  
    new Parallel (  
        new CSProcess[] {  
            new Producer1 (out),  
            new Producer2 (out),  
            new Producer3 (out)  
        }  
    ).run ();
```

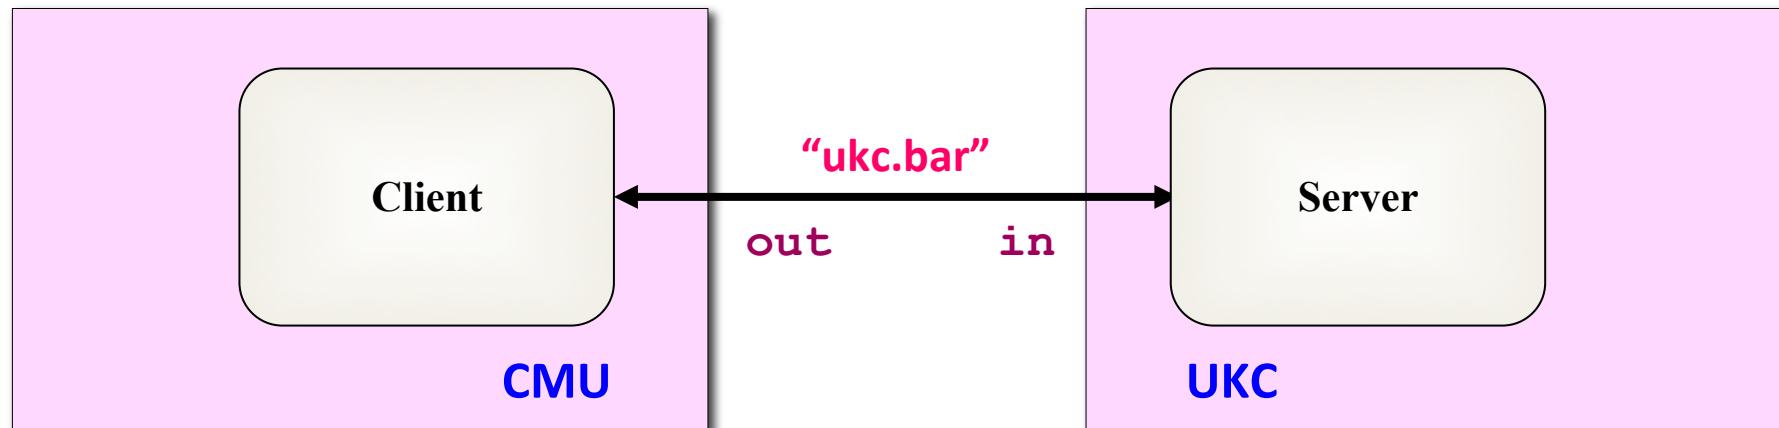
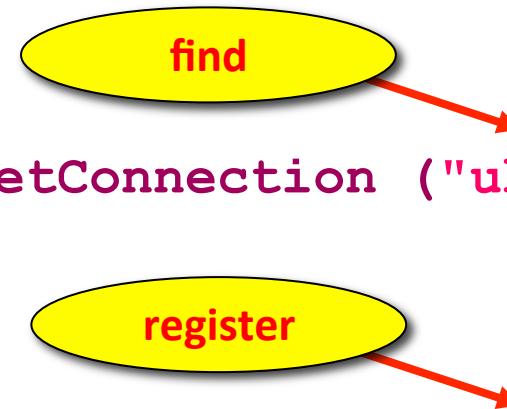
Connections (*two-way channels*)

- On the **CMU** machine:

```
One2NetConnection out = new One2NetConnection ("ukc.bar");  
new Client (out);
```

- On the **UKC** machine:

```
Net2OneConnection in = new Net2OneConnection ("ukc.bar");  
new Server (in);
```



Anonymous Channels

- Network channels *may be* connected by the **JCSP Channel Name Server (CNS)** ...
- *... but they don't have to be!*



- A network channel can be created (*always by the inputter*) without registering a name with the CNS:

```
Net2OneChannel in = new Net2OneChannel (); // no name!
```

- Remote processes cannot, of course, find it for themselves ... *but you can tell your friends ...*

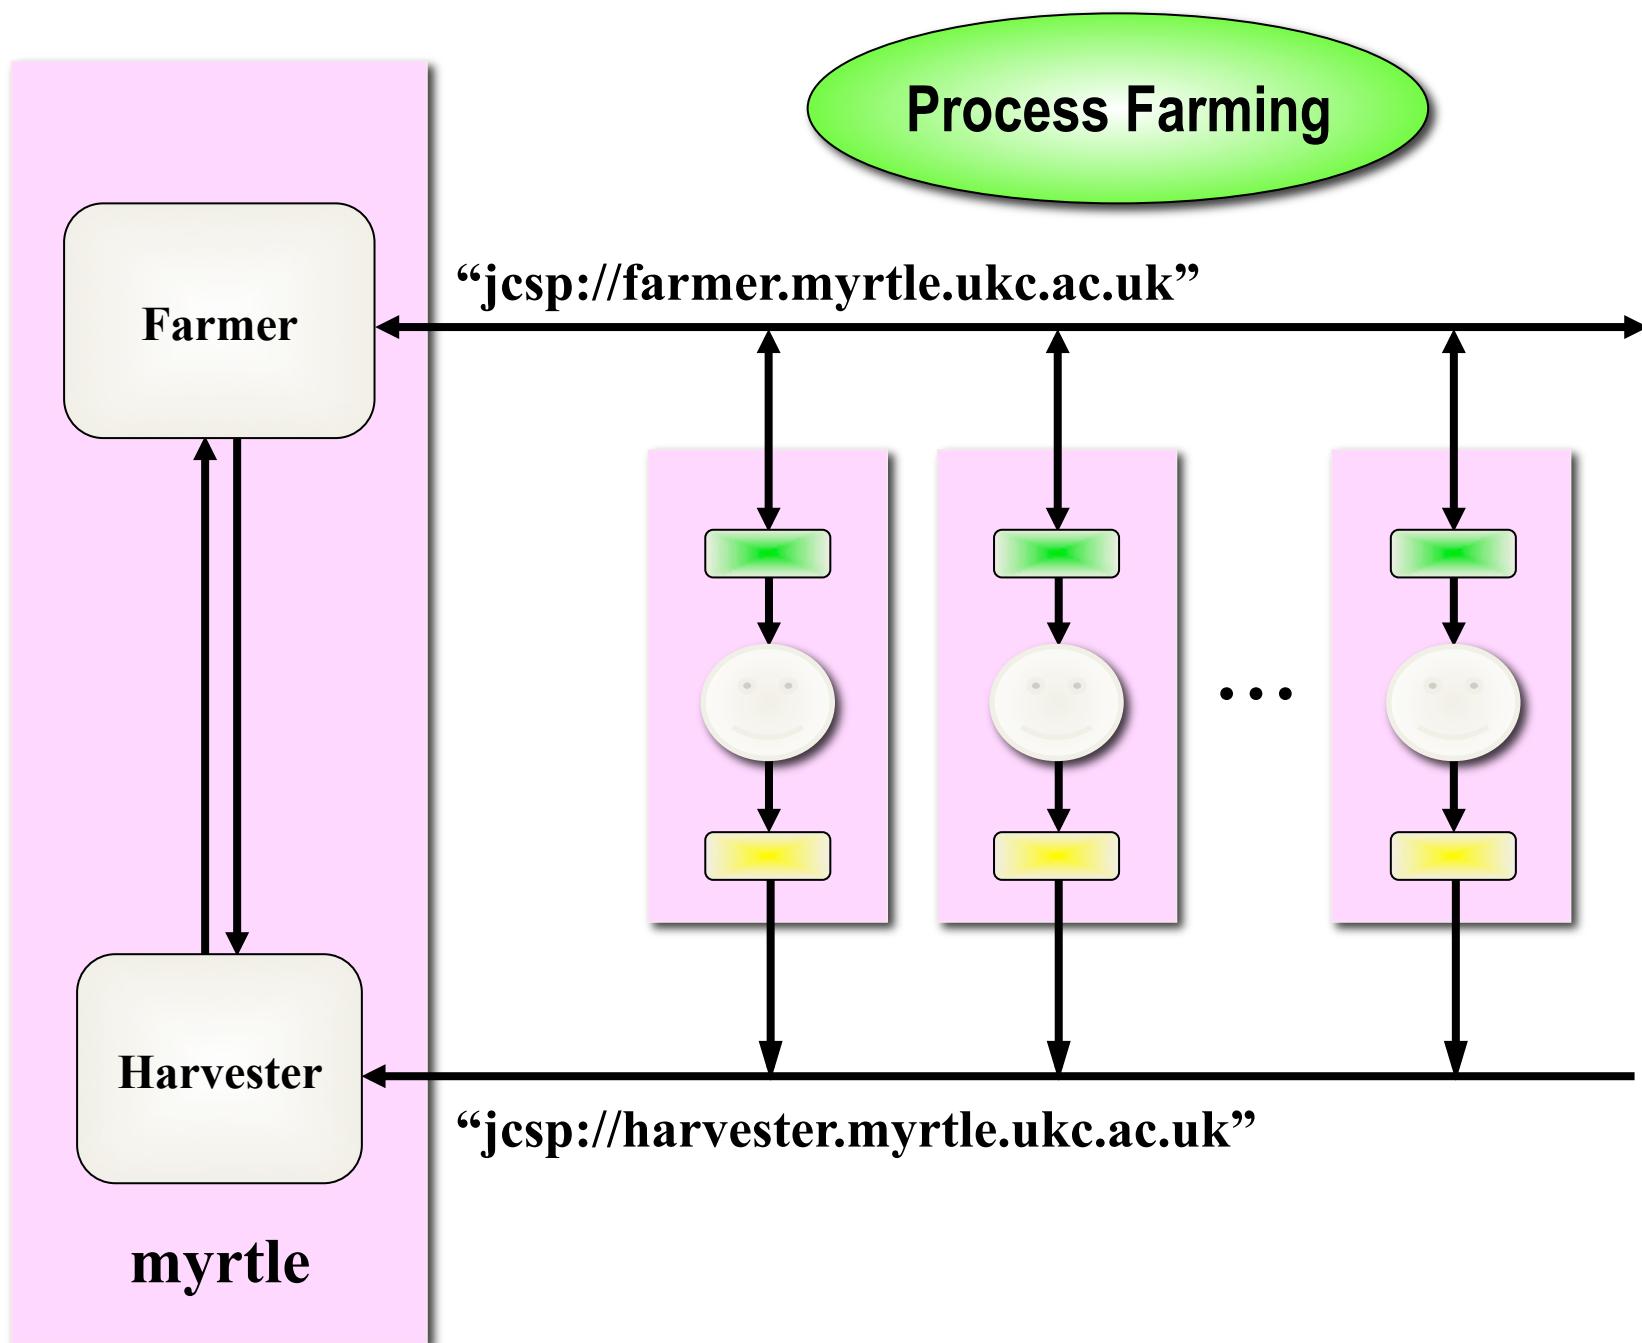
Anonymous Channels

- Location information (*<IP-address, port-number, virtual-channel-number>*) is held within the constructed network channel. This is the data registered with the CNS - if we had given it a name.
- Extract that information:

```
Net2OneChannel in = new Net2OneChannel () ; // no name !
NetChannelLocation inLocation = in.getLocation () ;
```

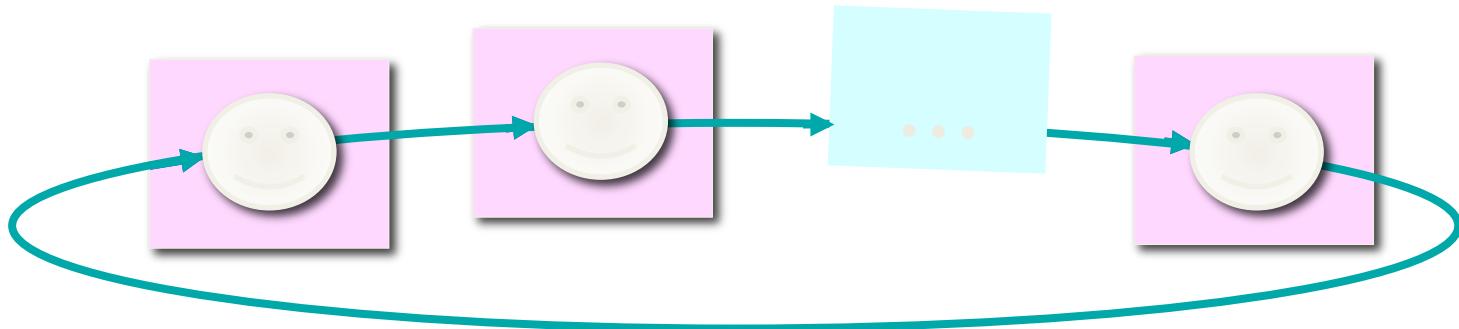
- The information can be distributed using existing (network) channels to those you trust:

```
toMyFriend.write (inLocation) ;
// remember your friend may distribute it further ...
```



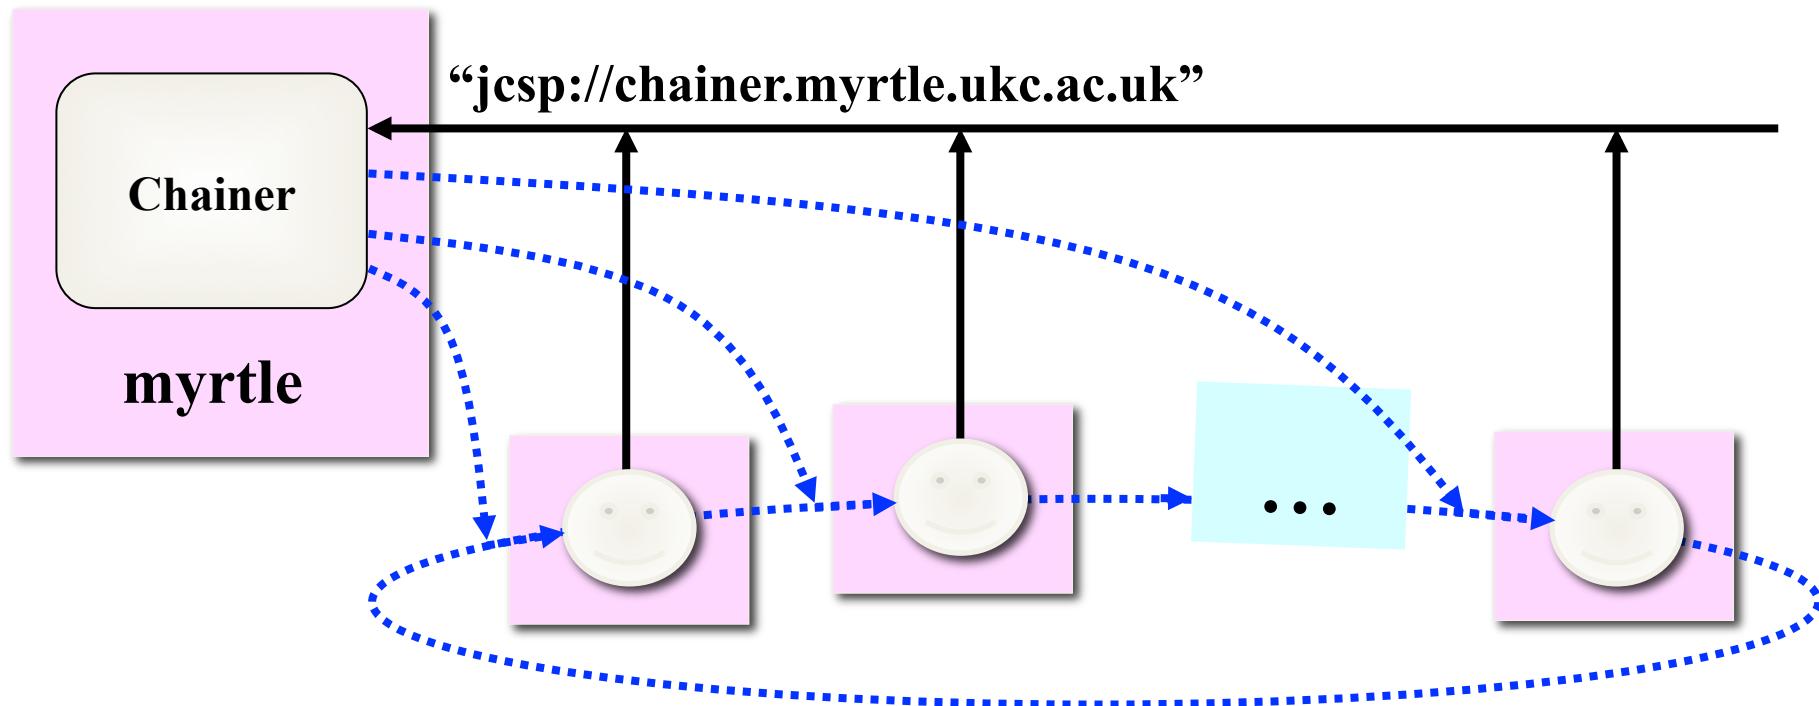
Process Chaining

```
final int MY_ID = ... ;  
final int N_NODES = ... ;  
final int NEXT_ID = ((MY_ID + 1) % N_NODES);  
  
Net2OneChannel in = new Net2OneChannel ("node-" + MY_ID);  
Net2OneChannel out = new One2NetChannel ("node-" + NEXT_ID);  
  
new WorkProcess (MY_ID, N_NODES, in, out).run ();
```



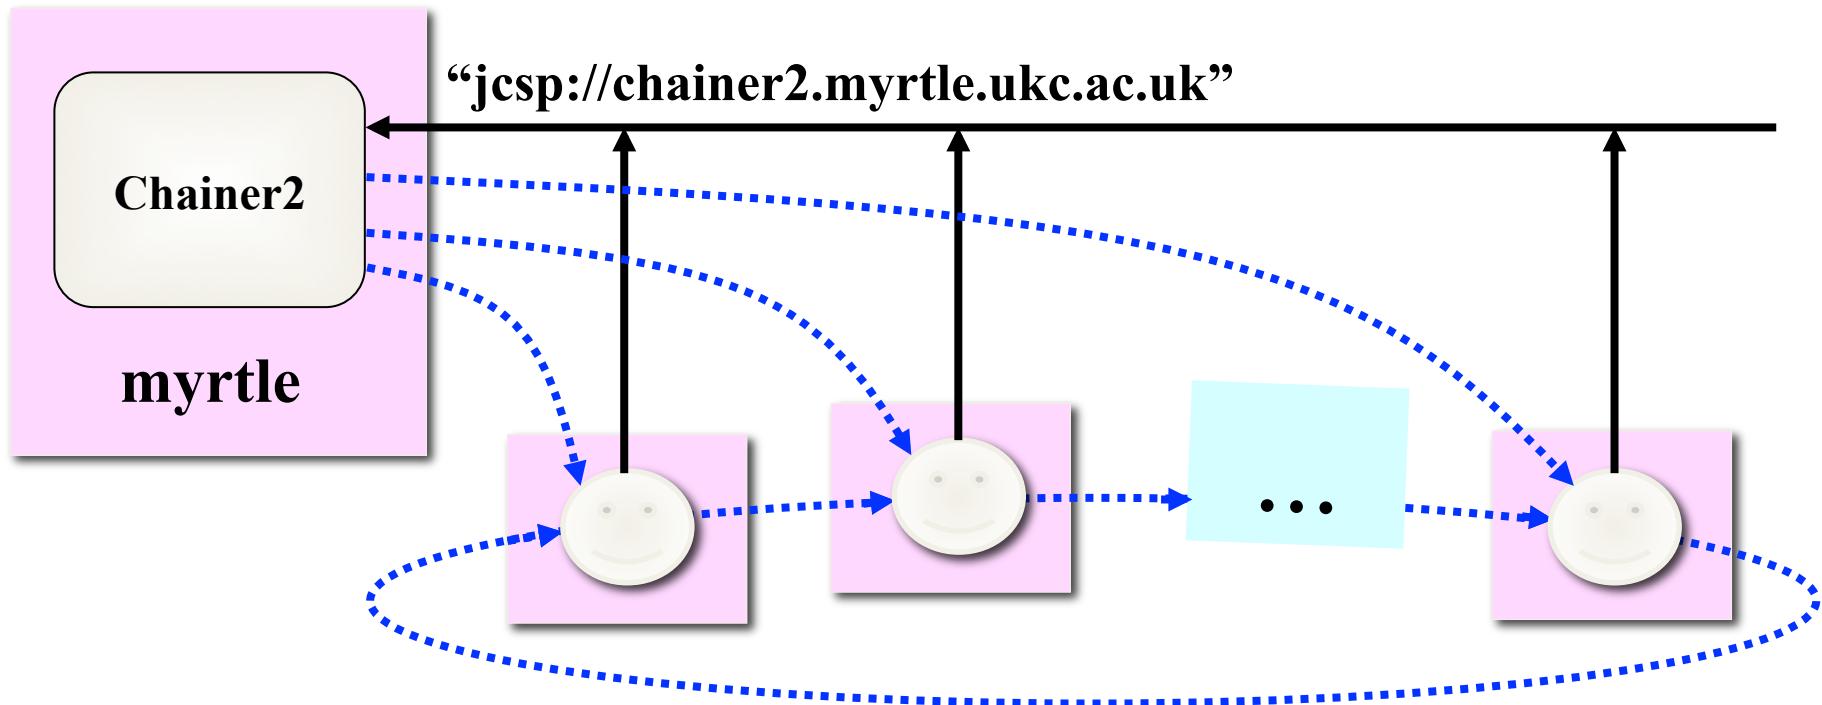
OK – so long as each worker knows the length of the chain and its place in it.

Process Chaining



A volunteer worker won't know this! But it can make its own network input channel *anonymously* and send its location to someone who does ...

Process Chaining



It's slightly easier if each node makes *two* network input channels – so that its *control* line is different from its *data* line from the chain ...

Ring worker code

```
One2NetChannel toChainer =  
  
    = new One2NetChannel ("jcsp://chainer.myrtle.ukc.ac.uk");  
  
Net2OneChannel in = new Net2OneChannel ();  
NetChannelLocation inLocation = in.getLocation ();  
toChainer.write (inLocation);  
  
NetChannelLocation outLocation = (NetChannelLocation) in.read ();  
One2NetChannel out = new One2NetChannel (outLocation);  
  
int[] info = (int[]) in.read ();                      // wait for ring sync  
final int MY_ID = info[0];                            // (optional)  
final int N_NODES = info[1];                          // (optional)  
  
info[0]++;  
if (info[0] < info[1]) out.write (info);      // pass on ring sync  
  
new WorkProcess (MY_ID, N_NODES, in, out).run ();
```

```
final int N_NODES = ... ;
```

Chainer (ringer) code

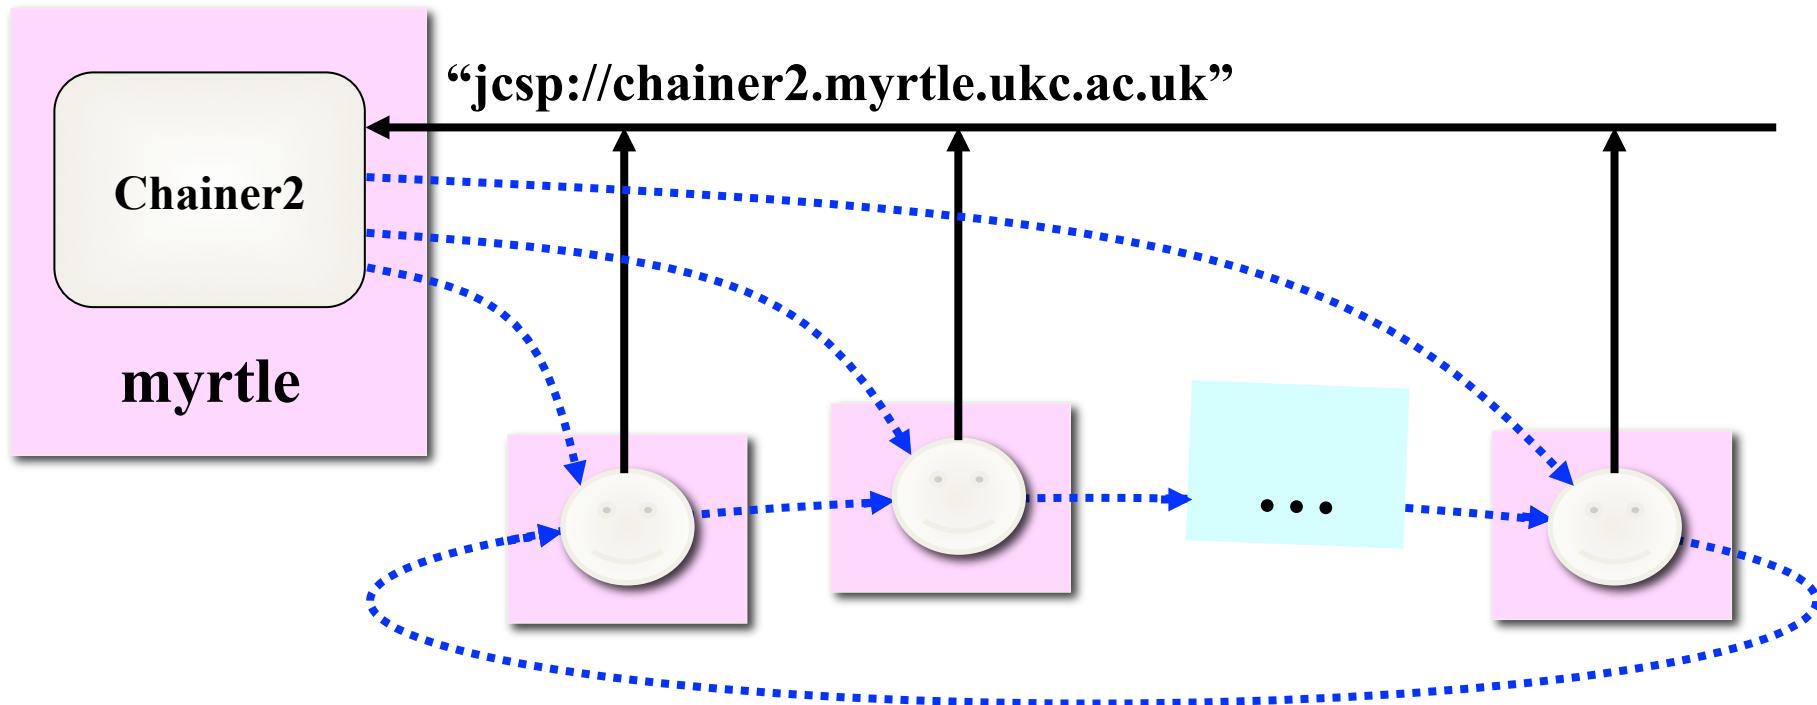
```
Net2OneChannel fromWorkers =
    = new Net2OneChannel ("jcsp://chainer.myrtle.ukc.ac.uk") ;

NetChannelLocation lastL =
    (NetChannelLocation) fromWorkers (read) ;
One2NetChannel lastC = new One2NetChannel (lastL) ;

for (int nWorkers = 1; nWorkers < N_NODES; nWorkers++) {
    NetChannelLocation nextL =
        (NetChannelLocation) fromWorkers (read) ;
    One2NetChannel nextC = new One2NetChannel (nextL) ;
    nextC.write (lastL) ;
    lastL = nextL;
}

lastC.write (lastL);                                // completes the network ring
lastC.write (new int[] {0, N_NODES});   // final ring synchronisation
```

Process Chaining



It's slightly easier if each node makes *two* network input channels – so that its *control* line is different from its *data* line from the chain ...

PyCSP

- Master
 - `sys.path.insert(0, "..")`
 - `Configuration().set(PYCSP_PORT, <portnr>)`
- Others
 - `addr = (<host>, <portnr>)`
- `@Multiprocess(<host>,<portnr>)`

Example Applications

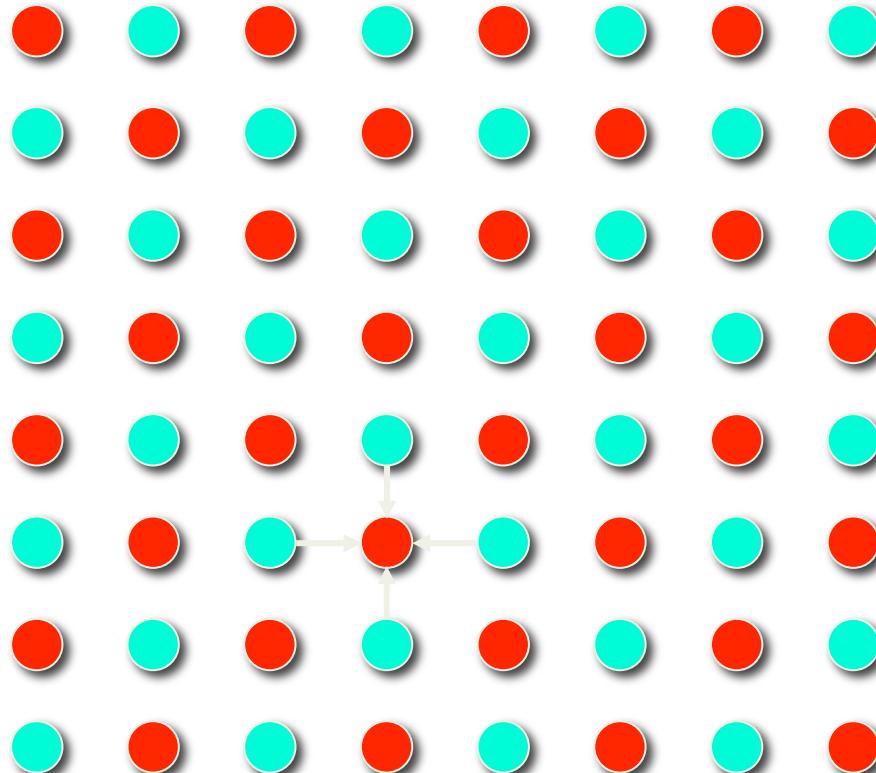
Process Farming

All ‘*embarrassingly parallel*’ ones, ray tracing, Mandelbrot, travelling salesman (needs dynamic control though), ...

Process Chaining

All space-division system modelling, n-body simulations, SORs, cellular automata, ... (some need *bi-directional* chains/rings)

SOR – red/black checker pointing

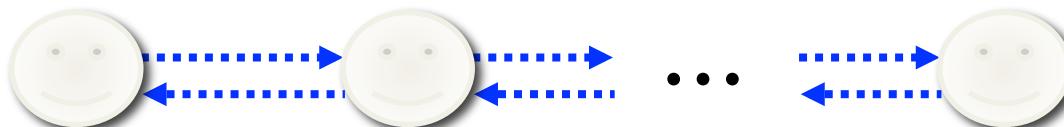


where = black

and = red

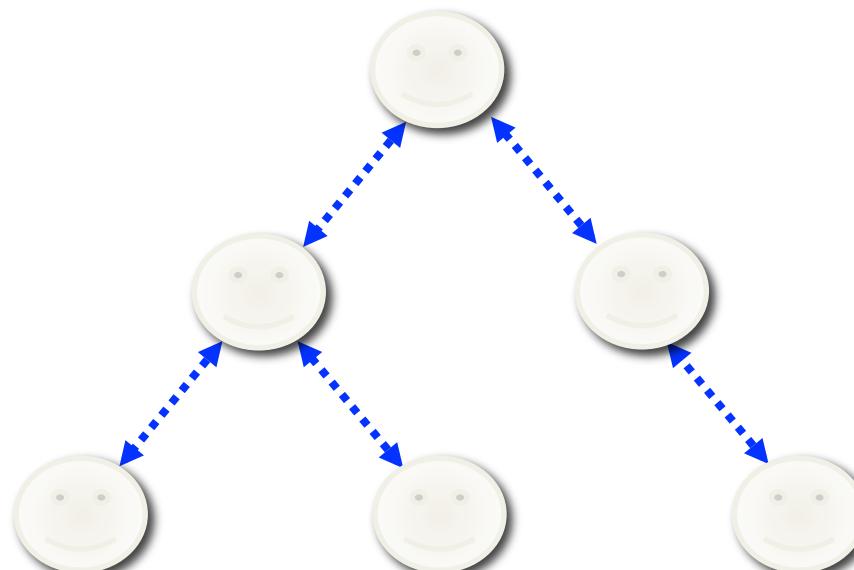
SOR – red/black checker pointing

This needs a *two-way* chain to exchange information on boundary regions being looked after by each worker ...



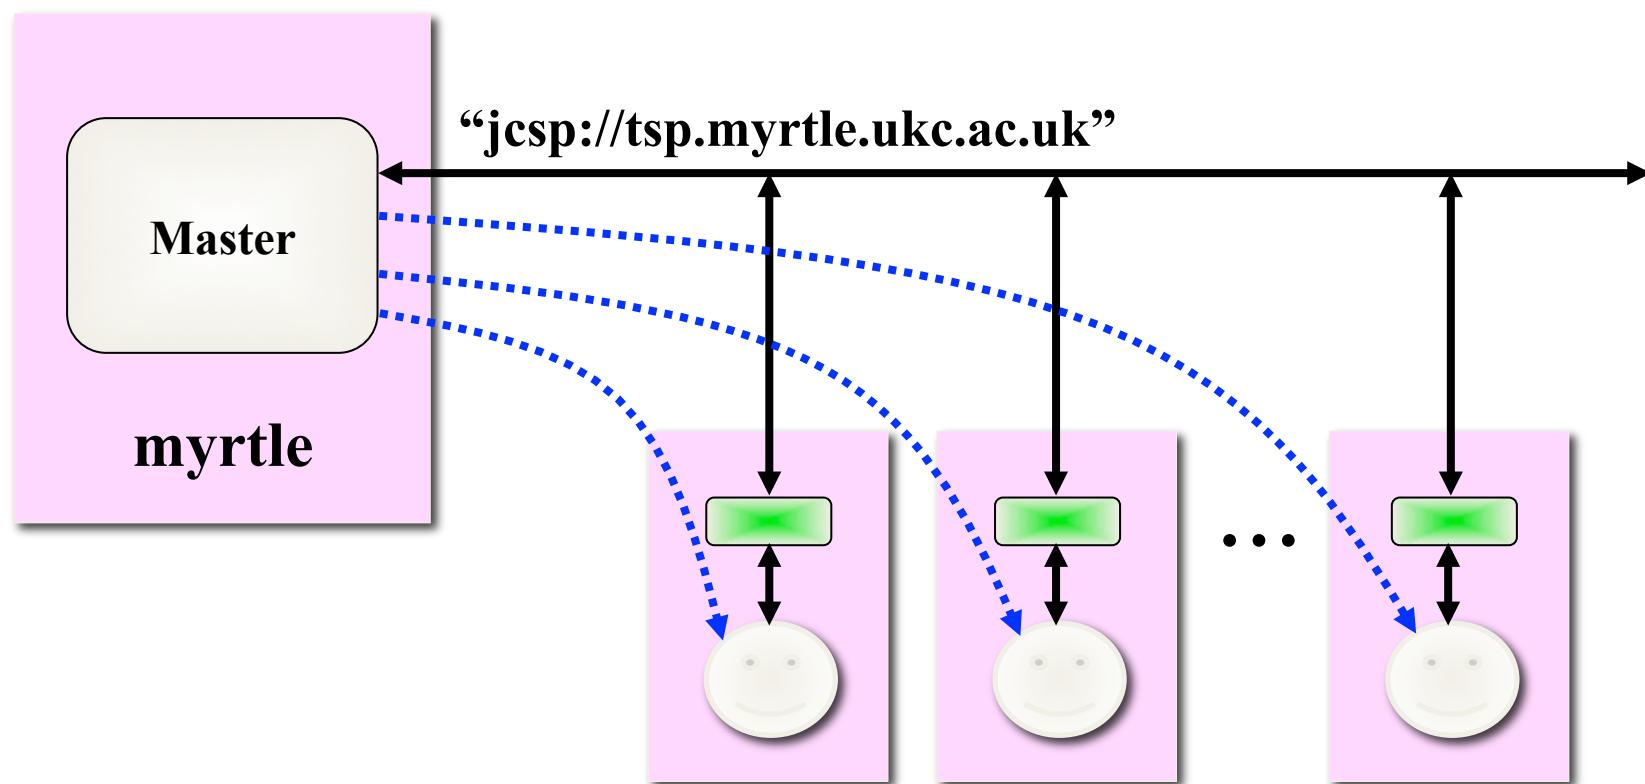
SOR – red/black checker pointing

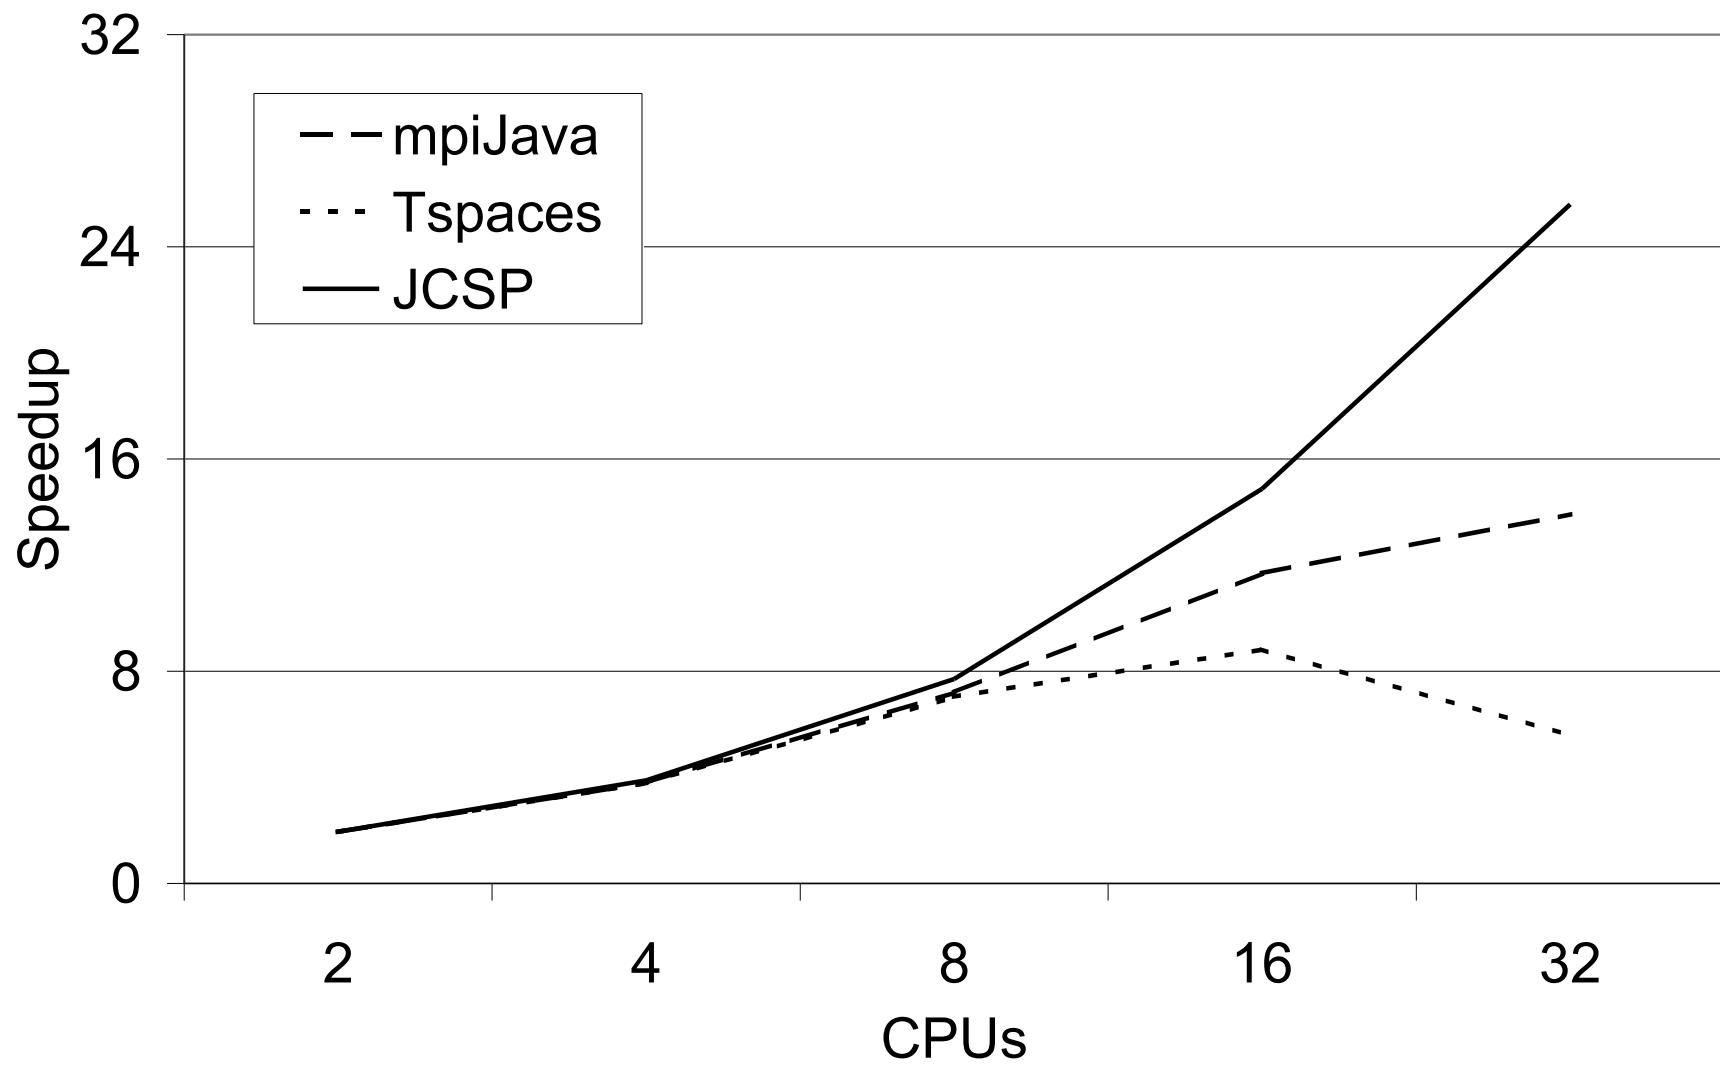
Also, a global *sum-of-changes* (found by each node) has to be computed each cycle to resolve halting criteria. This is speeded-up by connecting the nodes into a tree (so that *adds* and *communications* can take place in parallel).



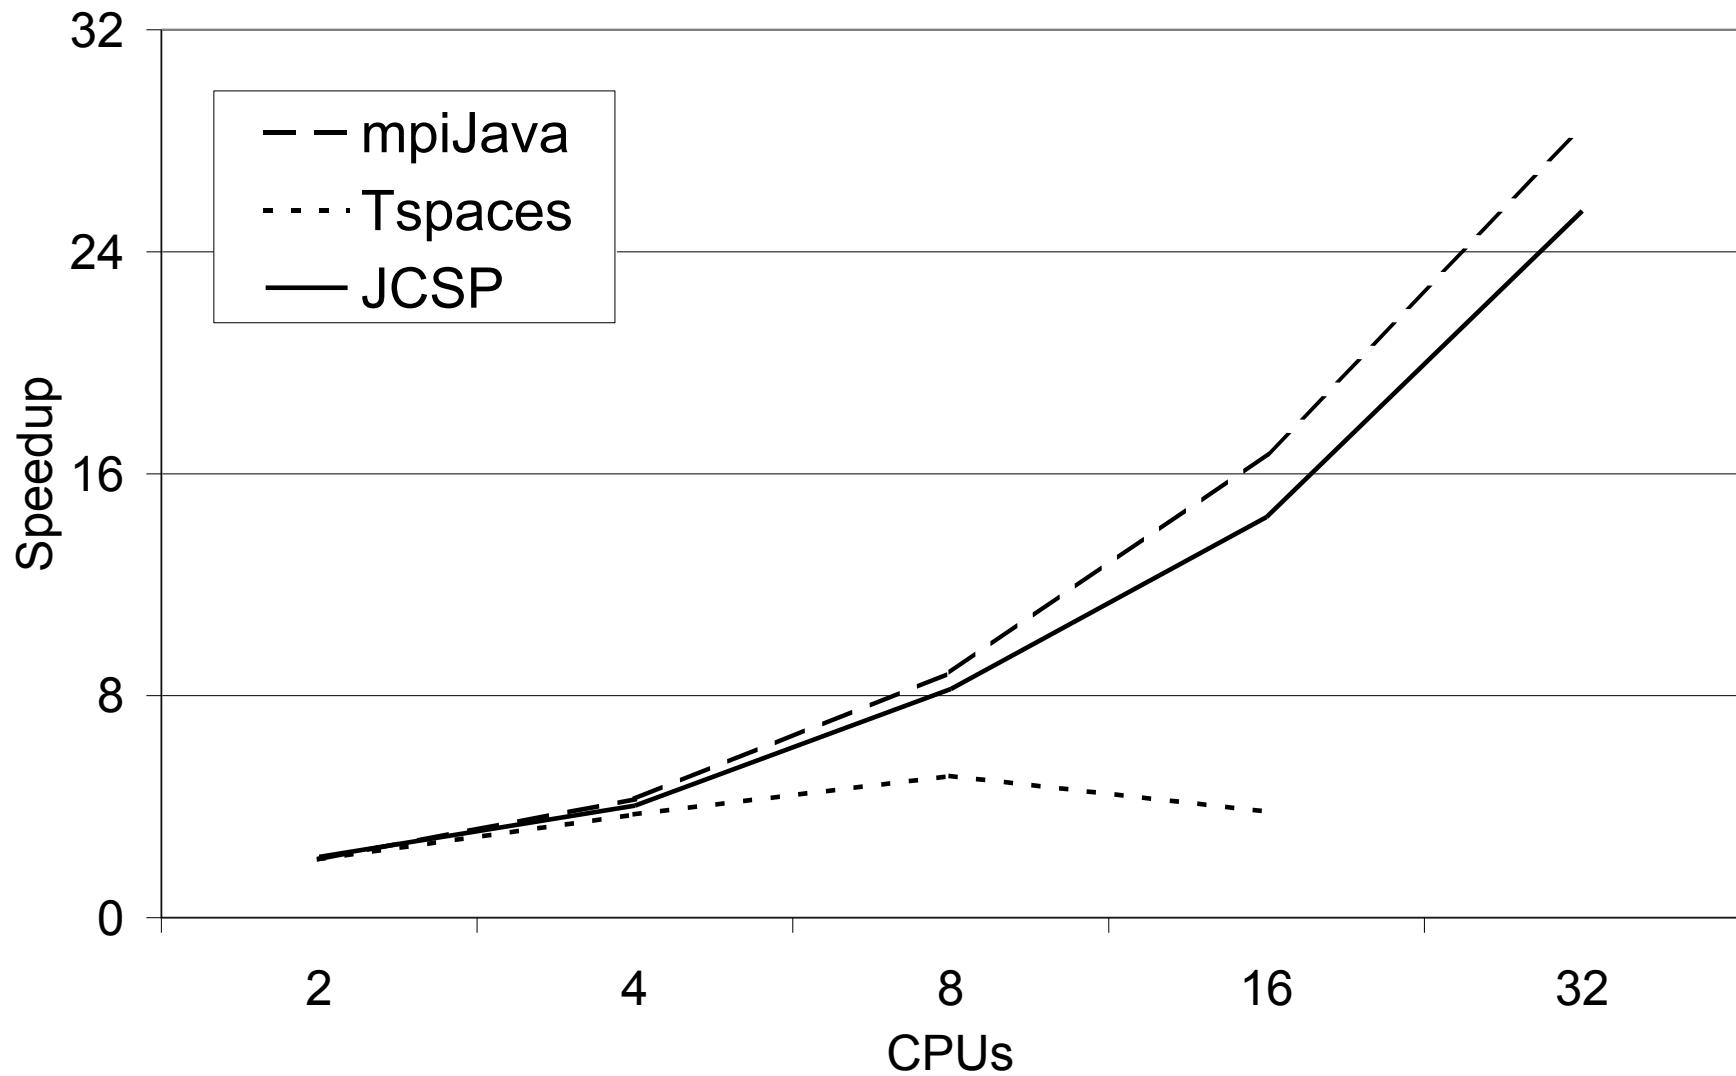
Travelling Salesman Problem

Basically a process farm ... but when better lower bounds arrive, they must be communicated to all concerned workers.

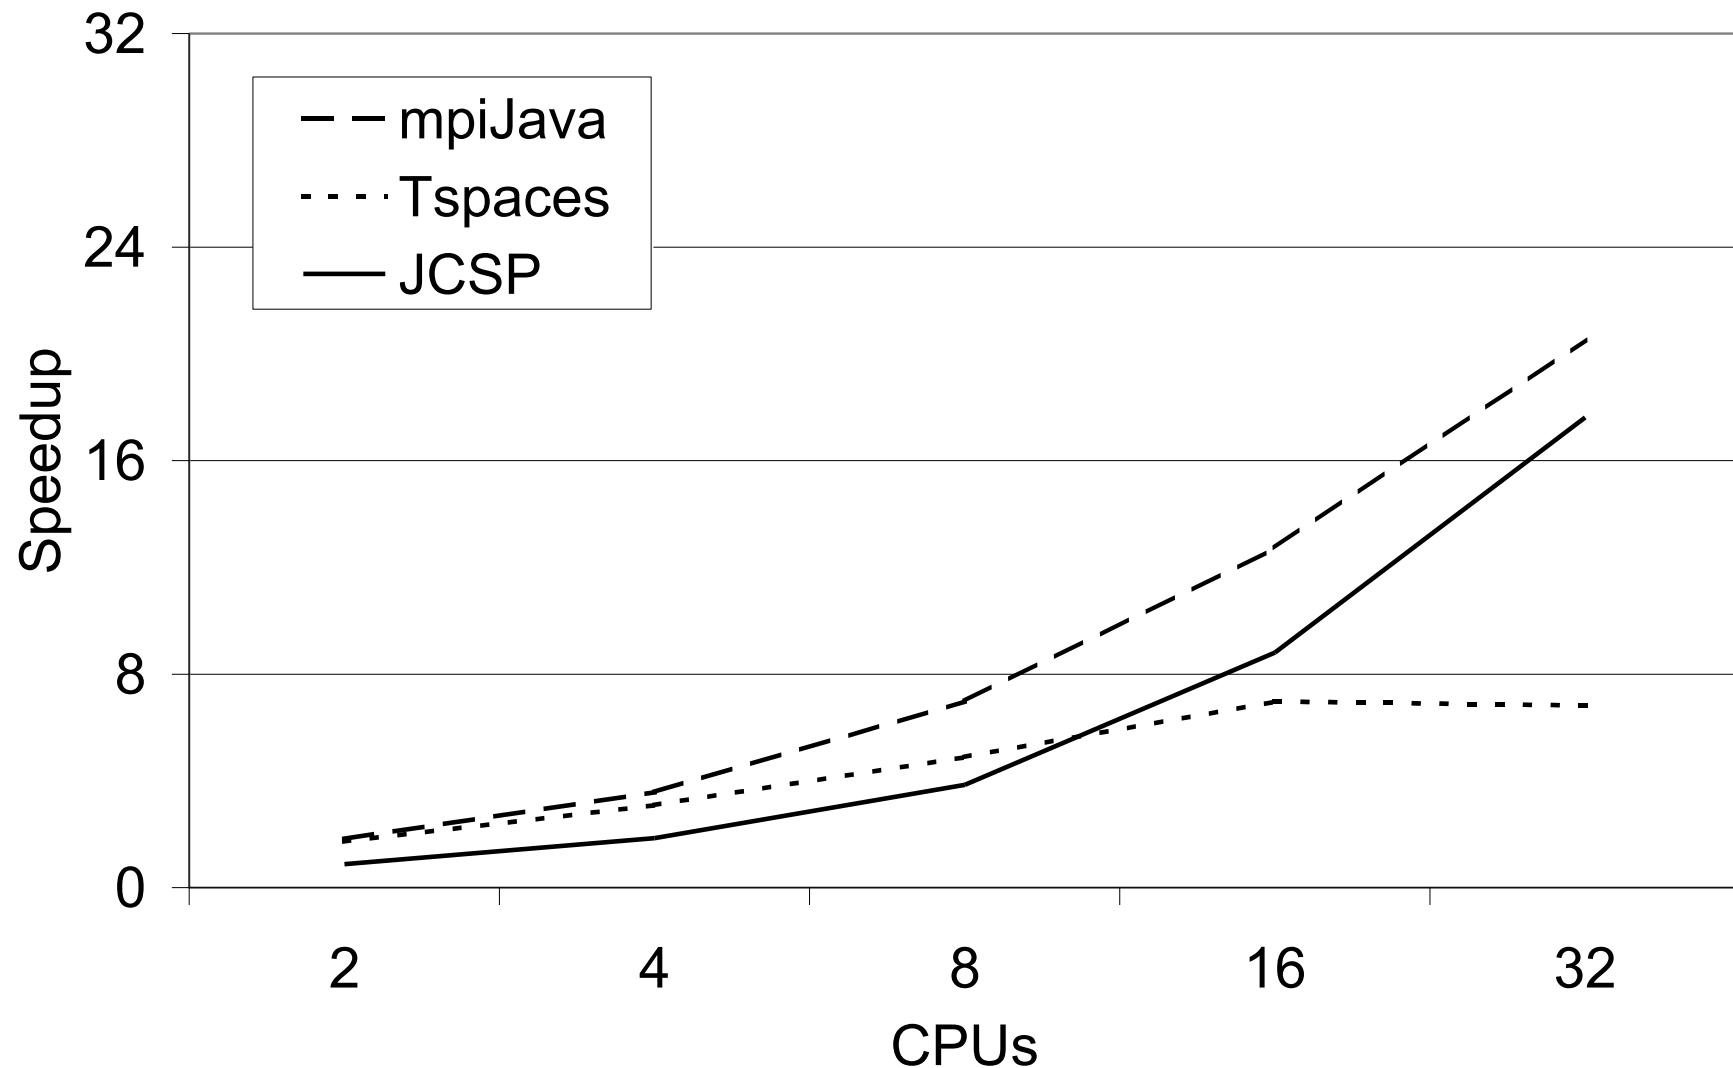




The n-Body benchmark ($n = 10000$)

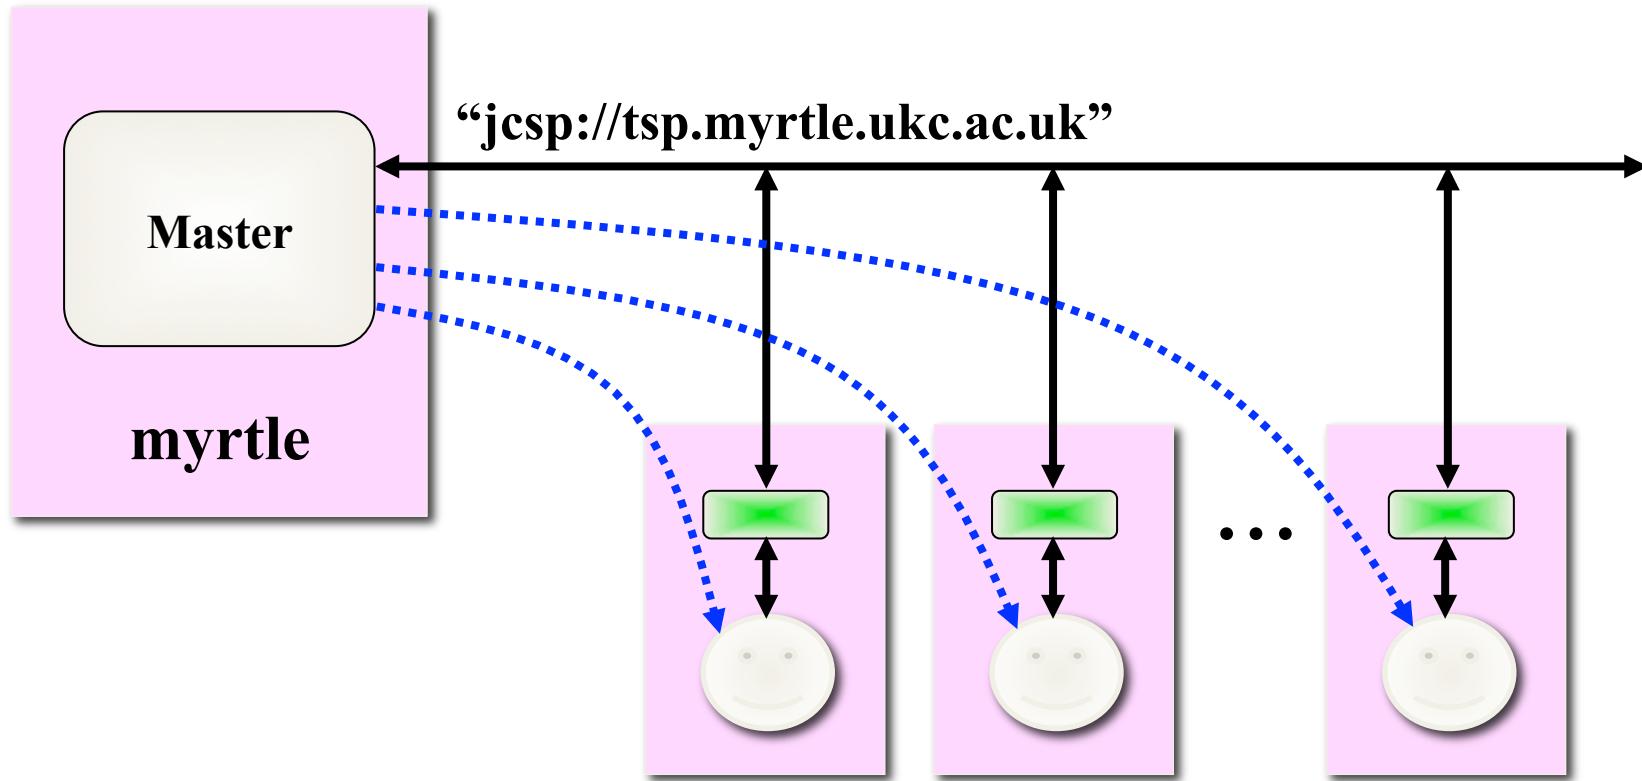


The SOR benchmark (7000 x 7000 matrix)



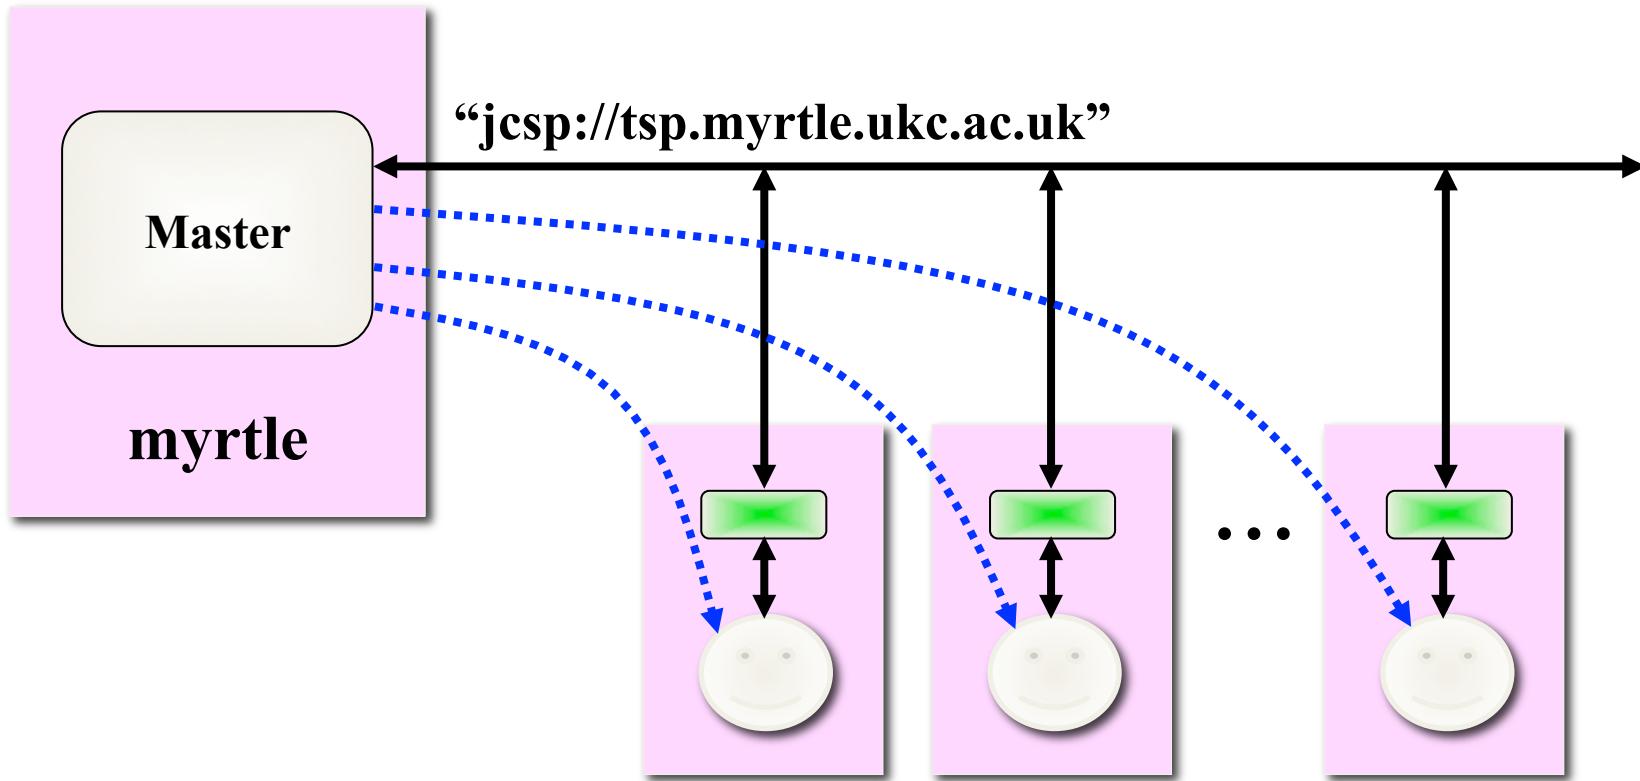
The Travelling Salesman Problem (15 cities)

Travelling Salesman Problem



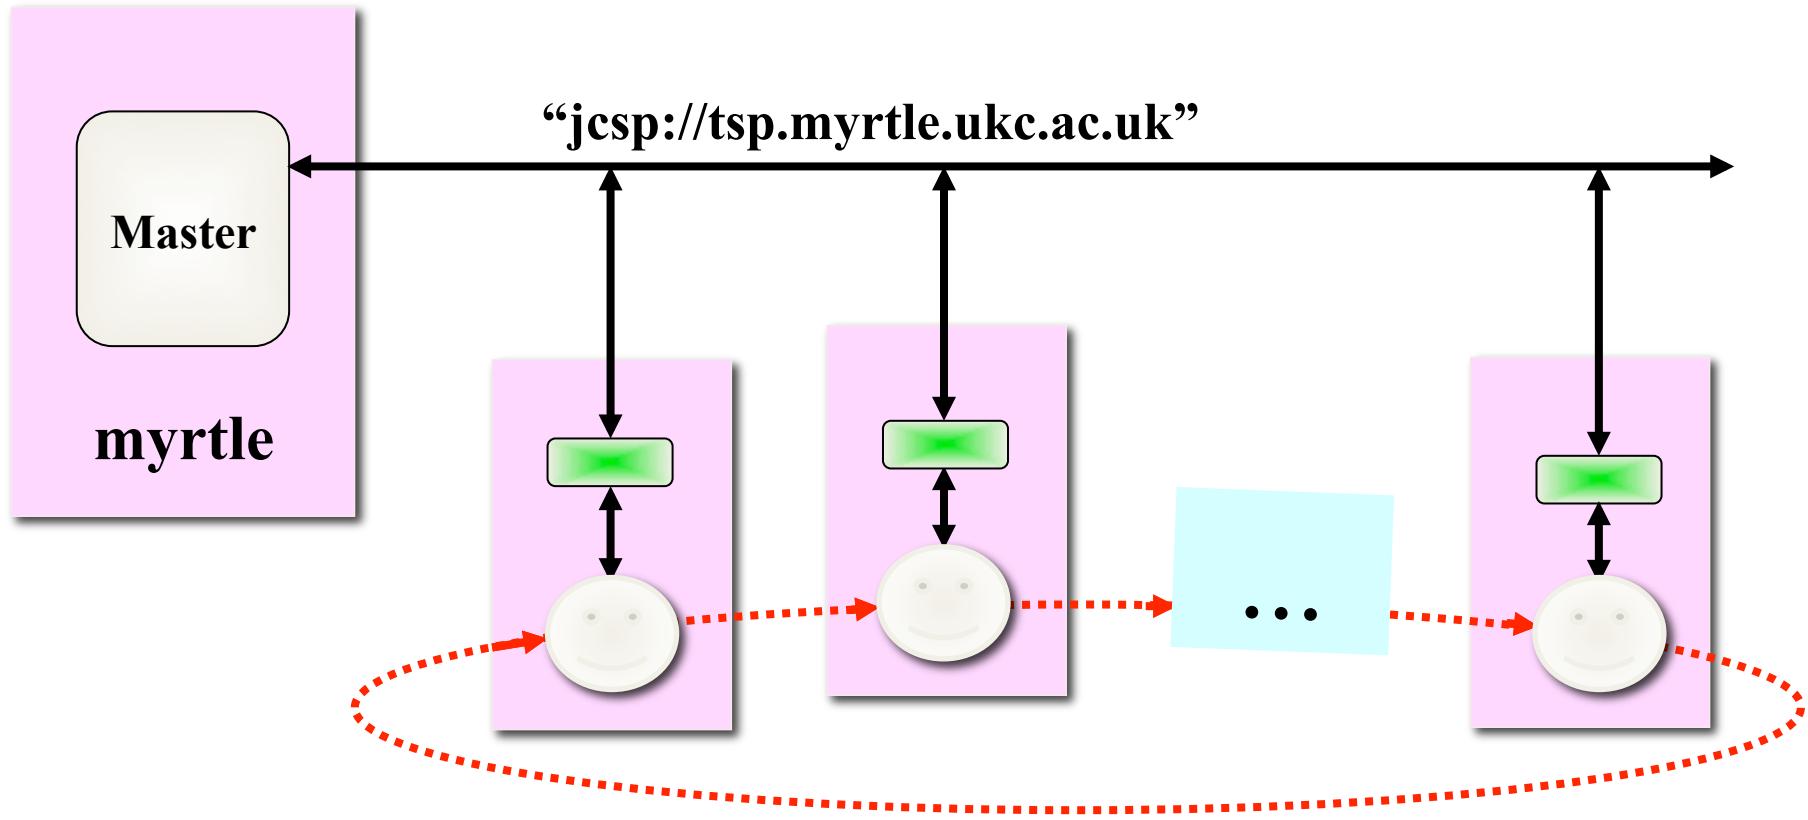
Currently, workers report newly discovered shorter paths back to the master (who maintains the global shortest). If master receives a better one, it broadcasts back to workers.

Travelling Salesman Problem



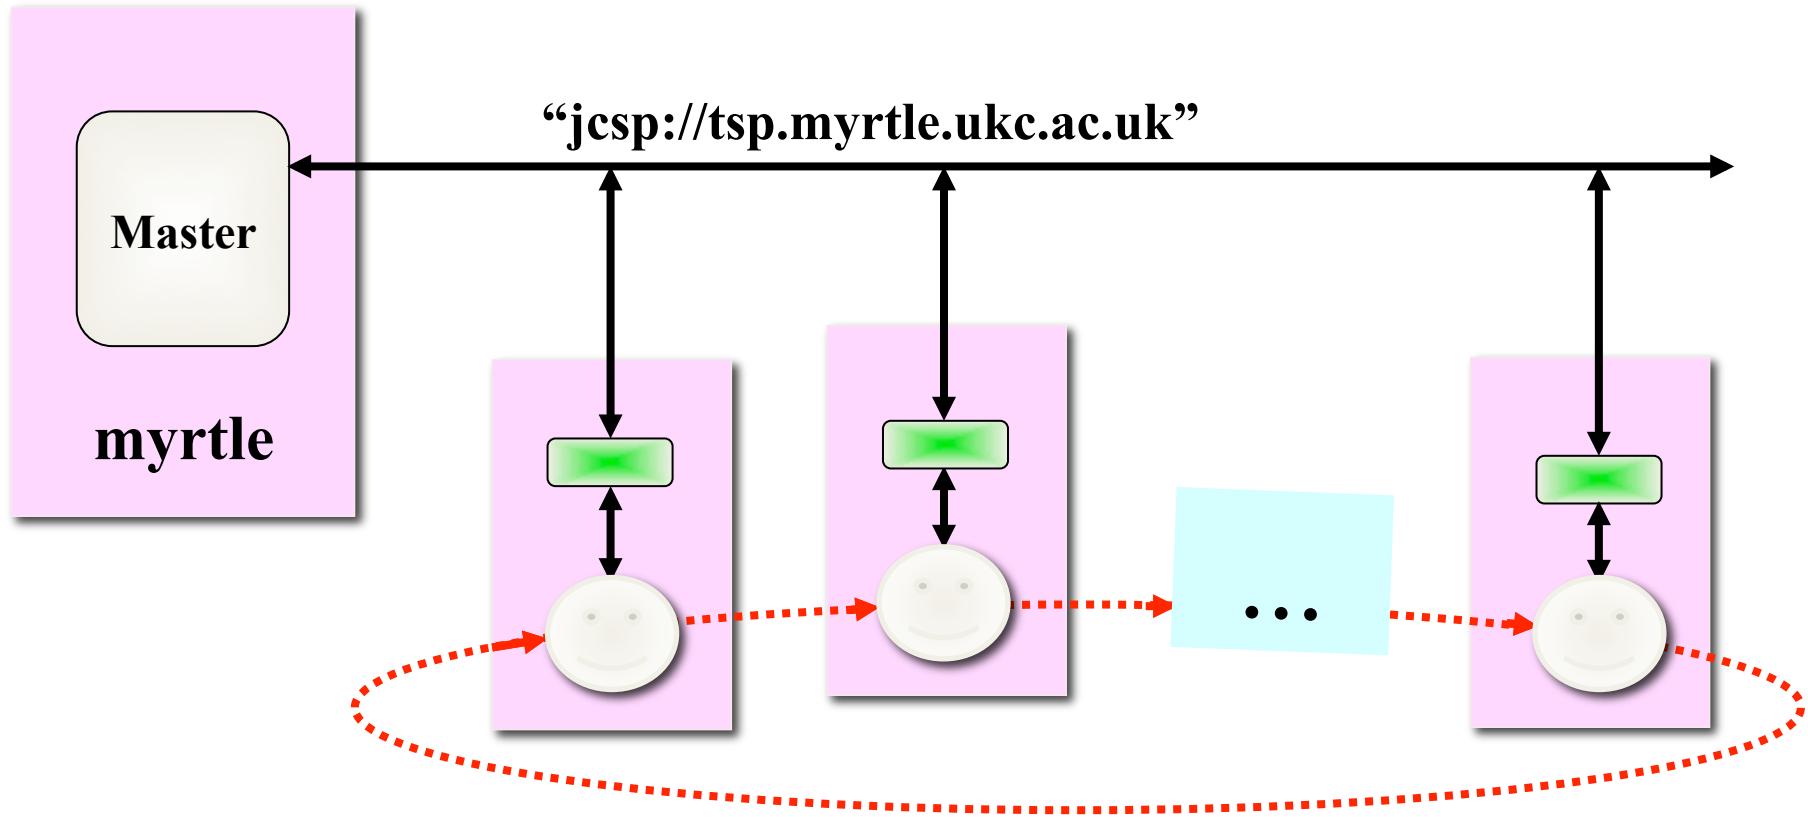
Massive swamping of network links in the early stages.
Also, generation of garbage *currently* provokes garbage collector
– and clobbers cache on dual-processor nodes.

Travelling Salesman Problem



Eliminate the broadcasting – control against swamping – stop generating garbage ... ☺ ☺ ☺

Travelling Salesman Problem

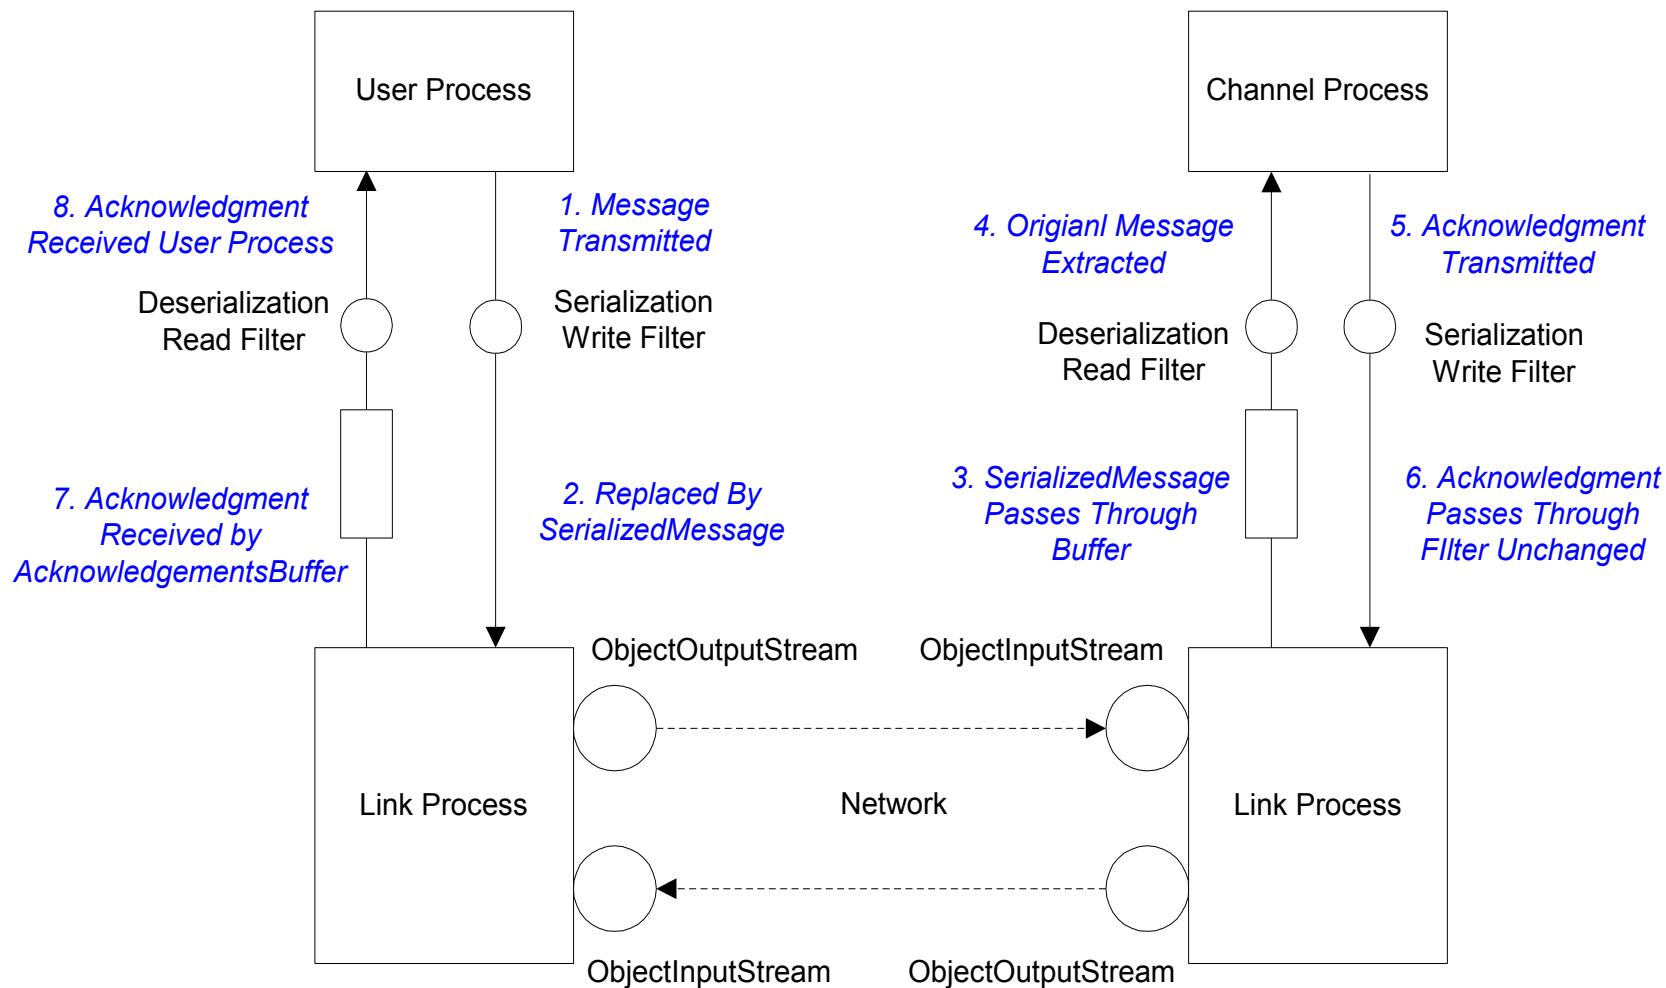


Global minimum maintained in ring (made with one-place
overwriting channel buffers) ... ***easy!!!***

Networked Class Loading

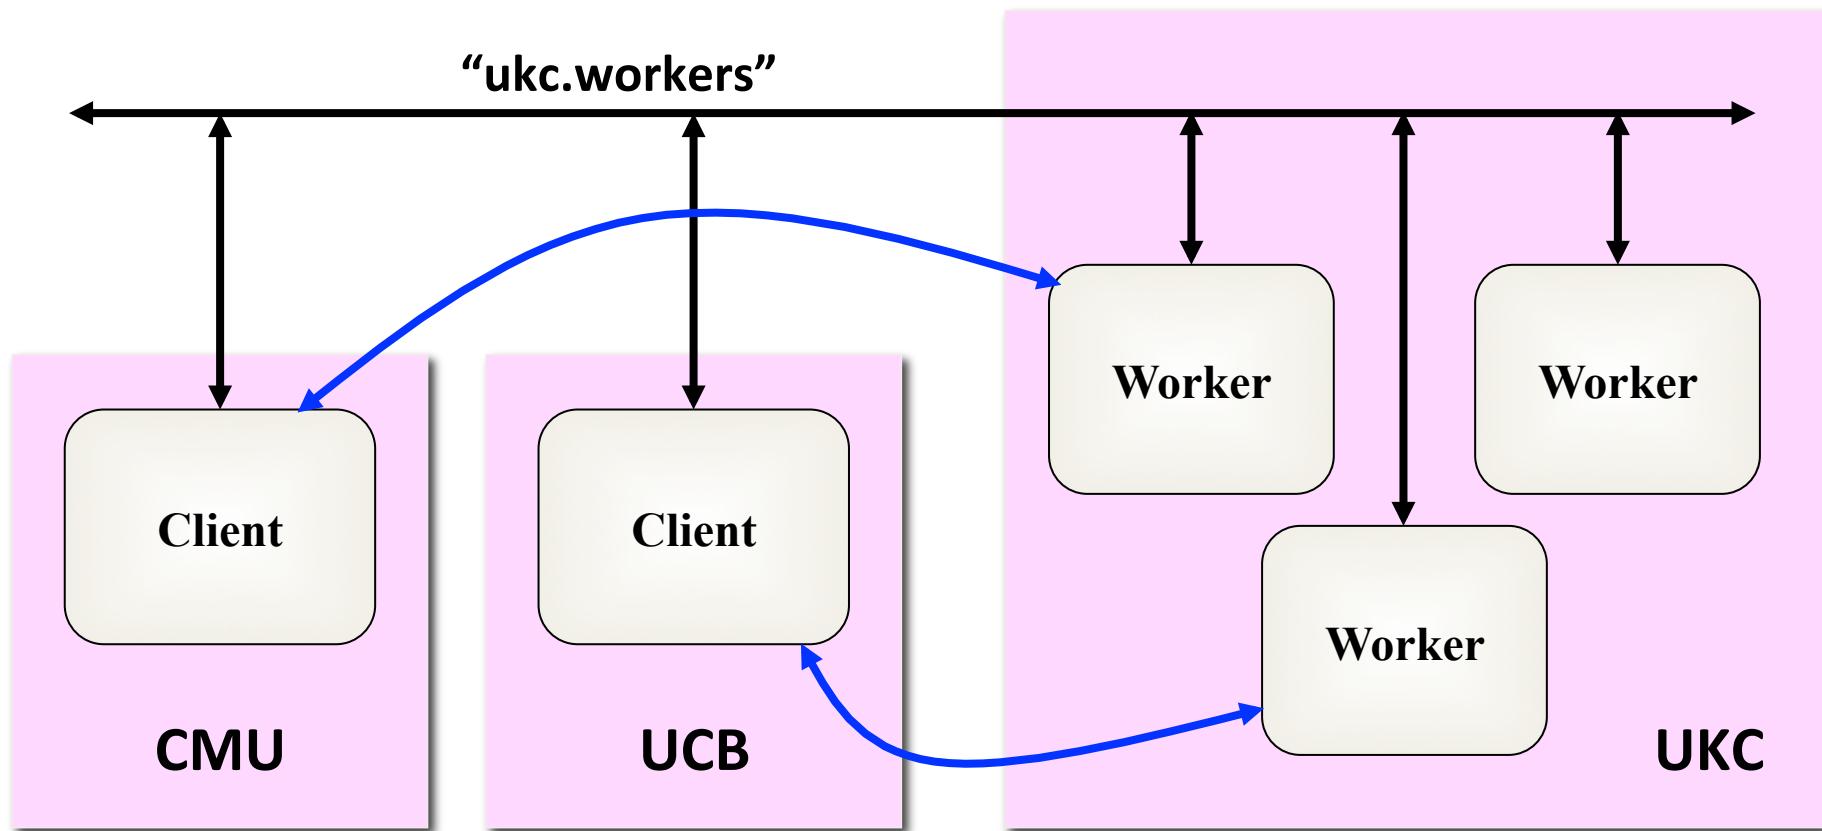
- By default, objects sent across a networked channel (or connection) use Java *serialization*.
- This means the receiving JVM is expected to be able to load (or already have loaded) the class files needed for its received objects.
- However, JCSP networked channels/connections can be set to communicate those class files *automatically* (if the receiver can't find them locally).
- Machine nodes cache those class files locally in case they themselves need to forward them.

Networked Class Loading



Remote Process Launching

- Example: UKC offers a simple *worker farm* ...
- *Clients* grab available *workers* ...

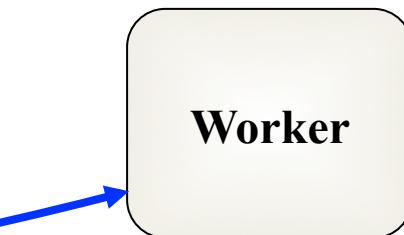
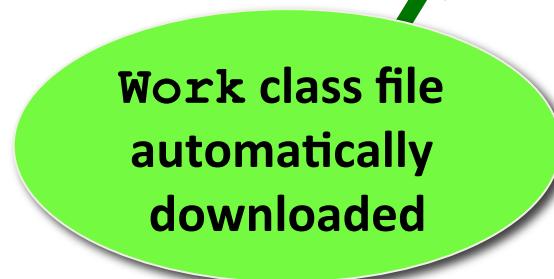


Remote Process Launching

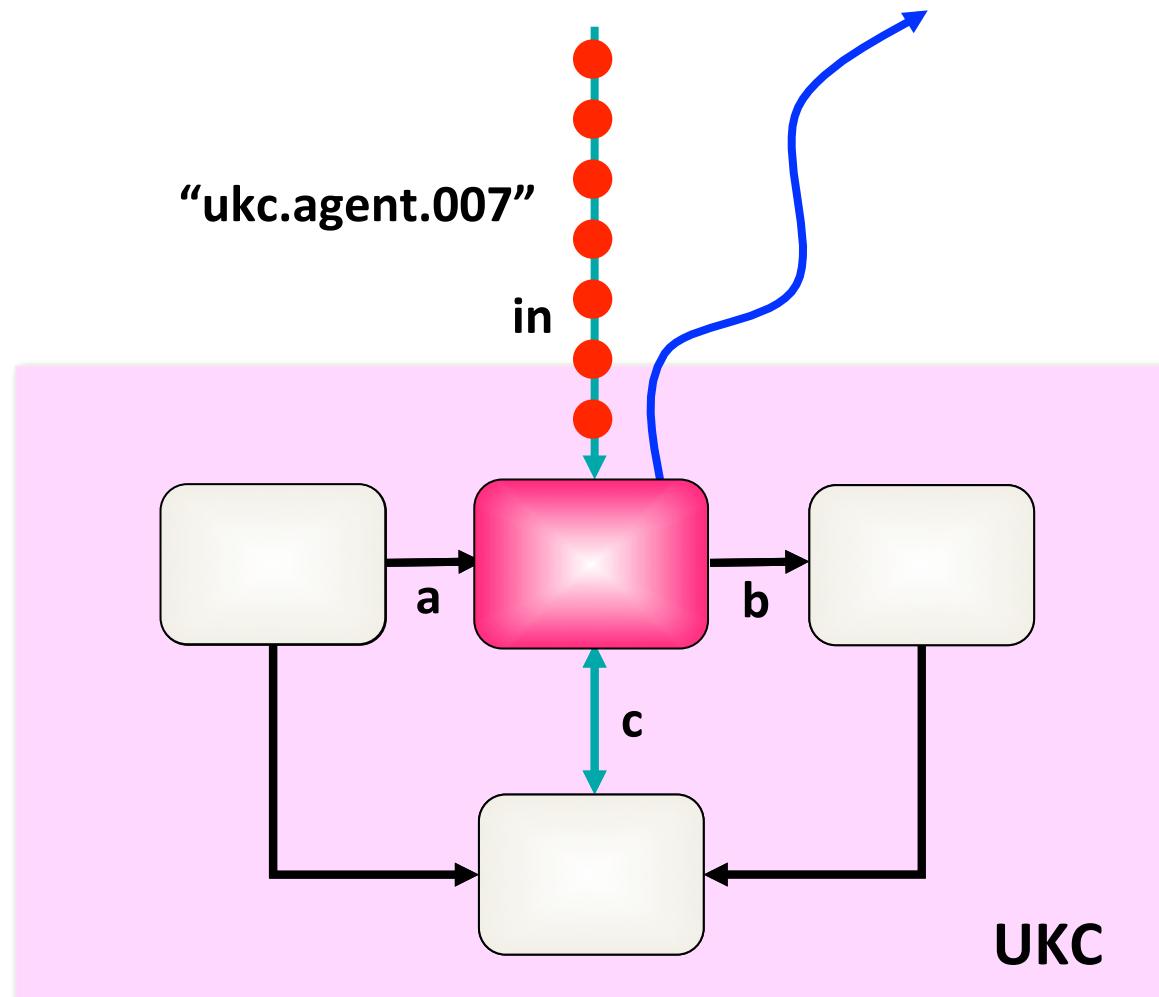
```
Work work = new Work () ; // CSProcess  
out.request (work) ;  
work = (Work) out.reply () ;
```



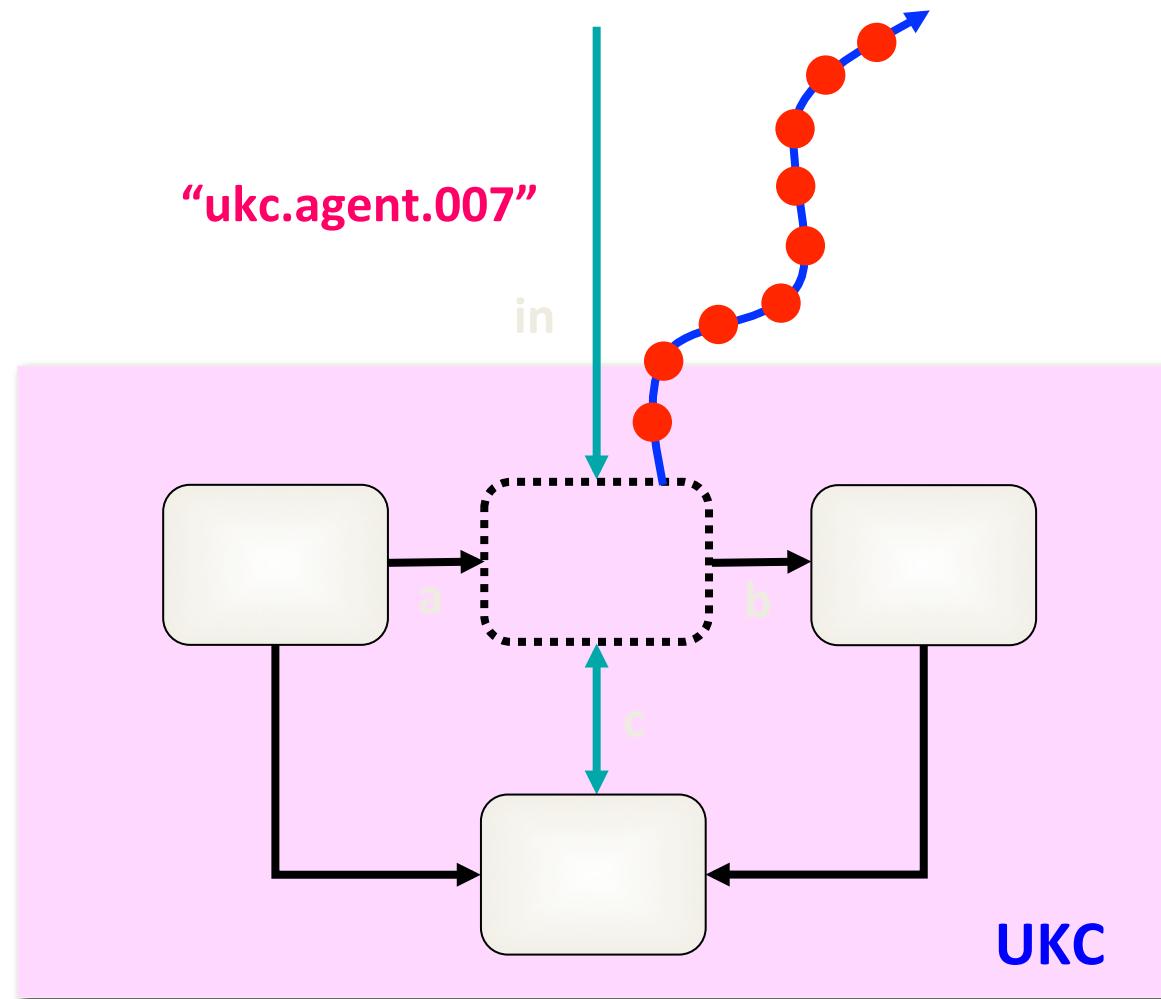
```
CSProcess work = (CSProcess) in.request () ;  
work.run () ;  
in.reply (work) ;
```



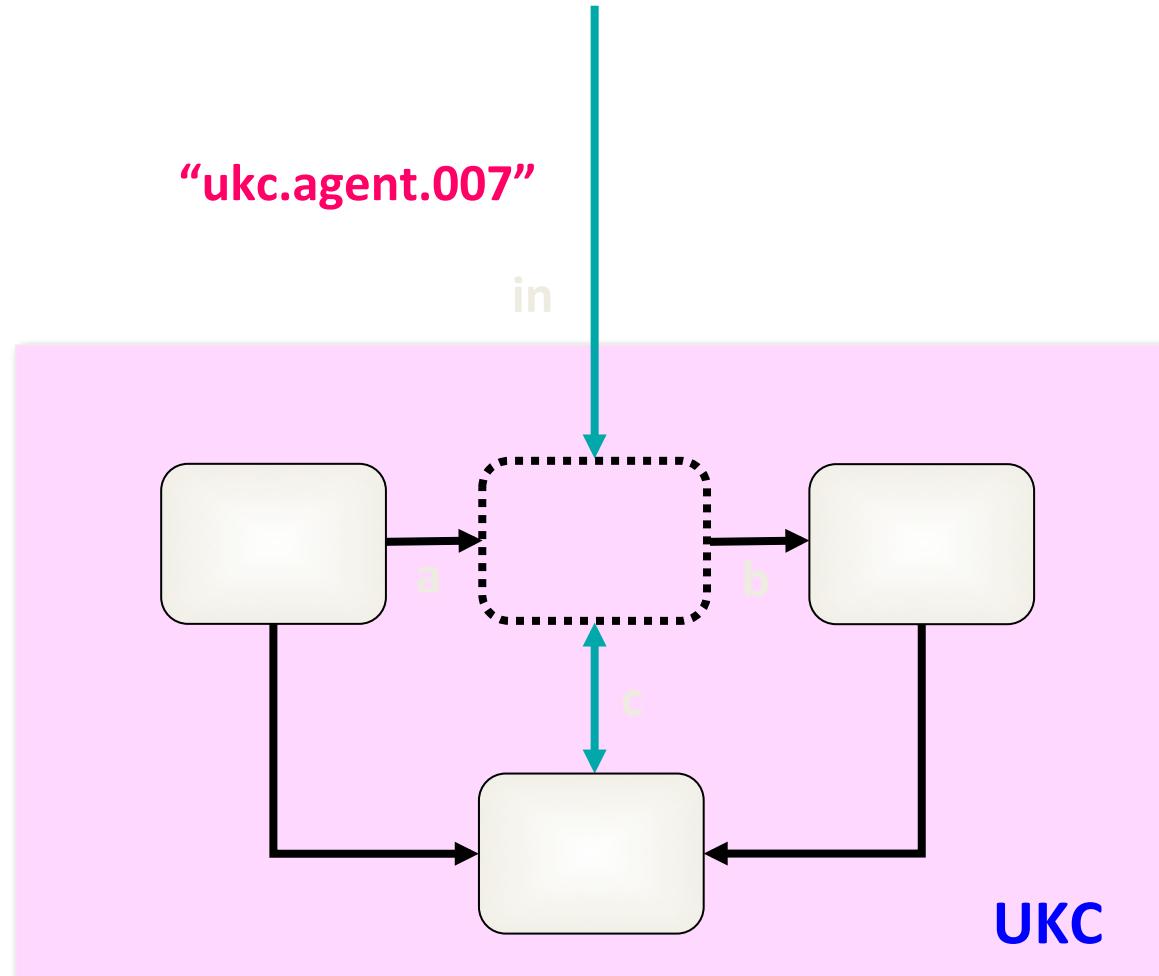
Mobile Processes (Agents)



Mobile Processes (Agents)

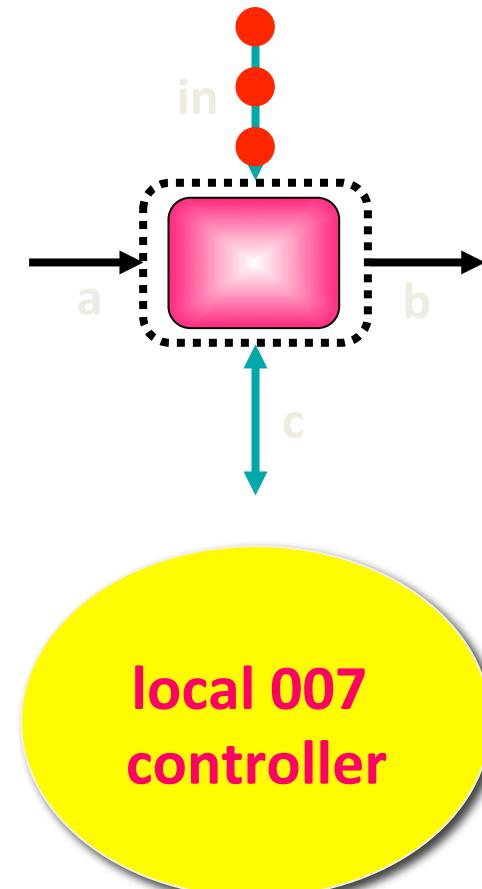


Mobile Processes (Agents)



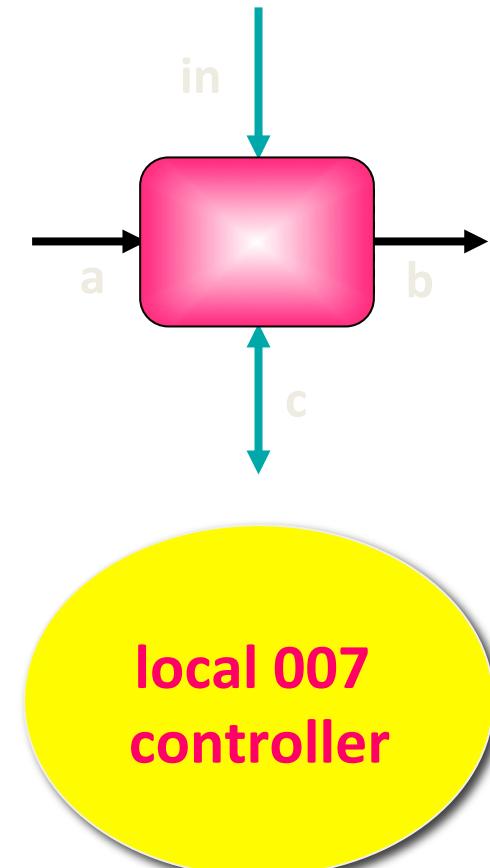
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}
```



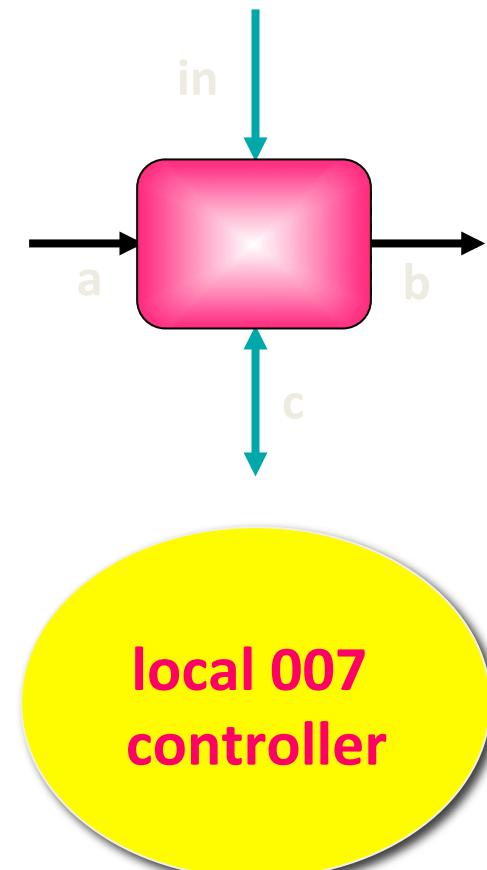
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}
```



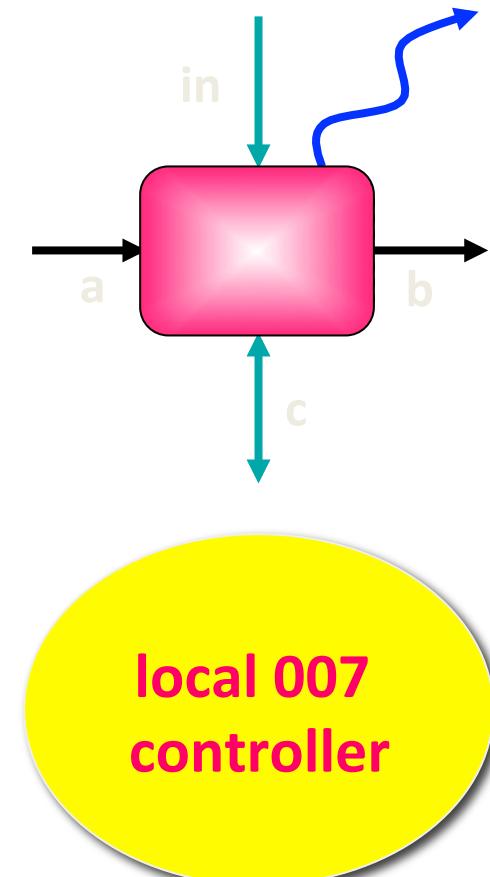
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}
```



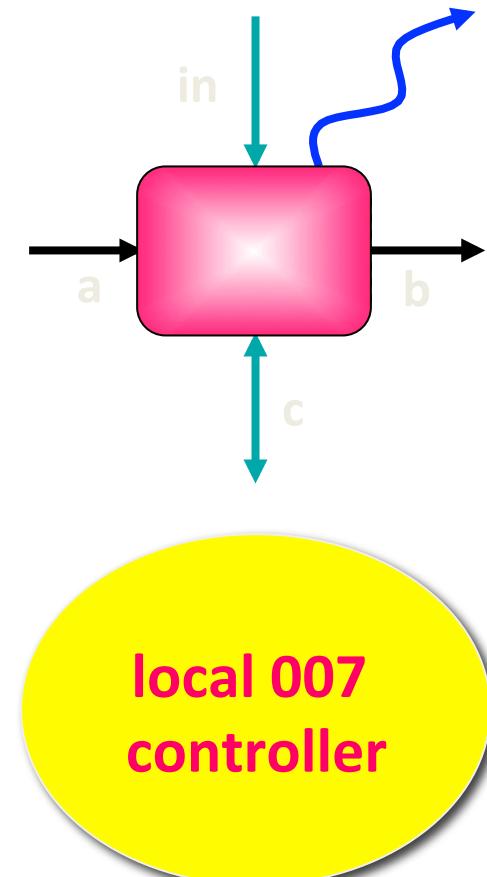
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}  
}
```



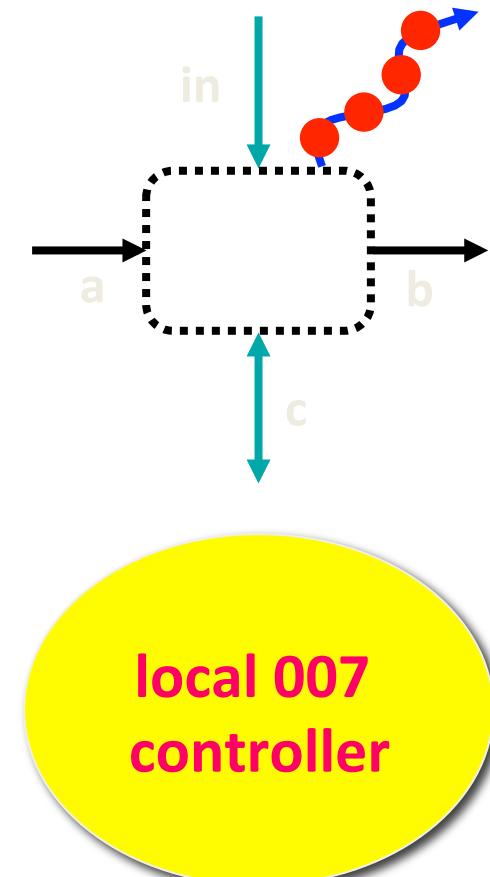
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}  
}
```



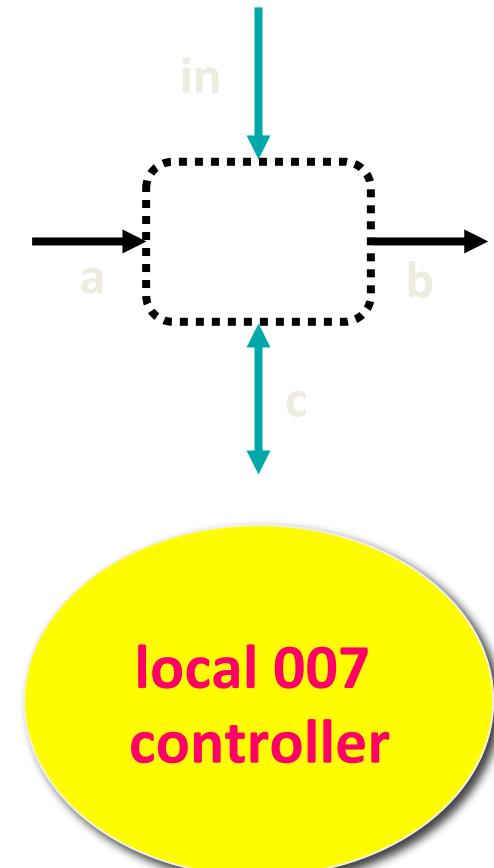
Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}  
}
```



Mobile Processes (Agents)

```
while (running) {  
    Bond james = (Bond) in.read ();  
    james.plugin (a, b, c);  
    james.run ();  
    NetChannelLocation escapeRoute =  
        james.getNextLocation ();  
    One2NetChannel escape =  
        new One2NetChannel (escapeRoute);  
    running = james.getNuke ();  
    escape.write (james);  
    escape.disconnect ();  
}  
}
```

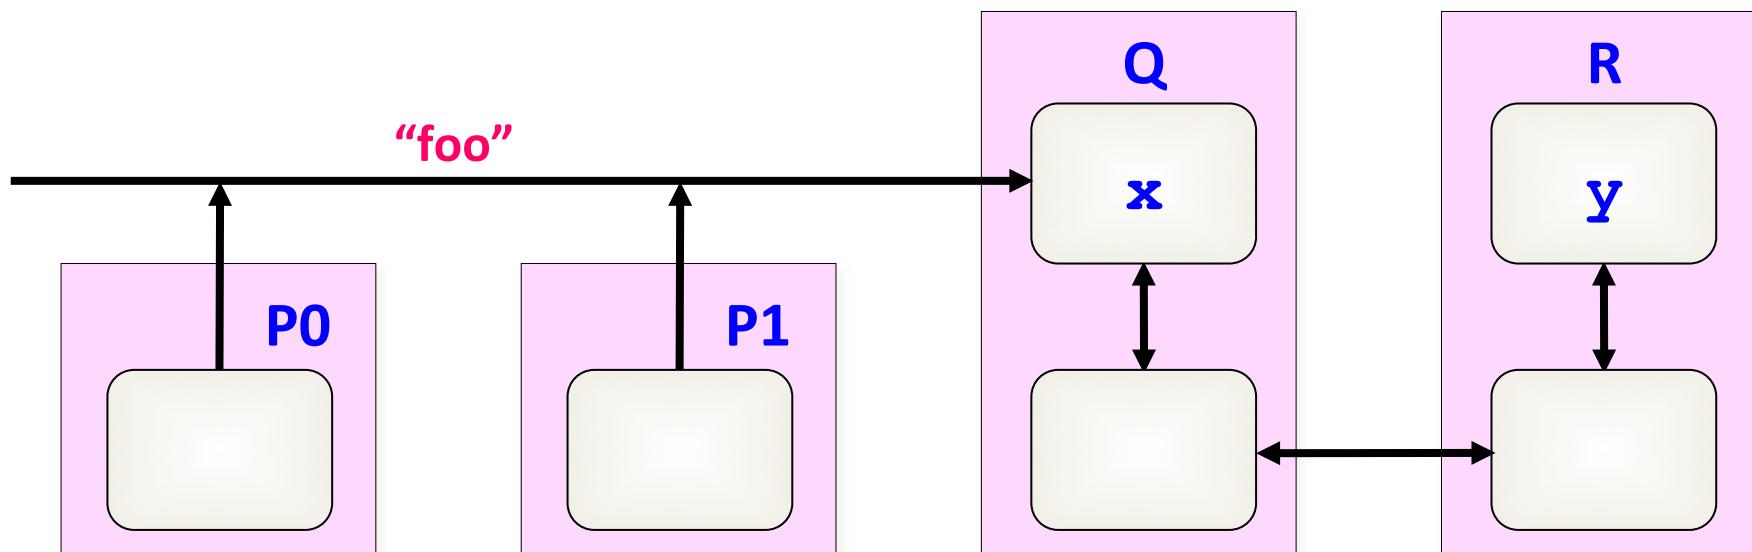


Mobile Network Channels

- Channel *ends* may be moved around a network.
- This is potentially dangerous as we are changing network topology, which may introduce deadlock - ***considerable care must be taken***.
- There is nothing special to do to migrate channel *write-ends*. Network channels are naturally *any-one*. All that is needed is to communicate the *CNS channel name* (or **NetChannelLocation**) to the new writer process.
- Migrating channel *read-ends* securely requires a special protocol ...

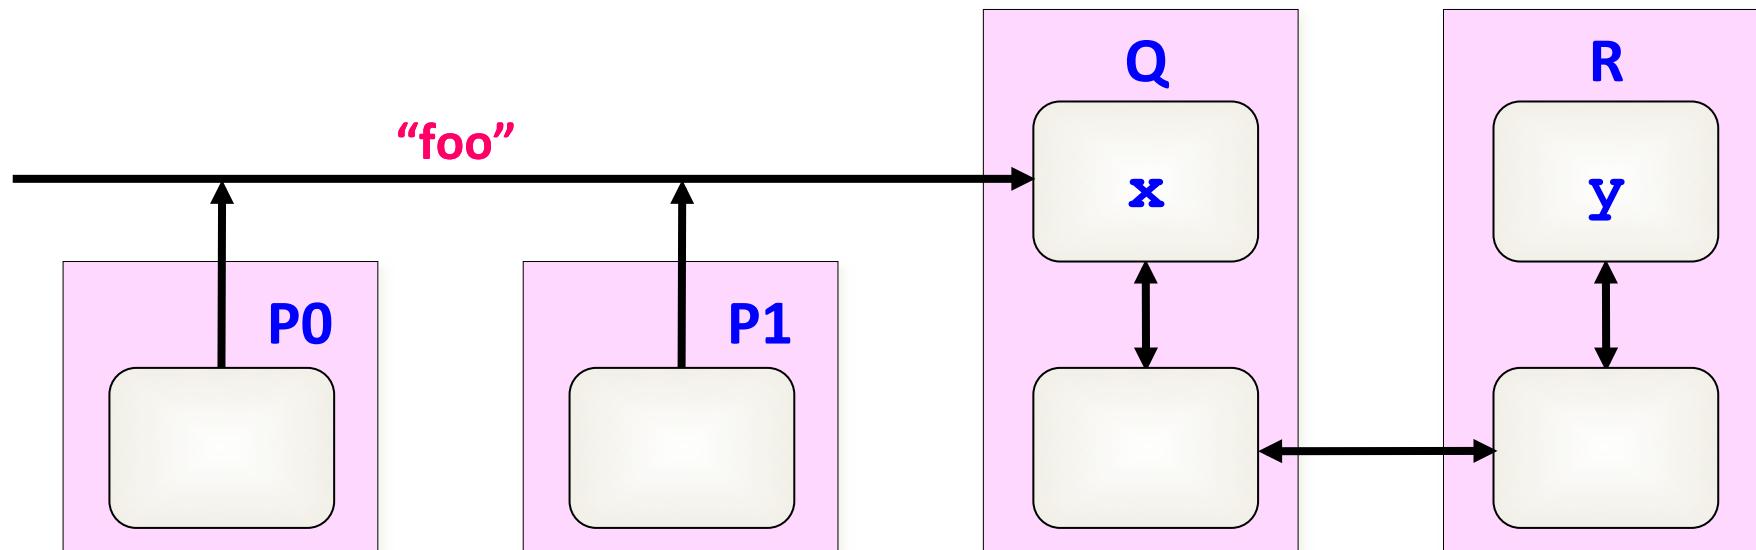
Mobile Network Channels

- Consider a process, **x**, on node **Q**, currently servicing the *CNS-registered* channel “**foo**”.
- It wants to pass on this responsibility to a (willing) process, **y**, in node **R**, with whom it is in contact.



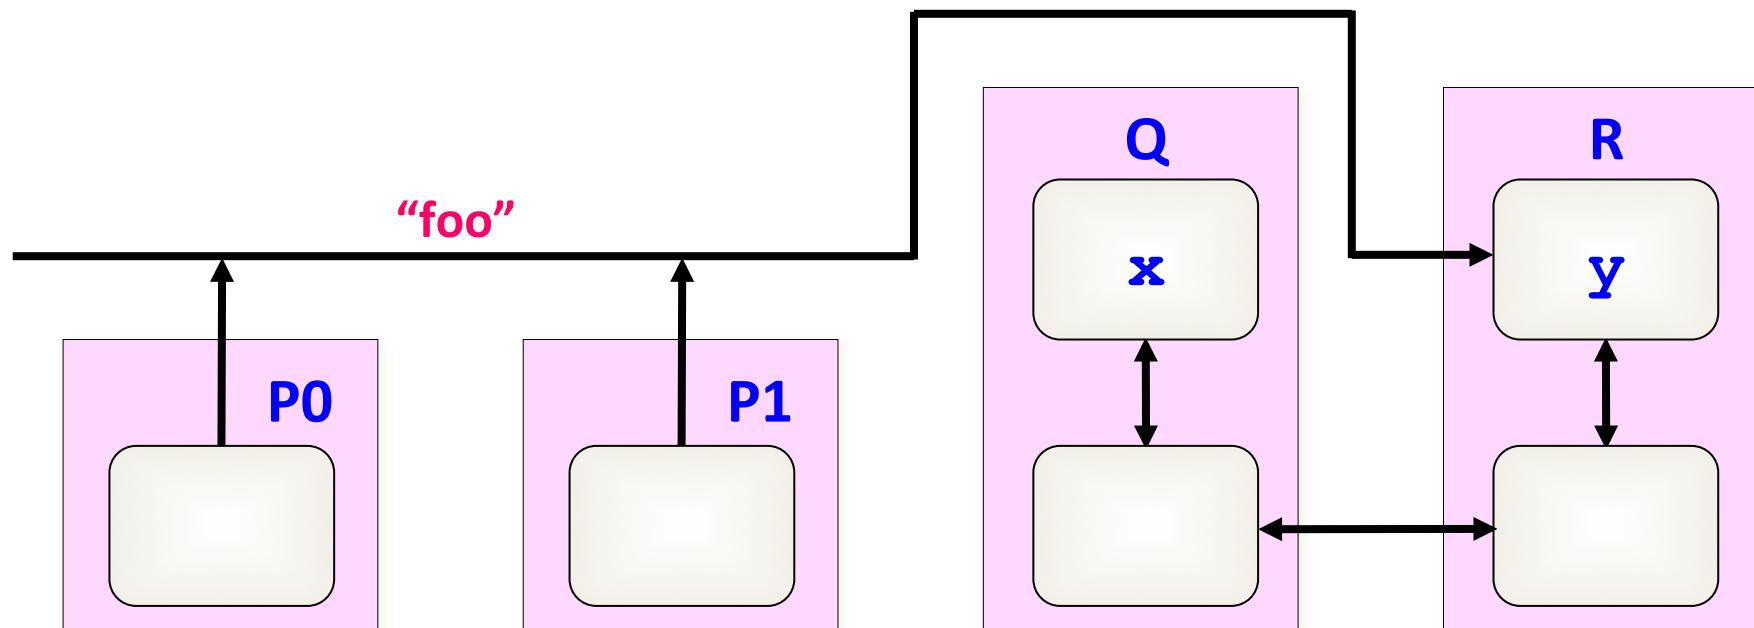
Mobile Network Channels

- Processes writing to “**foo**” are to be unaware of this channel migration.



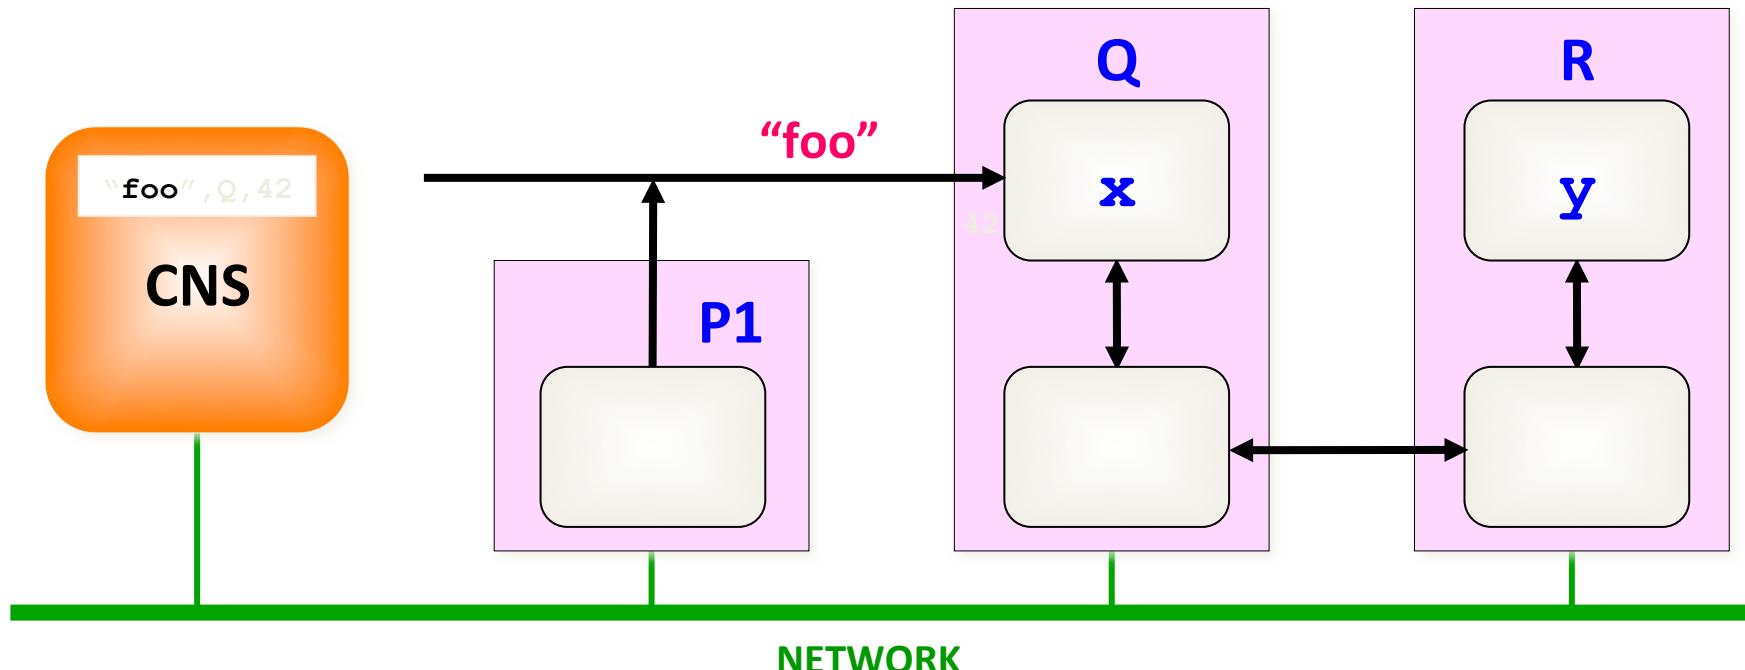
Mobile Network Channels

- Processes writing to “**foo**” are to be unaware of this channel migration.



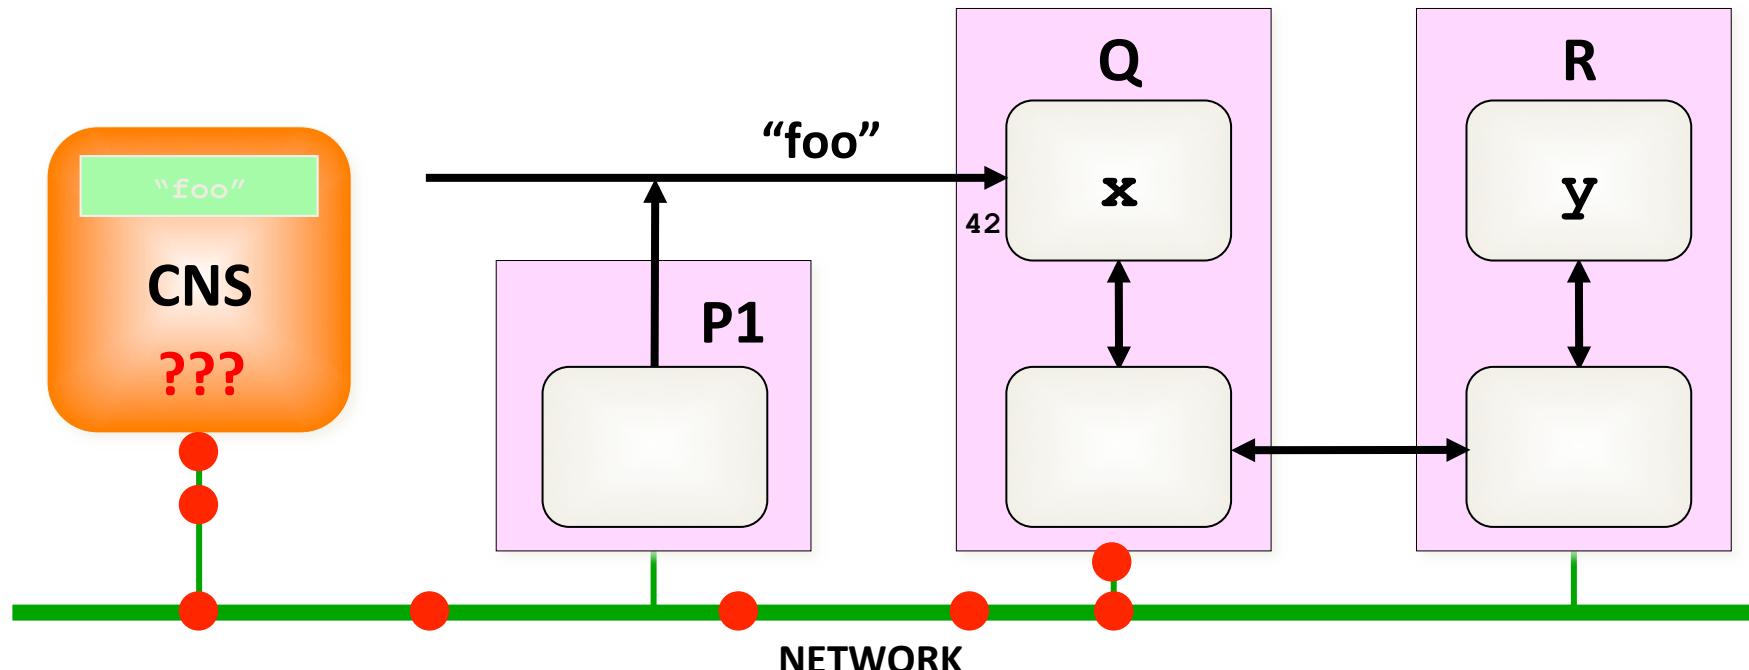
Mobile Network Channels

- Let's get back to the initial state ("foo" being serviced by **x** on node **Q**).
- Let's show the network ... and the CNS ...



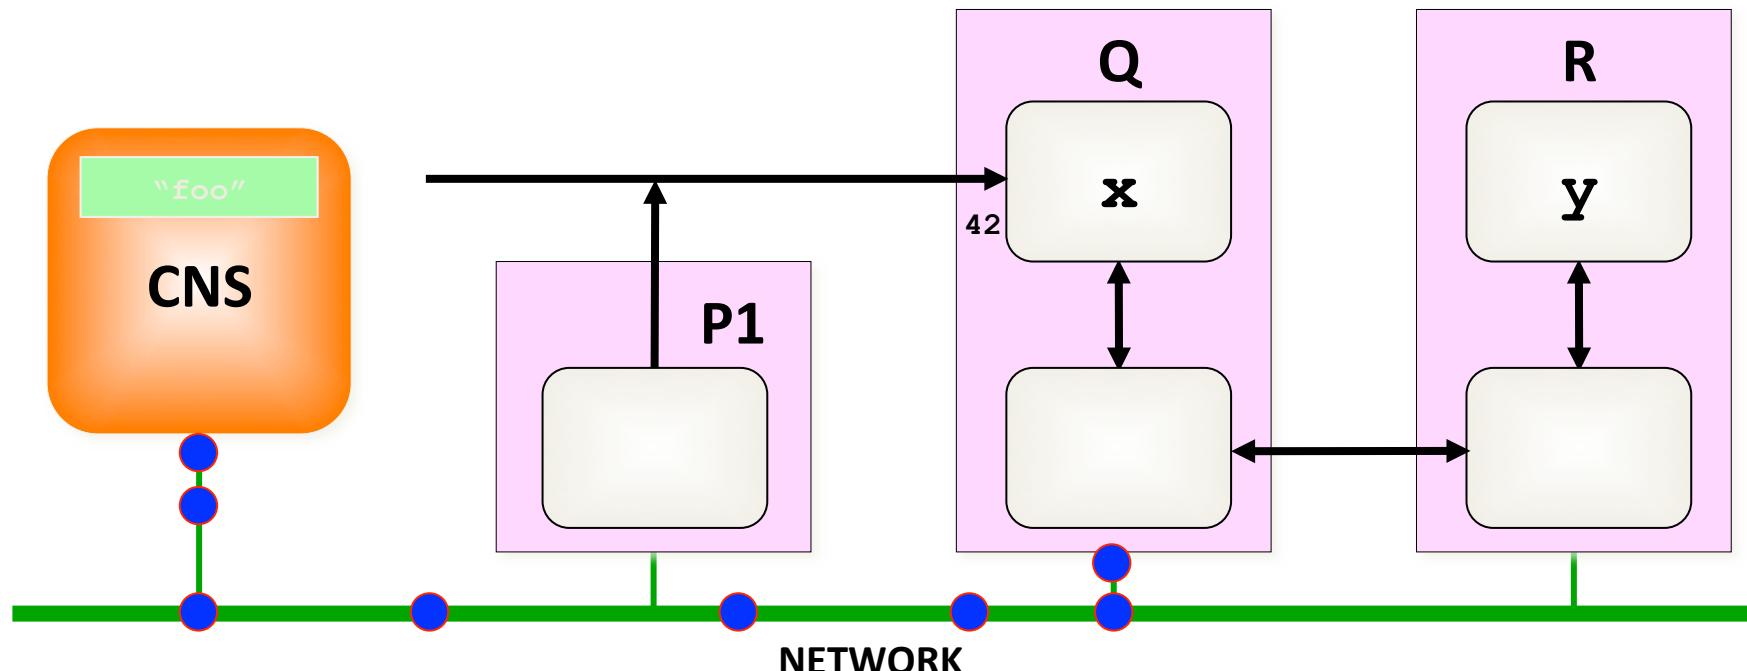
Mobile Network Channels

- First, process **x** *freezes* the name “**foo**” on the CNS ...



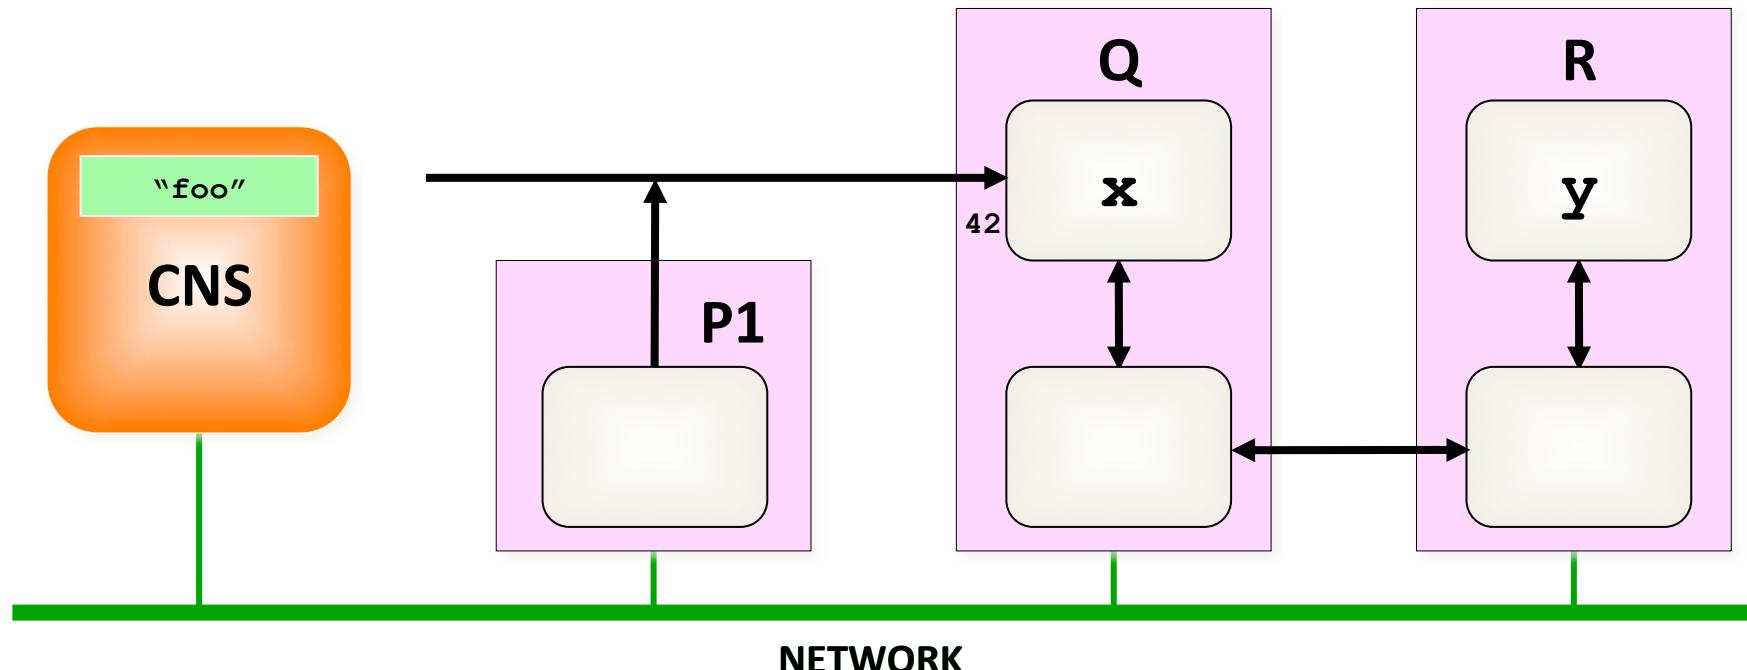
Mobile Network Channels

- First, process **x** *freezes* the name “**foo**” on the CNS ...
- The CNS returns an *unfreeze key* to process **x** ...



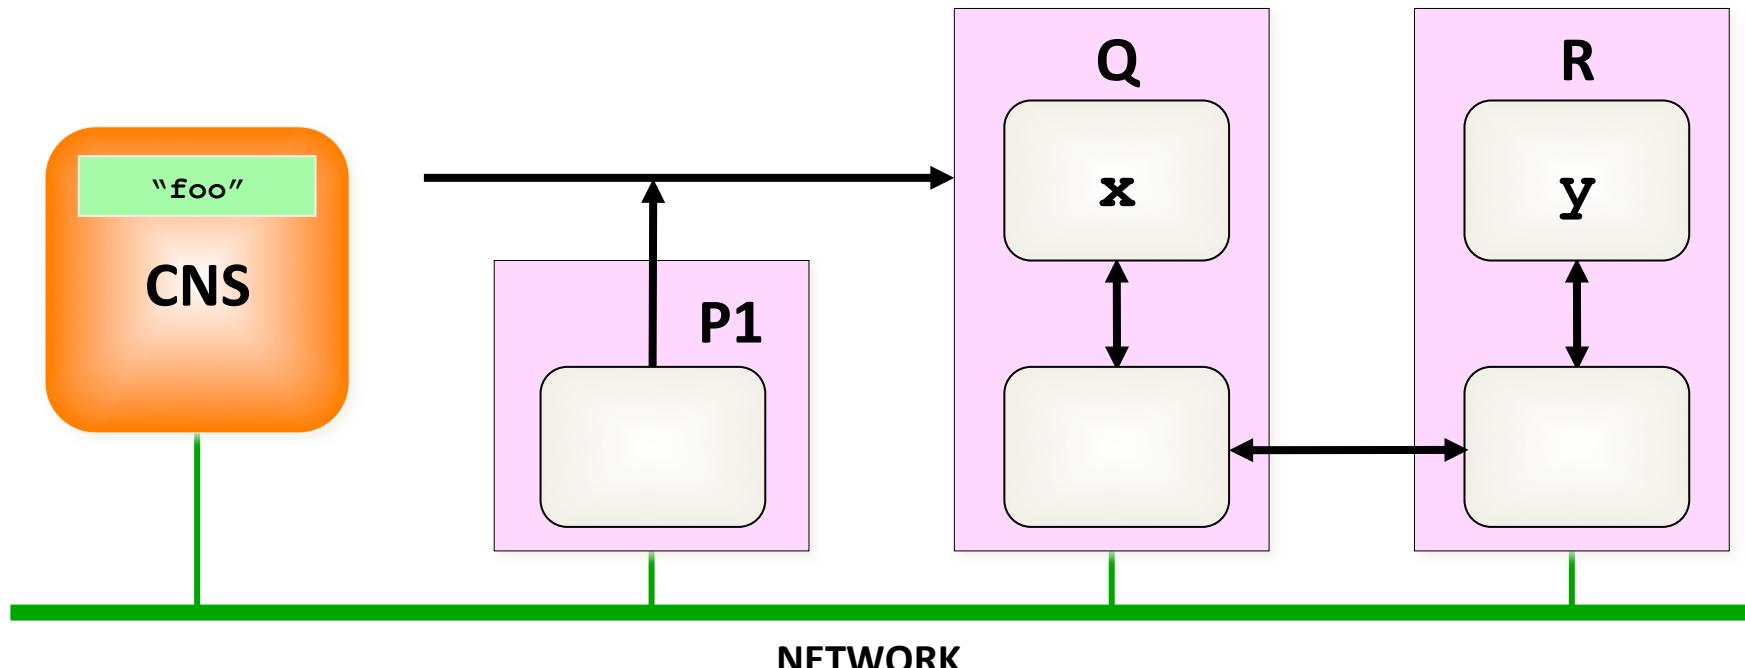
Mobile Network Channels

- The CNS no longer resolves “**foo**” for new writers and also disallows new registrations of the name.
- The network channel is deleted from processor **Q**.



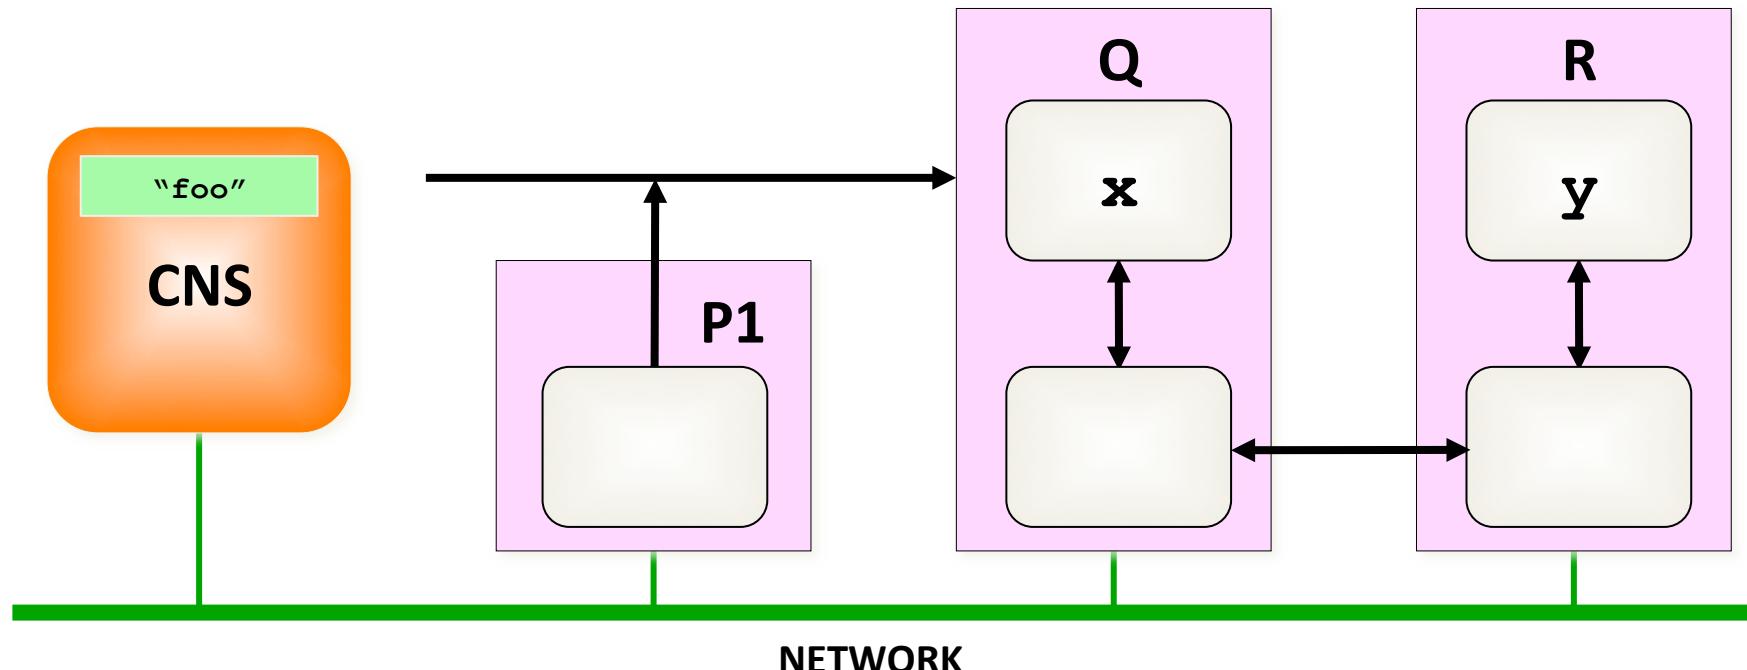
Mobile Network Channels

- The network channel is deleted from processor **Q**.
- Any *pending* and *future* messages for that channel (42) on **Q** are bounced (`NetChannelIndexException`).



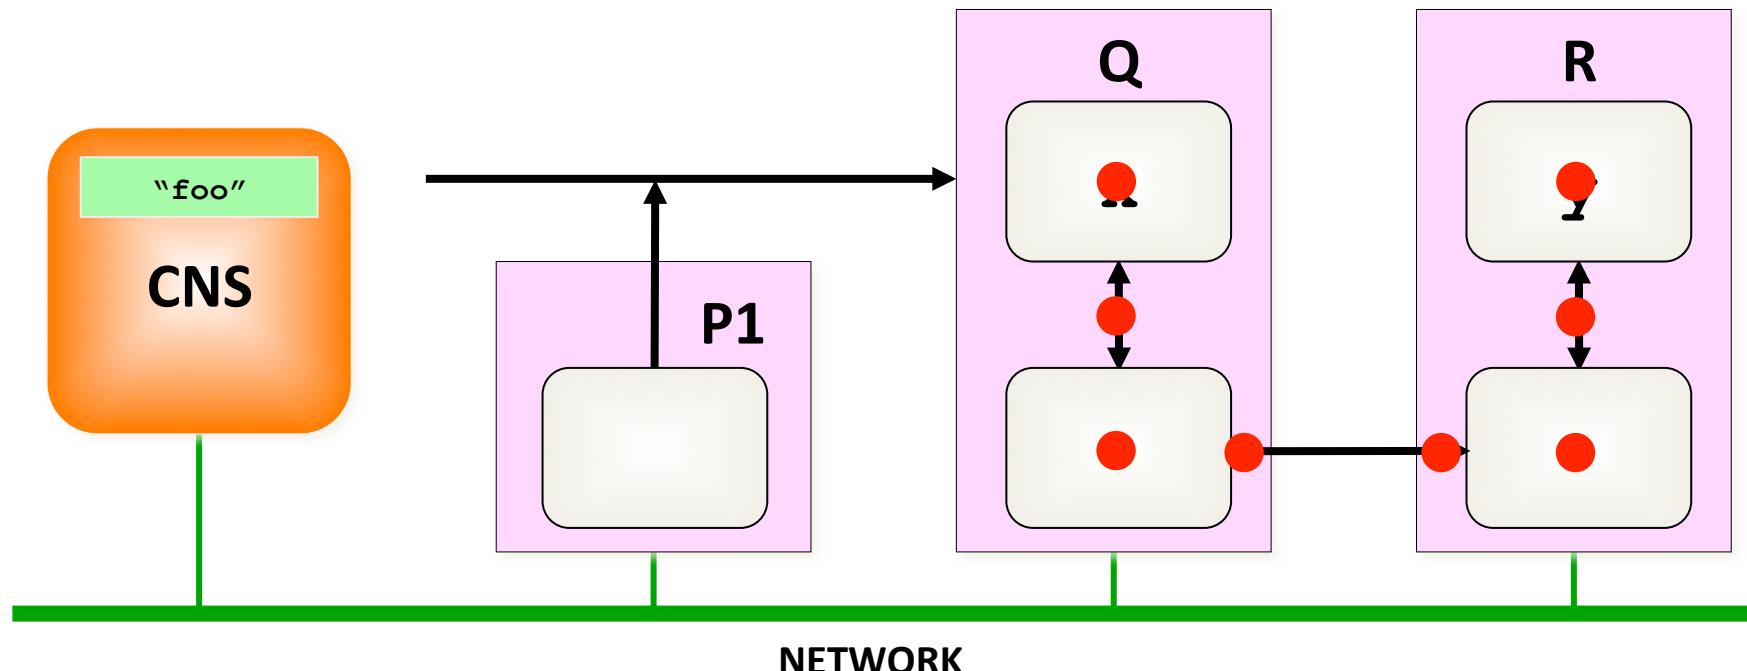
Mobile Network Channels

- The `write()` method at **P1** handles that bounce by appeal to the CNS for the new location of “**foo**”.
- This will not succeed until ...



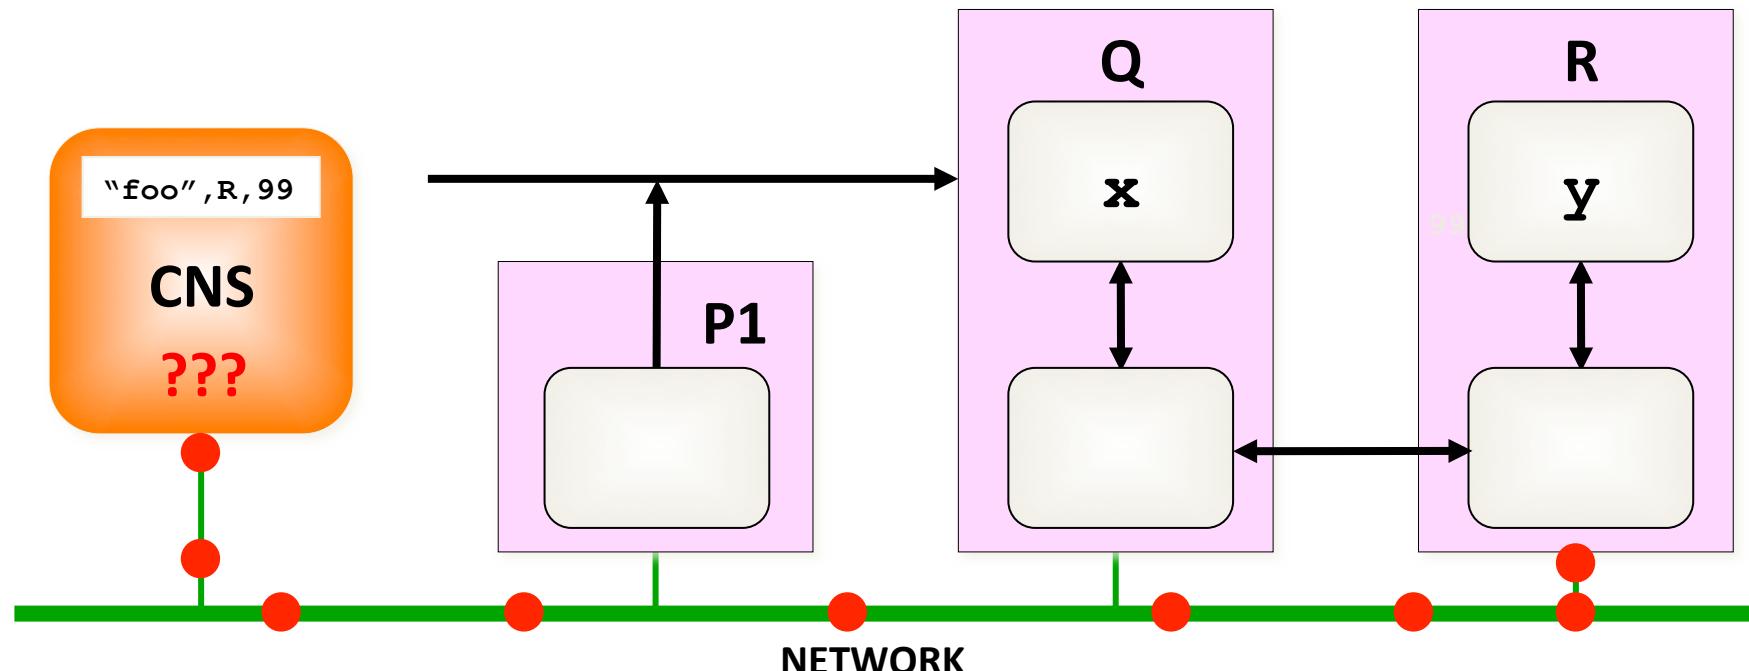
Mobile Network Channels

- ... process **x** (on node **Q**) passes on the channel name (“**foo**”) and *CNS unfreeze key* ...



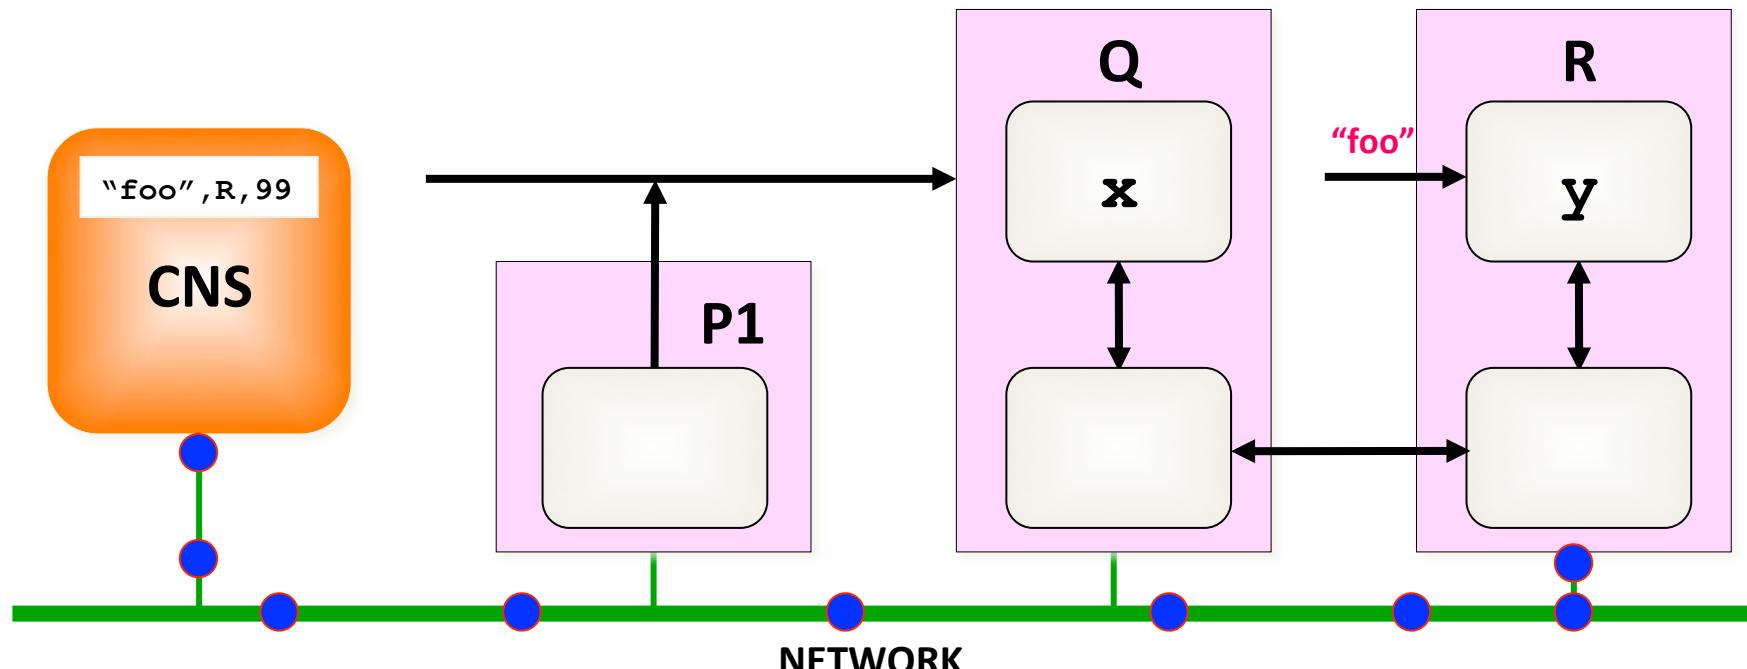
Mobile Network Channels

- ... process **x** (on node **Q**) passes on the channel name (“**foo**”) and CNS *unfreeze key* ...
- ... and the receiver (process **y** on **R**) unlocks the name “**foo**” (using the *key*) and re-registers it.



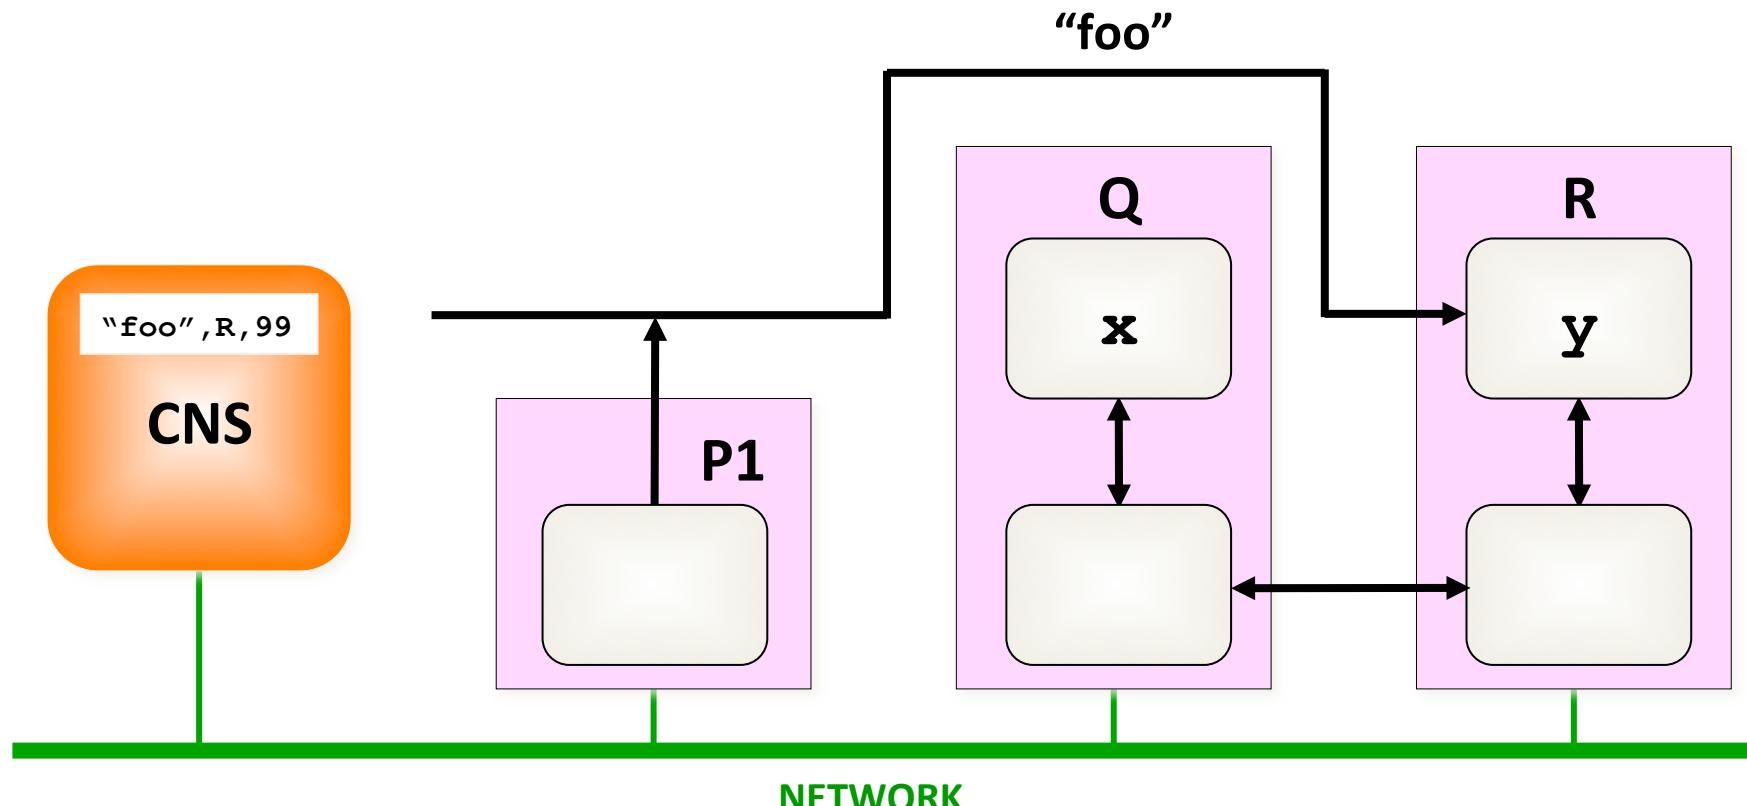
Mobile Network Channels

- ... and the receiver (process **y** on **R**) unlocks the name “**foo**” (using the **key**) and re-registers it.
- The **write()** method at **P1** now hears back from the CNS the new location of “**foo**” ...



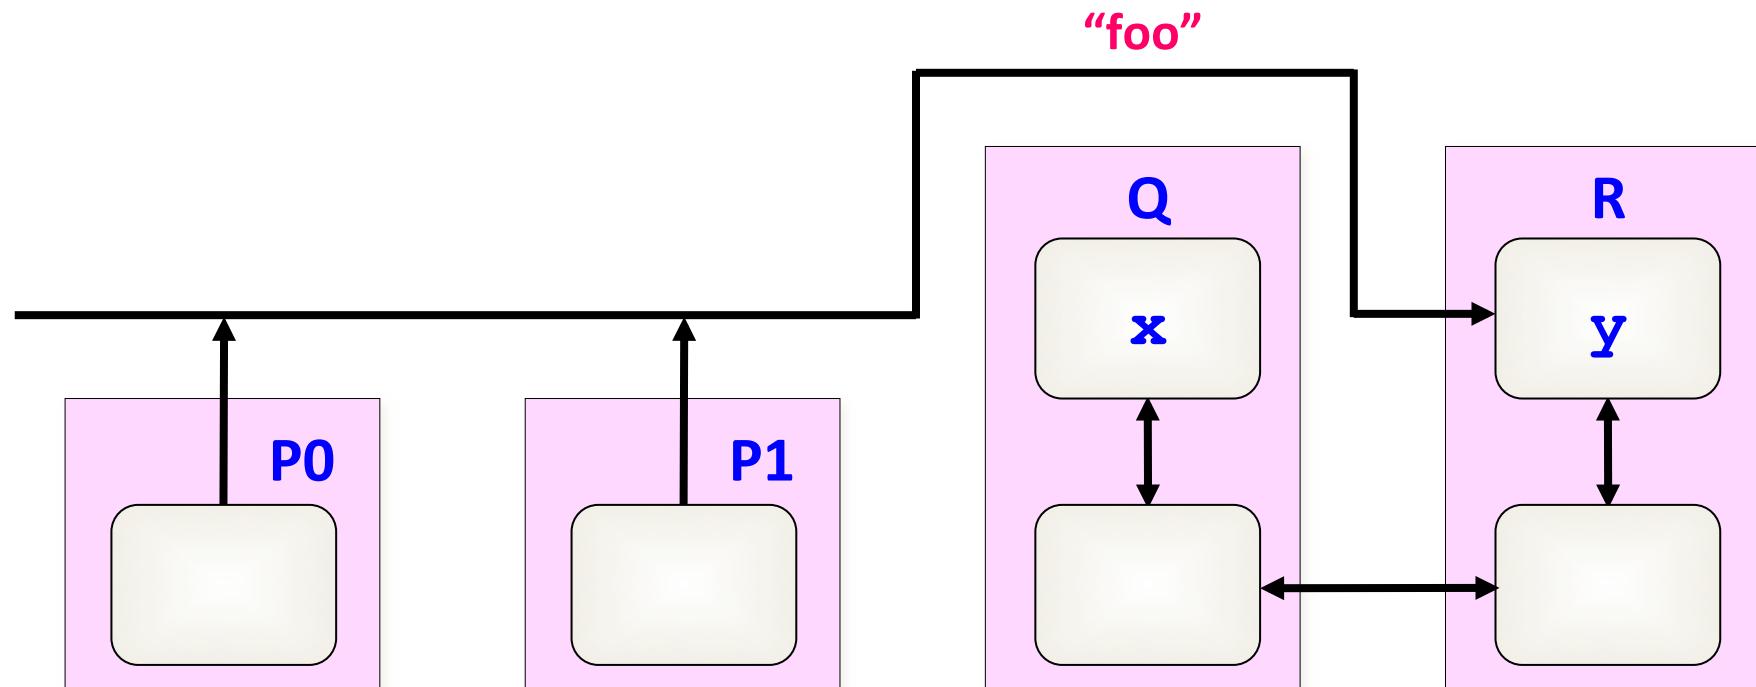
Mobile Network Channels

- ... and resends the message that was bounced.
- The writing process(es) at **P1** (and elsewhere) are unaware of the migration.



Mobile Network Channels

- ... and resends the message that was bounced.
- The writing process(es) at **P1** (and elsewhere) are unaware of the migration.



Mobile Network Connections

- Connection *ends* may be moved around a network.
- This is potentially dangerous as we are changing network topology, which may introduce deadlock - ***considerable care must be taken***.
- There is nothing special to do to migrate connection *client-ends*. Network connections are naturally *any-one*. All that is needed is to communicate the *CNS connection name* (or **NetConnectionLocation**) to the new writer process.
- Migrating *server-ends* safely requires a special protocol ... **the same as for channel write-ends**.

Summary

- **JCSP.net** enables *virtual channel communication* between processes on separate machines (JVMs).
- Application *channels/connections* between machines are set up (and taken down) dynamically.
- Channels/connections are multiplexed over *links*.
- Links can be developed for *any network protocol* and plugged into the **JCSP.net** infrastructure.
- No central management – *peer-to-peer* connections (bootstrapped off a basic *Channel Name Server*).
- Brokers for *user-definable matching services* are easy to set up as ordinary application servers.

Summary

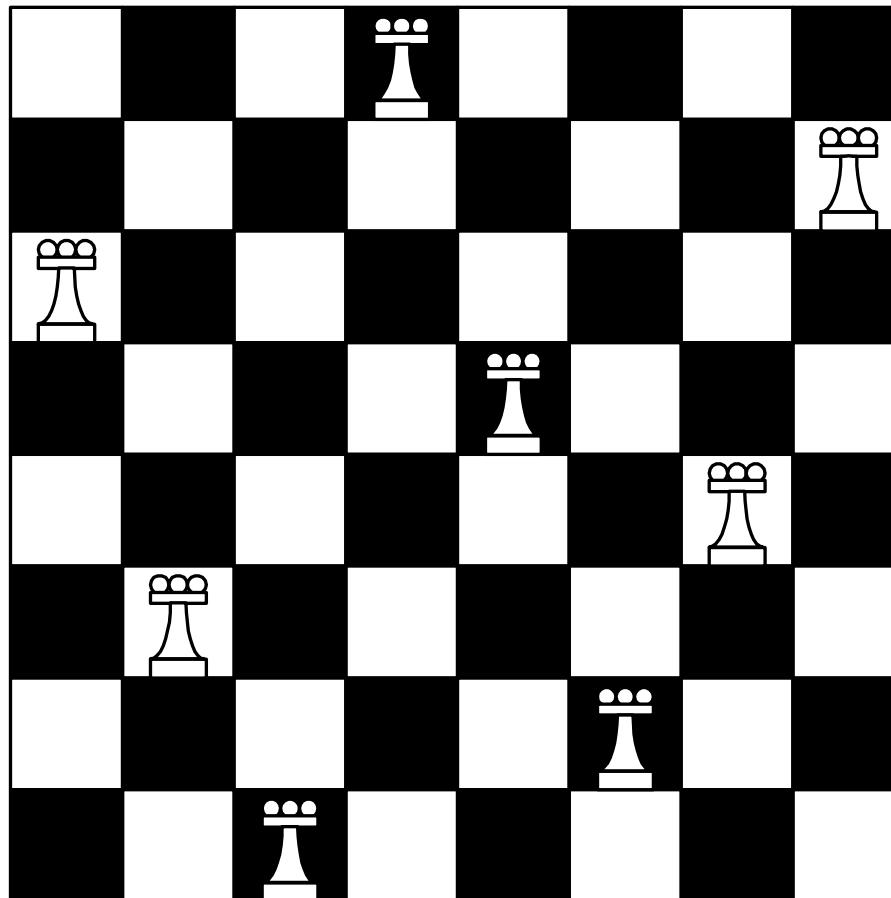
- Processes can *migrate* between processors (with classes loaded dynamically as necessary) – hence *mobile agents, worker farms, grid computation* ...
- *JCSP.net* provides *exactly the same (CSP/occam)* concurrency model for networked systems as *JCSP* provides within each physical node of that system.
- Network logic is *independent of physical distribution* (or even whether it is distributed).
- Major emphasis on *simplicity* – both in setting up application networks and in reasoning about them.
- *Lot's of fun* to be had – but still some work to do.

Examples

Bag-of-tasks programming

- When a problem can easily be split into independent sub-tasks
- Concurrent versions are very easy
 - But you can run into load-imbalance if your task size is too big

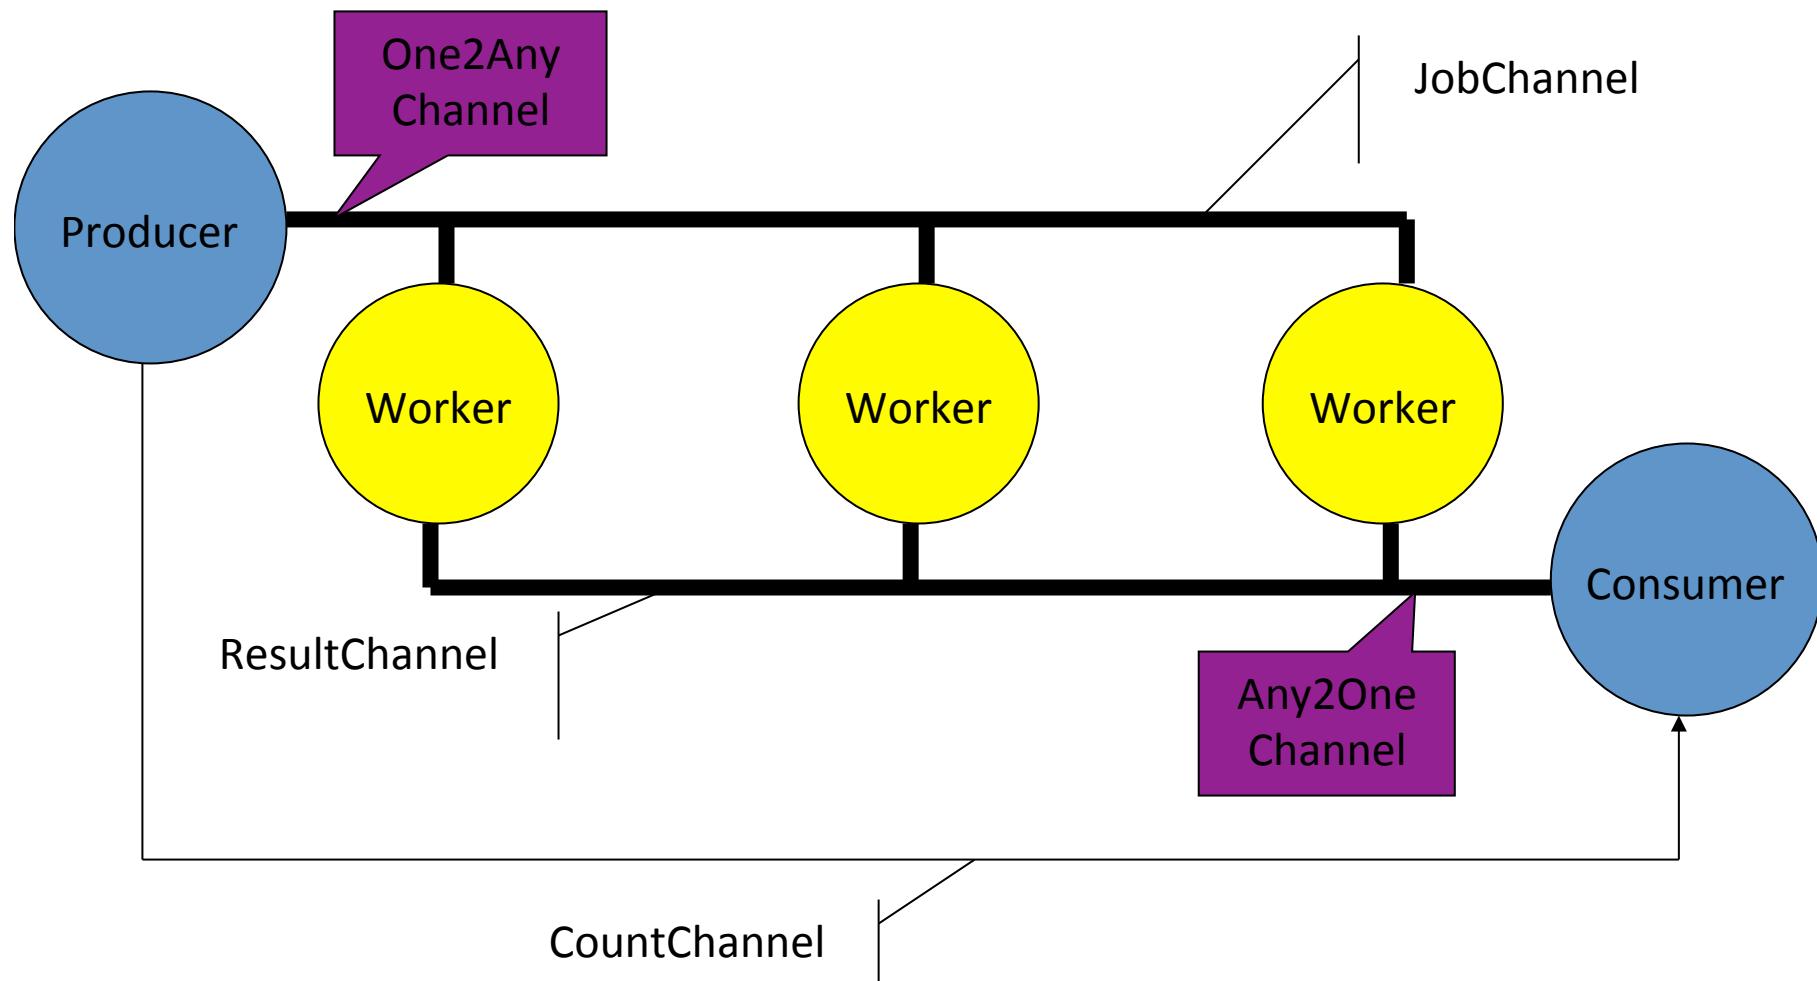
N-Queens Problem



Approach

- Generate a set of subtrees
- Put each subtree in a bag
 - Or simply make a list – it does not take long
- Have workers take tasks from the bag
 - And place results in another bag
- Sum the results
 - Or just let the bag do the summing

N-Queens example



Implementation

- **Master**

```
pool=[]
```

```
for i in range(2):
```

```
    for j in range(n):
```

```
        if(board(i,j)):
```

```
            pool.append([i,j])
```

```
CountChannel.Write(len(pool))
```

```
for x in pool:
```

```
    JobChannel.write(x)
```

Implementation

- **Worker**

```
while(true):  
    coor=JobChannel.read()  
    valid=checkboard(i,j)  
    ResultChannel.write(valid)
```

Implementation

- **Consumer**

```
n = CountChannel.read()
```

```
sum = 0
```

```
for i in range(n):
```

```
    sum = sum + ResultChannel.read()
```

```
print sum
```

Two concurrent web servers

- Two basic designs for a concurrent web server
 - High throughput web server
 - High performance web server

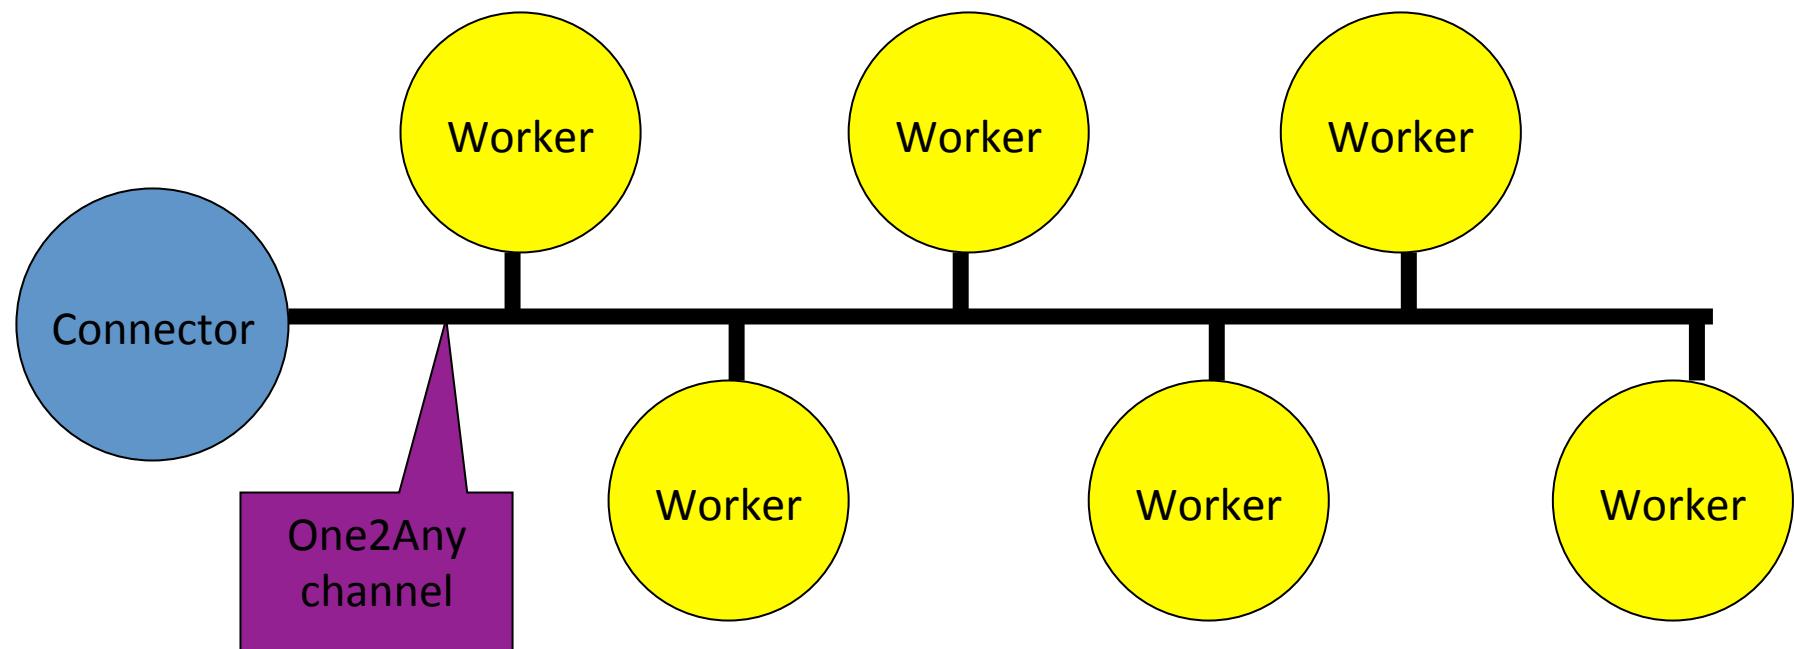
High throughput web server

- We desire a web server that can handle many concurrent sessions
- Each request is an operation in its own right
 - So each request can be implemented as a process in its own right
 - This is the same approach that is used in common web-servers

High throughput web server

- Two basic approaches to load-balancing
 - Use a semaphore to protect the listen/accept calls
 - Have one thread do all listen/accept calls and then pass the new socket to a worker thread
- The latter is far preferable

High throughput webserver



Psudo Code

Master:

```
x = new socket()  
bind(x, 80)  
while(true):  
    listen(x,1)  
    y=accept(x)  
    chan ! y
```

Worker:

```
while(true):  
    chan ? y  
    service(y)
```

High performance web server

- We desire a web server that can service a single webpage very quickly
- http 1.1 has support for 'keep-alive' and pipelined requests
 - So that when parsing a html page a client can send multiple requests towards the same page
 - Multiple images i.e.

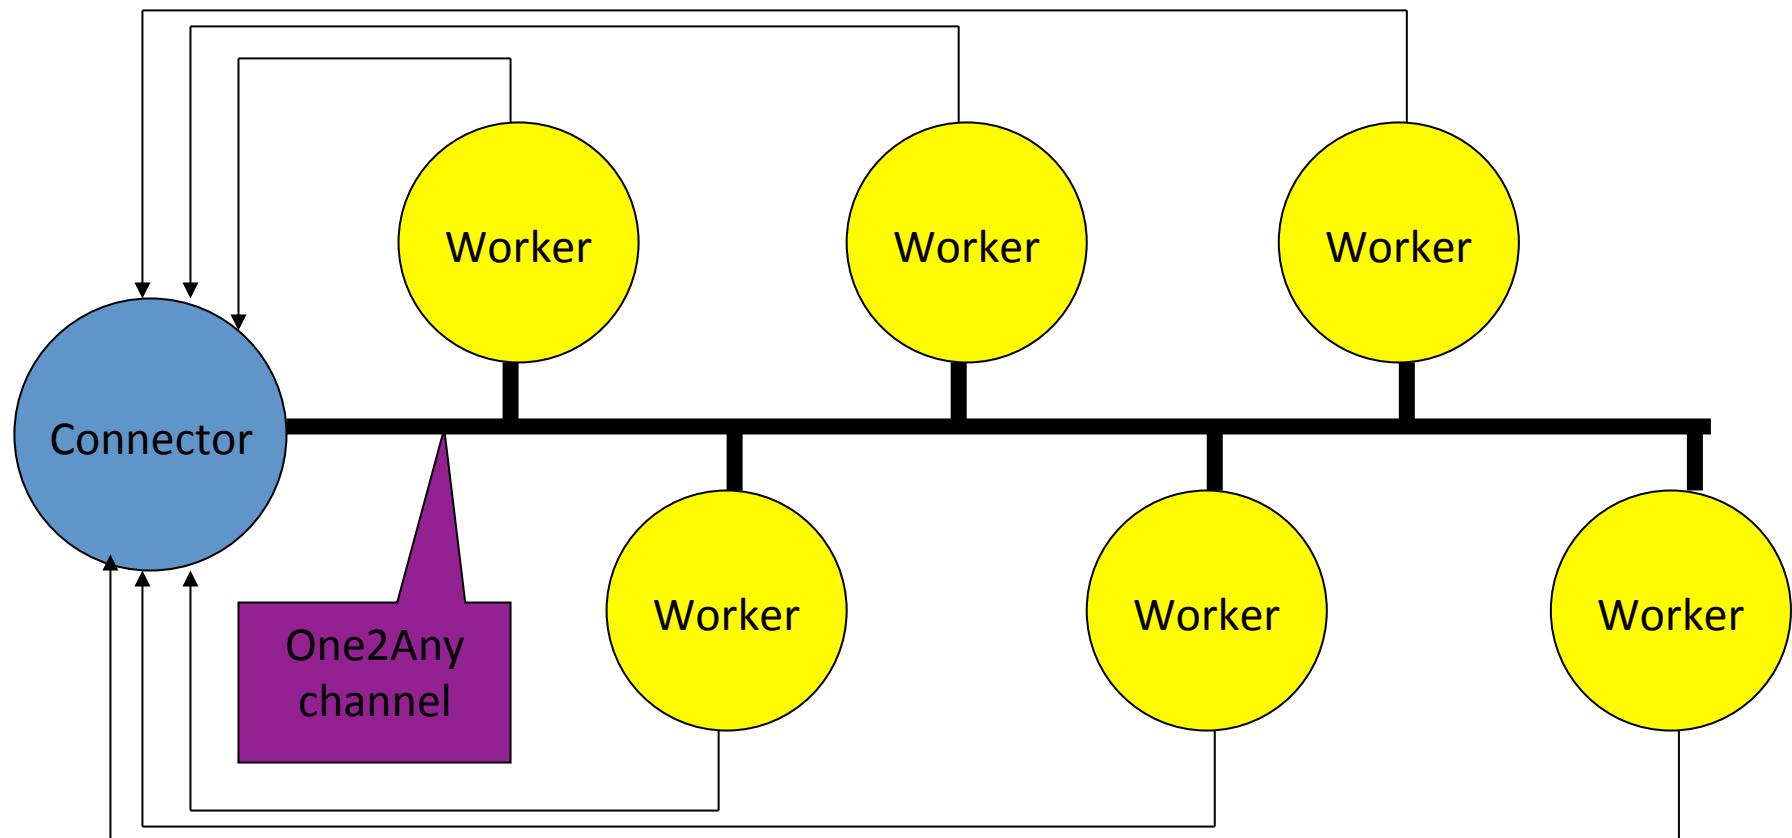
Approach

- Similar to the high throughput
 - But instead of passing sockets to the workers we pass requests
 - But the workers cannot write the answer directly to the socket any longer
- Thus the replies will have to be sent through the master

Approach

- We can use the same basic approach to distribute the requests
- But then we need each process to pass its results to the master who can then send it back

High performance web server



Psudo Code

Master:

```
x = new socket()
```

```
bind(x, 80)
```

```
while(true):
```

```
    listen(x,1)
```

```
    y=accept(x)
```

```
PRIALT:
```

```
    y ? req; chan ! req
```

```
    w[0] ? ans; y ! ans
```

```
...
```

```
    w[n] ? ans; y ! ans
```

Worker:

```
while(true):
```

```
    chan?y
```

```
    chan ! service(y)
```